

## 2. Distributed DBMS Architecture

---

Chapter 4

# Distributed DBMS Architecture

# Outline

---

- ❖ Introduction
- ❖ Top-Down Design of DDBMS Architecture
  - ◆ Schema and Distribution Transparency
- ❖ Bottom-up Design of DDBMS Architecture
  - ◆ Architectural Alternatives for DDBMSs
  - ◆ Reference Architectures for a DDBMS
- ❖ Global Directory/Dictionary

# Outline

---

## Introduction

- ❖ Top-Down Design of DDBMS Architecture
  - ◆ Schema and Distribution Transparency
- ❖ Bottom-up Design of DDBMS Architecture
  - ◆ Architectural Alternatives for DDBMSs
  - ◆ Reference Architectures for a DDBMS
- ❖ Global Directory/Dictionary

# Introduction

---

- ❖ Architecture defines the structure of the system
  - ◆ components identified
  - ◆ functions of each component defined
  - ◆ interrelationships and interactions between components defined

# Reference Model（参考模型）

---

A **conceptual framework** whose purpose is to divide standardization work into manageable pieces and to show at a general level how these pieces are related to one another.

# Three Approaches to Define a Reference Model

---

## ① Component-based

- Components of the system are defined together with the interrelationships between components.
- Good for design and implementation of the system.

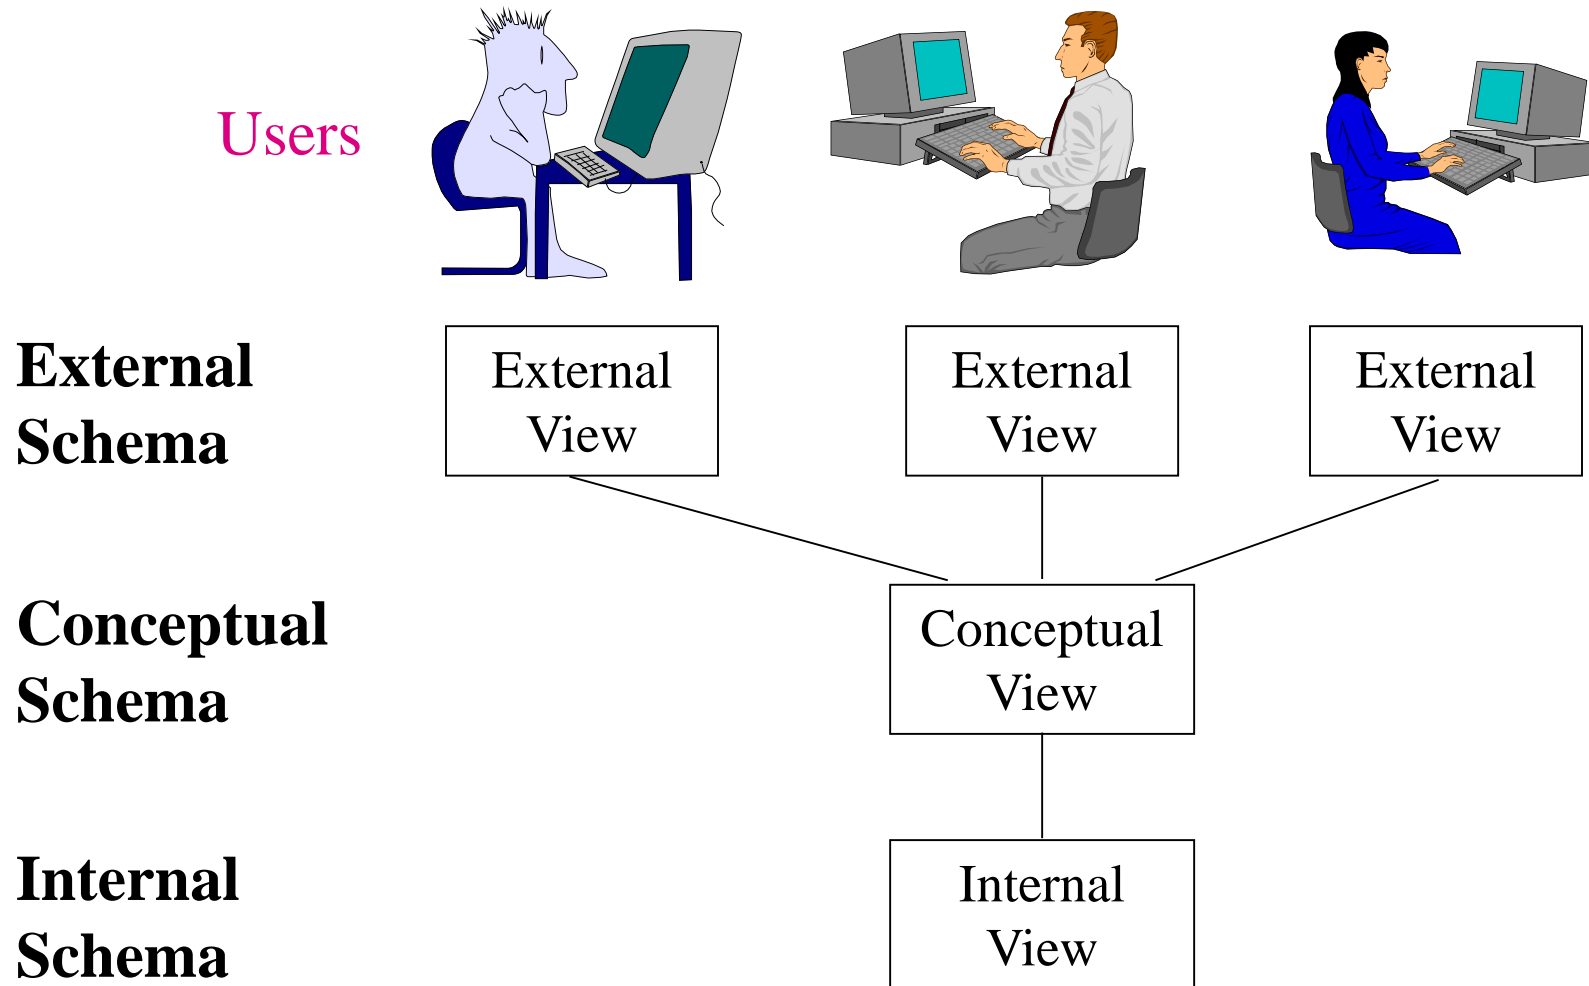
## ② Function-based

- Classes of users are identified together with the functionality that the system will provide for each class.
- The objectives of the system are clearly identified. But how do you achieve these objectives?

## ③ Data-based

- Identify different types of data and specify the functional units that will realize and/or use data according to these views.
- The ANSI/SPARC architecture discussed below belongs to this category.

# ANSI/SPARC Data-based Architecture



# Conceptual Schema (概念模式)

**RELATION EMP [**

**KEY = {ENO}**

**ATTRIBUTES = {**

ENO : CHARACER(9)

ENAME : CHARACER(15)

TITLE : CHARACER(10)

**}**

**]**

**RELATION PAY [**

**KEY = {TITLE}**

**ATTRIBUTES = {**

TITLE : CHARACER(10)

SAL : NUMERIC(6)

**}**

**]**

**RELATION PROJECT [**

**KEY = {PNO}**

**ATTRIBUTES = {**

PNO : CHARACER(7)

PNAME : CHARACER(20)

BUDGET : NEMERIC(7)

**}**

**]**

**RELATION ASG [**

**KEY = {ENO,PNO}**

**ATTRIBUTES = {**

ENO : CHARACER(9)

PNO : CHARACER(7)

RESP : CHARACER(10)

DUR : NUMERIC(6)


**}**

**]**



# Internal Schema (内部模式)

```
RELATION EMP [  
  KEY = {ENO}  
  ATTRIBUTES = {  
    ENO      : CHARACER(9)  
    ENAME    : CHARACER(15)  
    TITLE    : CHARACER(10)  
  }  
]
```



```
INTERNAL_REL EMP [  
  INDEX ON ENO CALL EMINX  
  FIEDLS = {  
    HEADER : BYTE(1)  
    ENO     : BYTE(9)  
    ENAME   : BYTE(15)  
    TITLE   : BYTE(10)  
  }  
]
```

# External View（外部模式） – Example 1

Create a BUDGET view from the PROJ relation

```
CREAT VIEW      BUDGET(PNAME, BUD)  
AS      SELECT  PNAME, BUDGET  
          FROM    PROJ
```

```
RELATION PROJECT [  
  KEY = {PNO}  
  ATTRIBUTES = {  
    PNO      : CHARACER(7)  
    PNAME    : CHARACER(20)  
    BUDGET   : NEMERIC(7)  
  }  
]
```

# External View (外部模式) – Example 2

Create a Payroll view from relations EMP and PAY

```
CREATE VIEW      PAYROLL(ENO, ENAME, SAL)
AS      SELECT  EMP.ENO, EMP.ENAME, PAY.SAL
              FROM    EMP, PAY
              WHERE   EMP.TITLE = PAY.TITLE
```

```
RELATION EMP [
  KEY = {ENO}
  ATTRIBUTES = {
    ENO      : CHARACER(9)
    ENAME    : CHARACER(15)
    TITLE    : CHARACER(10)
  }
]
```

```
RELATION PAY [
  KEY = {TITLE}
  ATTRIBUTES = {
    TITLE : CHARACER(10)
    SAL   : NUMERIC(6)
  }
]
```

# Outline

---

- ❖ Introduction

- ☞ **Top-Down Design of DDBMS Architecture**

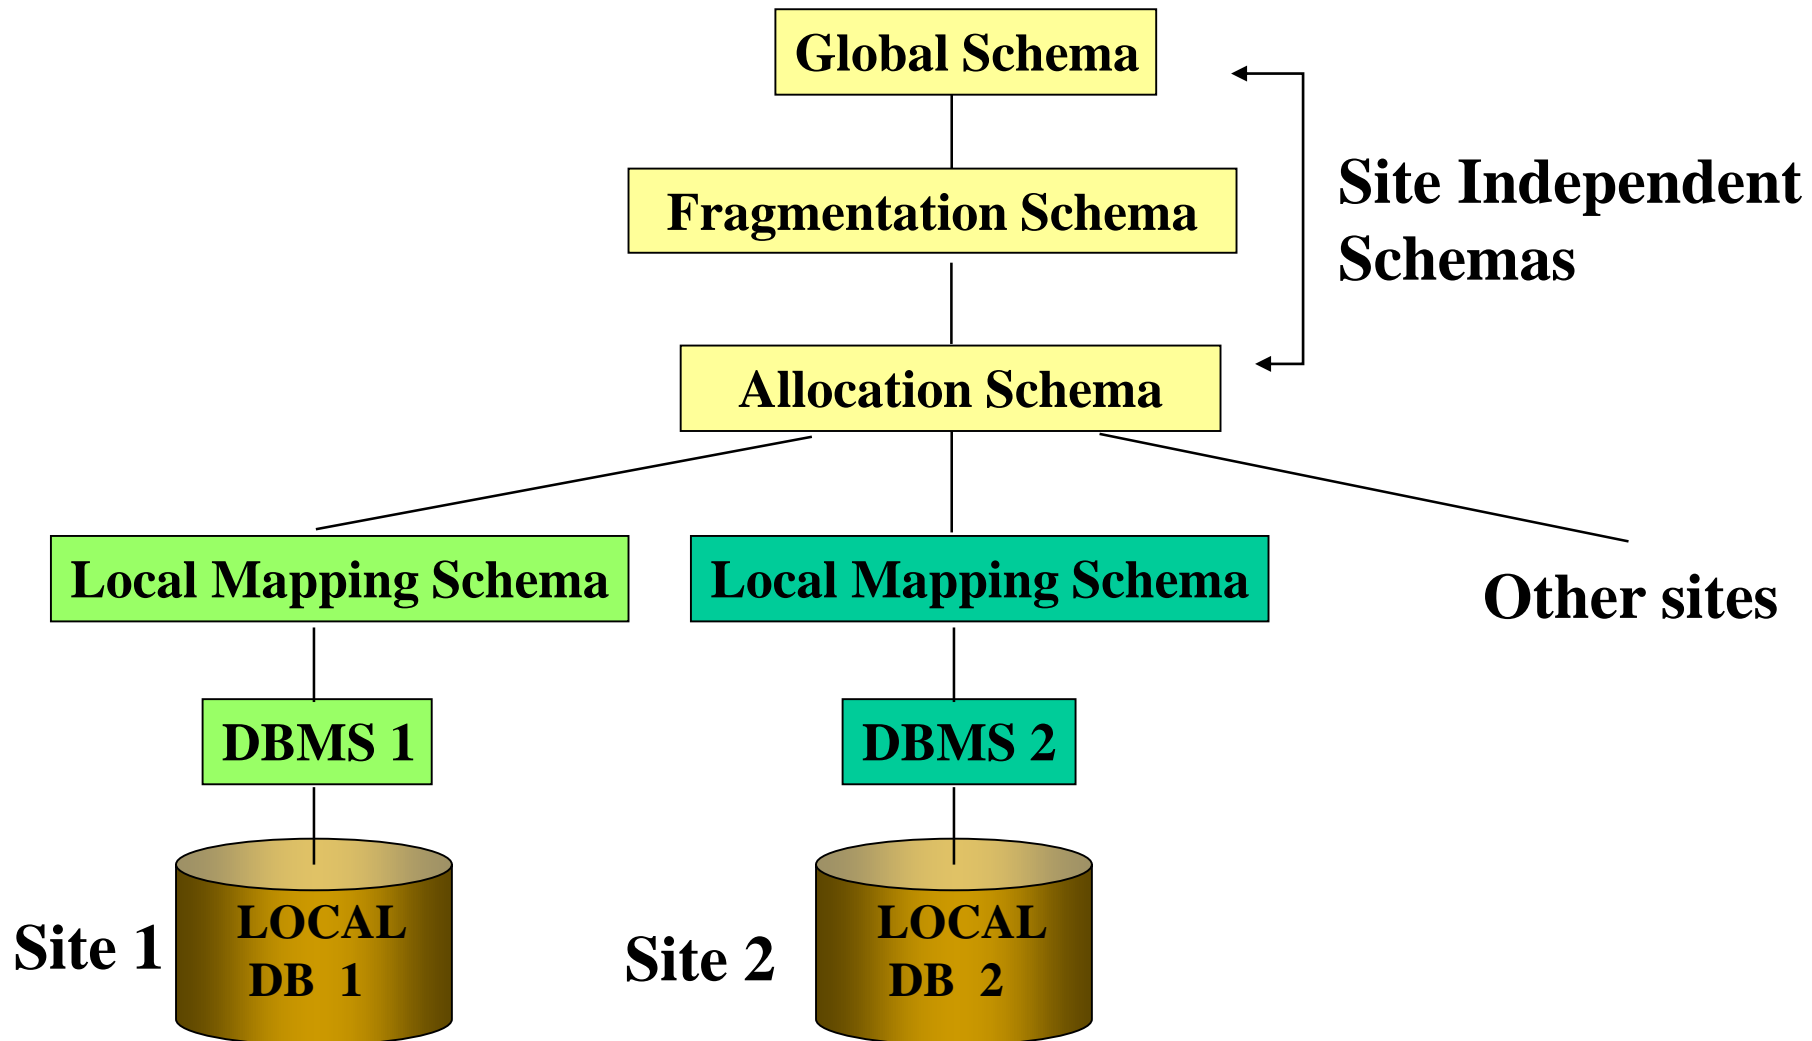
- ◆ Schema and Distribution Transparency

- ❖ **Bottom-up Design of DDBMS Architecture**

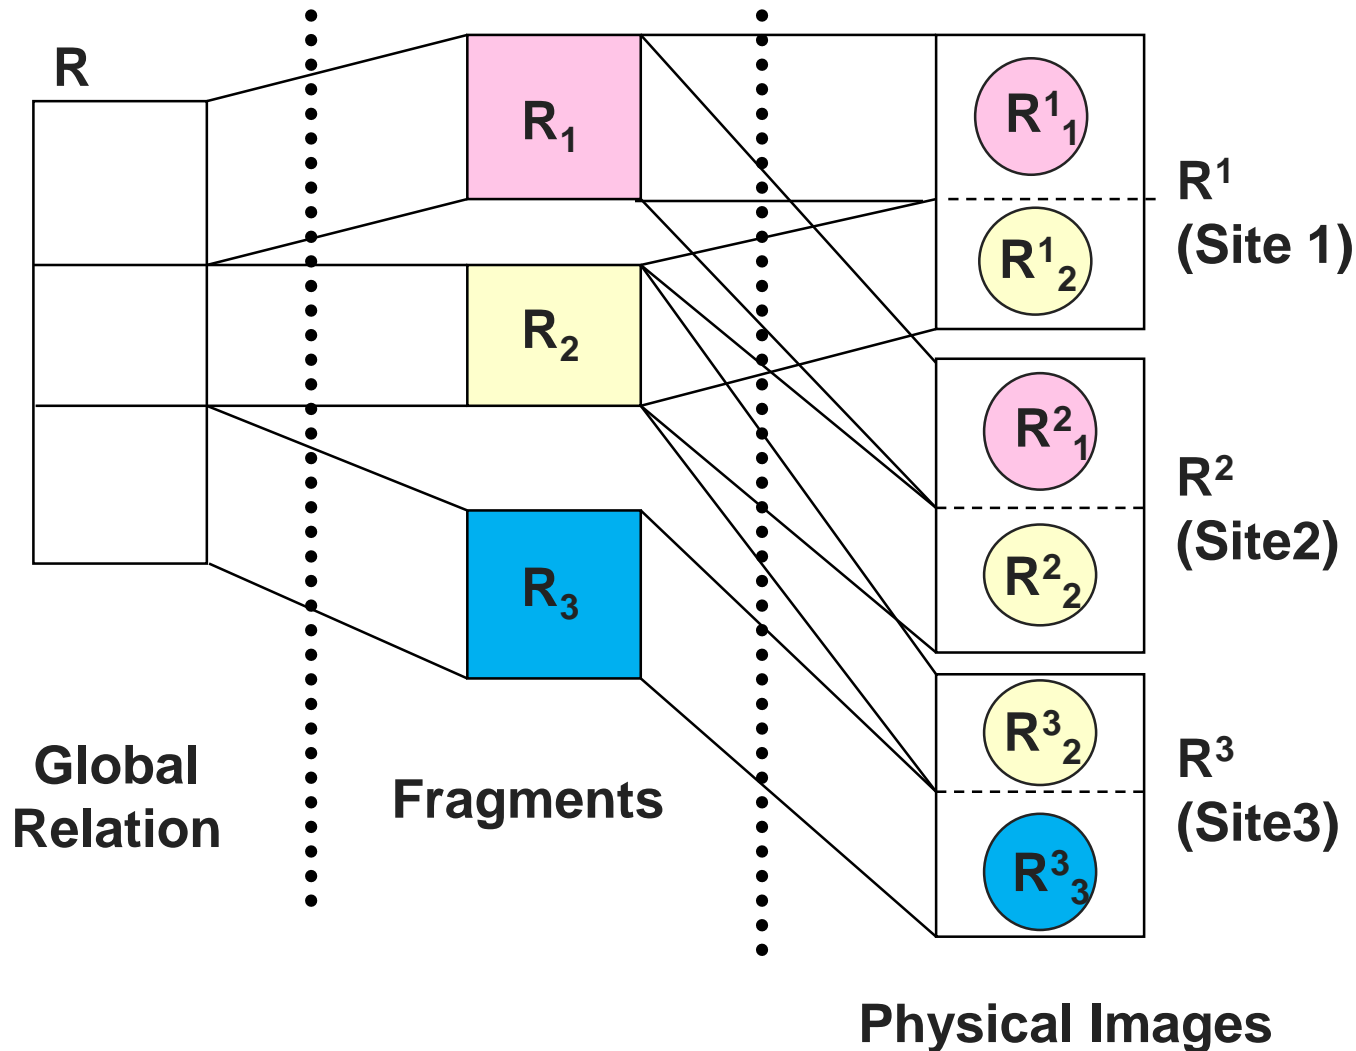
- ◆ Architectural Alternatives for DDBMSs
- ◆ Reference Architectures for a DDBMS

- ❖ **Global Directory/Dictionary**

# Top-Down Classical DDBMS Architecture



# Global Relations, Fragments, and Physical Images



# DDBMS Schemas

- ❖ **Global Schema:** a set of global relations as if database were not distributed at all
- ❖ **Fragmentation Schema:** global relation is split into “non-overlapping” (logical) fragments.  $1:n$  mapping from relation  $R$  to fragments  $R_i$ .
- ❖ **Allocation Schema:**  $1:1$  or  $1:n$  (redundant) mapping from fragments to sites. All fragments corresponding to the same relation  $R$  at a site  $j$  constitute the physical image  $R^j_i$ . A copy of a fragment is denoted by  $R^j_i$ .
- ❖ **Local Mapping Schema:** a mapping from physical images to physical objects, which are manipulated by local DBMSs.

# Motivation for this Architecture

---

- ❖ Separating the concept of data fragmentation from the concept of data allocation
- ❖ Fragmentation transparency
- ❖ Location transparency
- ❖ Explicit control of redundancy
- ❖ Independence from local databases allows local mapping transparency



# Rules for Data Fragmentation

---

- ❖ Completeness

All the data of the global relation must be mapped into the fragments.


- ❖ Reconstruction

It must always be possible to reconstruct each global relation from its fragments.

- ❖ Disjointedness


It is convenient that fragments are disjoint, so that the replication of data can be controlled explicitly at the allocation level.

# Types of Data Fragmentation



## Vertical Fragmentation

- Projection on relation (subset of attributes)
- Reconstruction by join
- Updates require no tuple migration



## Horizontal Fragmentation

- Selection on relation (subset of tuples)
- Reconstruction by union
- Updates may require tuple migration

## Mixed Fragmentation

- A fragment is a Select-Project query on relation

# Horizontal Fragmentation (水平划分)

- ❖ Partitioning the tuples of a global relation into subsets

Example:

**Supplier (SNum, Name, City)**

Horizontal Fragmentation can be:

**Supplier<sub>1</sub> =  $\sigma_{\text{City} = \text{"HK"}} \text{Supplier}$**

**Supplier<sub>2</sub> =  $\sigma_{\text{City} \neq \text{"HK"}} \text{Supplier}$**

Reconstruction is possible:

**Supplier = Supplier<sub>1</sub>  $\cup$  Supplier<sub>2</sub>**

- ❖ The set of predicates defining all the fragments must be complete, and mutually exclusive

# Derived Horizontal Fragmentation

$\text{Supplier}_1 = \sigma_{\text{City} = \text{``HK''}} \text{Supplier}$   
 $\text{Supplier}_2 = \sigma_{\text{City} \neq \text{``HK''}} \text{Supplier}$

- ❖ The horizontal fragmentation is derived from the horizontal fragmentation of another relation

Example:

**Supply (SNum, PNum, DeptNum, Quan)**

**SNum** is a supplier number

**Supply<sub>1</sub> = Supply ⋈<sub>SNum=SNum</sub> Supplier<sub>1</sub>**

**Supply<sub>2</sub> = Supply ⋈<sub>SNum=SNum</sub> Supplier<sub>2</sub>**

 **semijoin operation**

The predicates defining derived horizontal fragments are:

**(Supply.SNum = Supplier.SNum) and (Supplier.City = ``HK``)**

**(Supply.SNum = Supplier.SNum) and (Supplier.City != ``HK``)**

# Vertical Fragmentation (垂直划分)

- ❖ The vertical fragmentation of a global relation is the subdivision of its attributes into groups; fragments are obtained by projecting the global relation over each group

Example

**EMP (ENum, Name, Sal, Tax, MNum, DNum)**

A vertical fragmentation can be

**$EMP_1 = \Pi_{ENum, Name, MNum, DNum} EMP$**

**$EMP_2 = \Pi_{ENum, Sal, Tax} EMP$**

Reconstruction:

**$EMP = EMP_1 \bowtie_{ENum = ENum} EMP_2$**

# Distribution Transparency (分布透明)

- ❖ Different levels of distribution transparency can be provided by DDBMS for applications.

## A Simple Application

**Supplier(SNum, Name, City)**

Horizontally fragmented into:

**Supplier<sub>1</sub> =  $\sigma_{City = \text{``HK"}}$  Supplier at Site1**

**Supplier<sub>2</sub> =  $\sigma_{City \neq \text{``HK"}}$  Supplier at Site2, Site3**

**Application:**


*Read the supplier number from the user and return the name of the supplier with that number.*

# Level 1 of Distribution Transparency

Fragmentation transparency:

```
read (terminal, $SNum);  
  Select      Name into $Name  
  from        Supplier  
  where       SNum = $SNum;  
write (terminal, $Name).
```

Read the supplier number from the user  
and return the name of the supplier with  
that number.




The DDBMS interprets the database operation by accessing the databases at different sites in a way which is completely determined by the system.

# Level 2 of Distribution Transparency

## Location Transparency (but fragmentation not)

```
read (terminal, $SNum);  
  Select    Name into $Name  
  from      Supplier1  
  where      SNum = $SNum;  
If not FOUND then  
  Select    Name into $Name  
  from      Supplier2  
  where      SNum = $SNum;  
write (terminal, $Name).
```




The application is independent from changes in allocation schema, but is not independent from changes to fragmentation schema.



# Level 3 of Distribution Transparency

## Local Mapping Transparency (but distribution not)

```
read (terminal, $SNum);  
  Select Name into $Name  
  from S1.Supplier1  
  where SNum = $SNum;  
If not FOUND then  
  Select Name into $Name  
  from S3.Supplier2  
  where SNum = $SNum;  
write (terminal, $Name).
```




The applications have to specify both the fragment names and the sites where they are located. The mapping of database operations specified in applications to those in DBMSs at sites is transparent.

# Level 4 of Distribution Transparency

❖ No transparency at all !!!

```
read (terminal, $SNum);  
  $SupIMS($Snum,$Name,$Found) at S1;  
If not FOUND then  
  $SupCODASYL($Snum,$Name,$Found) at S3;  
write (terminal, $Name).
```



# Distribution Transparency for Updates

## Difficulties

- broadcasting updates to all copies
- migration of tuples because of change of fragment defining attributes

$$\text{EMP1} = \Pi_{\text{Enum,Name,Sal,Tax}} \sigma_{\text{Dnum} \leq 10} (\text{EMP})$$

$$\text{EMP2} = \Pi_{\text{Enum,Mnum,Dnum}} \sigma_{\text{Dnum} \leq 10} (\text{EMP})$$

$$\text{EMP3} = \Pi_{\text{Enum,Name,Dnum}} \sigma_{\text{Dnum} > 10} (\text{EMP})$$

$$\text{EMP4} = \Pi_{\text{Enum,Mnum,Sal,Tax}} \sigma_{\text{Dnum} > 10} (\text{EMP})$$

**EMP1**

Enum	Name	Sal	Tax
100	Ann	100	10

**EMP2**

Enum	Mnum	Dnum
100	20	3

Update Dnum=15  
for Employee with  
Enum=100

**EMP3**

Enum	Name	Dnum
100	Ann	15

**EMP4**

Enum	Mnum	Sal	Tax
100	20	100	10

# An Update Application

```
UPDATE EMP  
SET Dnum = 15  
WHERE Enum = 100;
```

With Level 1  
Fragmentation  
Transparency

With Level 2 Location  
Transparency only

```
Select Name, Tax, Sal into $Name, $Sal, $Tax  
From EMP 1
```

```
Where Enum = 100;
```

```
Select Mnum into $Mnum  
From EMP 2
```

```
Where Enum = 100;
```

```
Insert into EMP 3 (Enum, Name, Dnum)  
(100, $Name, 15);
```

```
Insert into EMP 4 (Enum, Sal, Tax, Mnum)  
(100, $Sal, $Tax, $Mnum);
```

```
Delete EMP 1 where Enum = 100;
```

```
Delete EMP 2 where Enum = 100;
```

# Levels of Distribution Transparency

---

- ❖ Fragmentation Transparency
  - ◆ Just like using global relations
- ❖ Location Transparency
  - ◆ Need to know fragmentation schema; but no need to know where fragments are located
  - ◆ Applications access fragments (no need to specify sites where fragments are located).
- ❖ Local Mapping Transparency
  - ◆ Need to know both fragmentation and allocation schema; no need to know what the underlying local DBMSs are.
  - ◆ Applications access fragments explicitly specifying where the fragments are located.
- ❖ No Transparency
  - ◆ Need to know local DBMS query languages, and write applications using functionality provided by the Local DBMS

# On Distribution Transparency

---

- ❖ More distribution transparency requires appropriate DDBMS support, but makes end-application developers' work easy.
- ❖ The less distribution transparency, the more the end-application developer needs to know about fragmentation and allocation schemes, and how to maintain database consistency.
- ❖ There are tough problems in query optimization and transaction management that need to be tackled (in terms of system support and implementation) before fragmentation transparency can be supported.

# Some Aspects of the Classical DDBMS Architecture

---

- ❖ Distributed database technology is an “add-on” technology, and most users already have populated centralized DBMSs, whereas **top-down** design assumes implementation of new DDBMS from scratch.
- ❖ In many application environments, such as semi-structured databases, continuous streaming multimedia data, the notion of fragment is difficult to define.

# Outline

---

- ❖ Introduction
- ❖ Top-Down Design of DDBMS Architecture
  - ◆ Schema and Distribution Transparency
- ☞ Bottom-up Design of DDBMS Architecture
  - ◆ Architectural Alternatives for DDBMSs
  - ◆ Reference Architectures for a DDBMS
- ❖ Global Directory/Dictionary



# Bottom-up Distributed Architectural Models

---

- ❖ Possible ways in which multiple data(bases) are put together for sharing, which are characterized according to three dimensions
  - ◆ Distribution
  - ◆ Heterogeneous
  - ◆ Autonomy

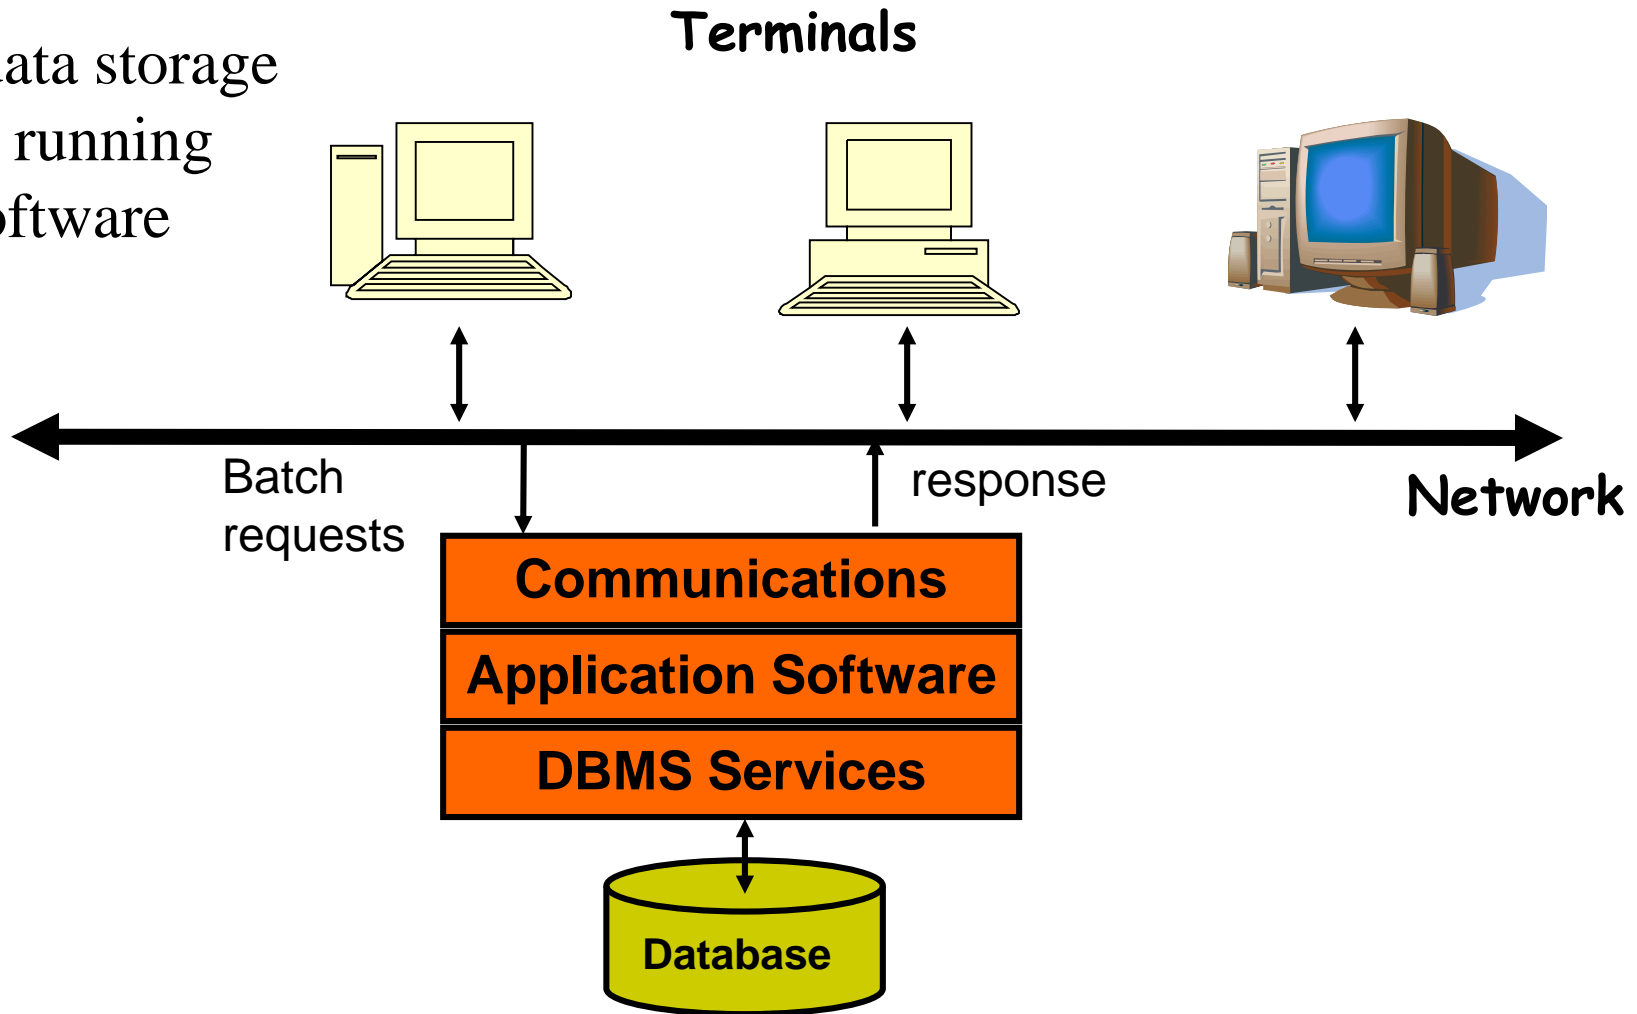
# Dimension 1: Distribution (分布)

---

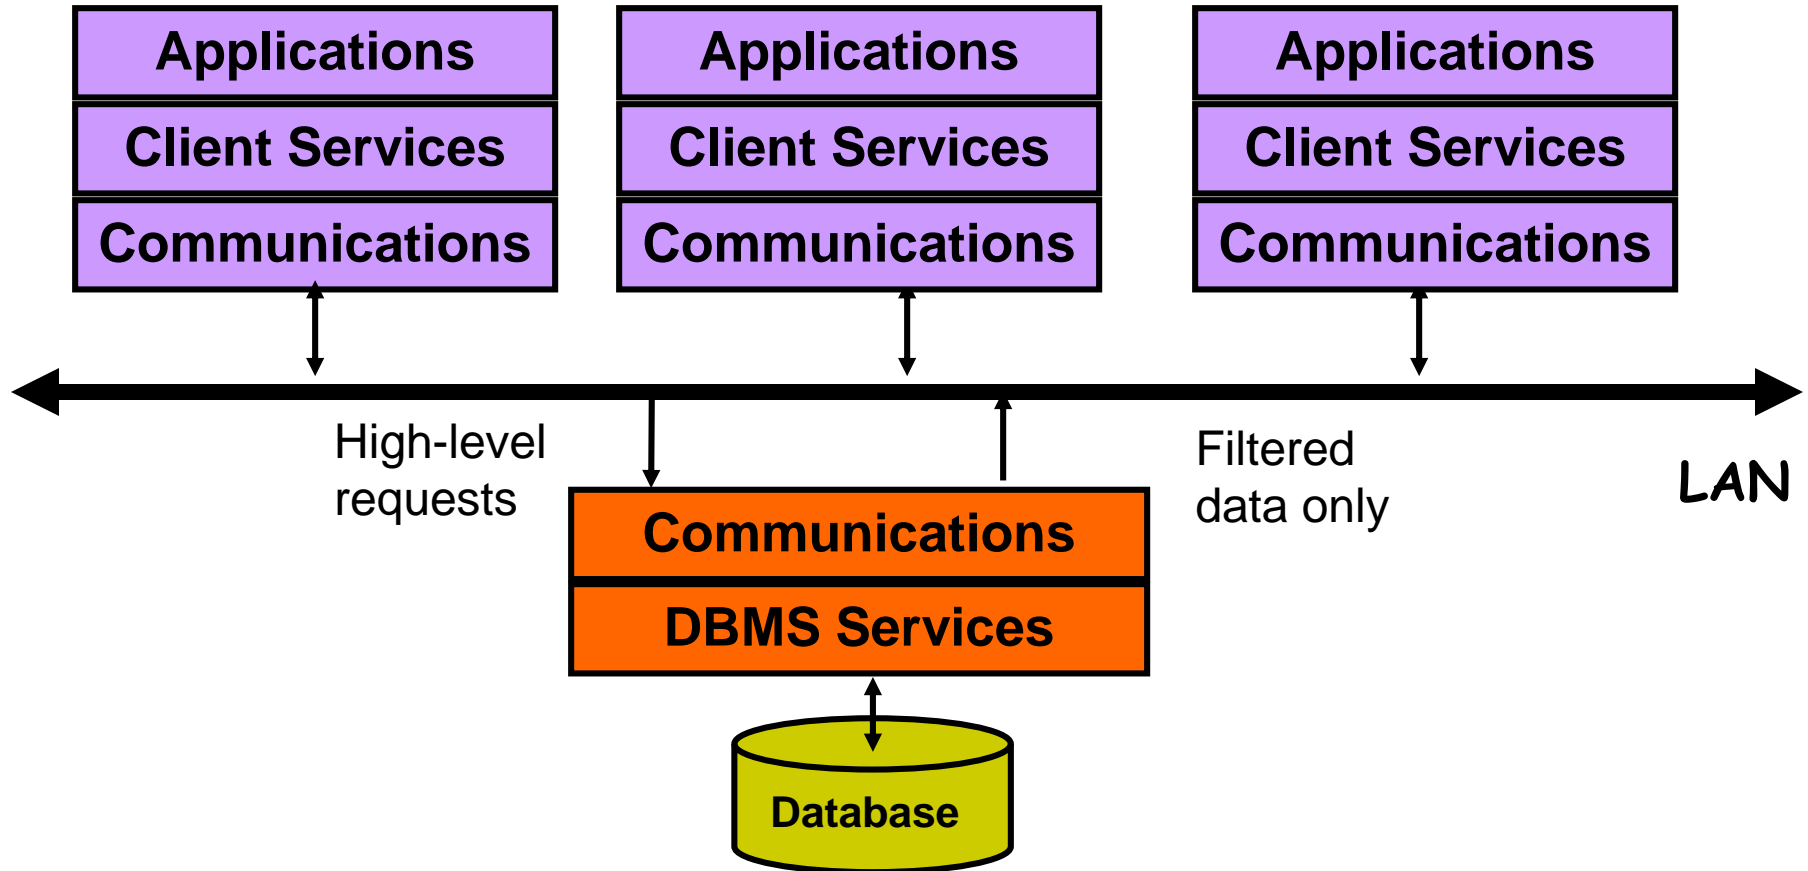
- ❖ Whether the components of the system are located on the same machine or not
  - ◆ 0 - no distribution - single site (central database)
  - ◆ 1 - client-server - distribution of DBMS functionalities
  - ◆ 2 - master-slaves - distribution of DBMS functionalities
  - ◆ 3 - peer to peer

# 0 – No Distribution (Time Sharing Access to a Central Database)

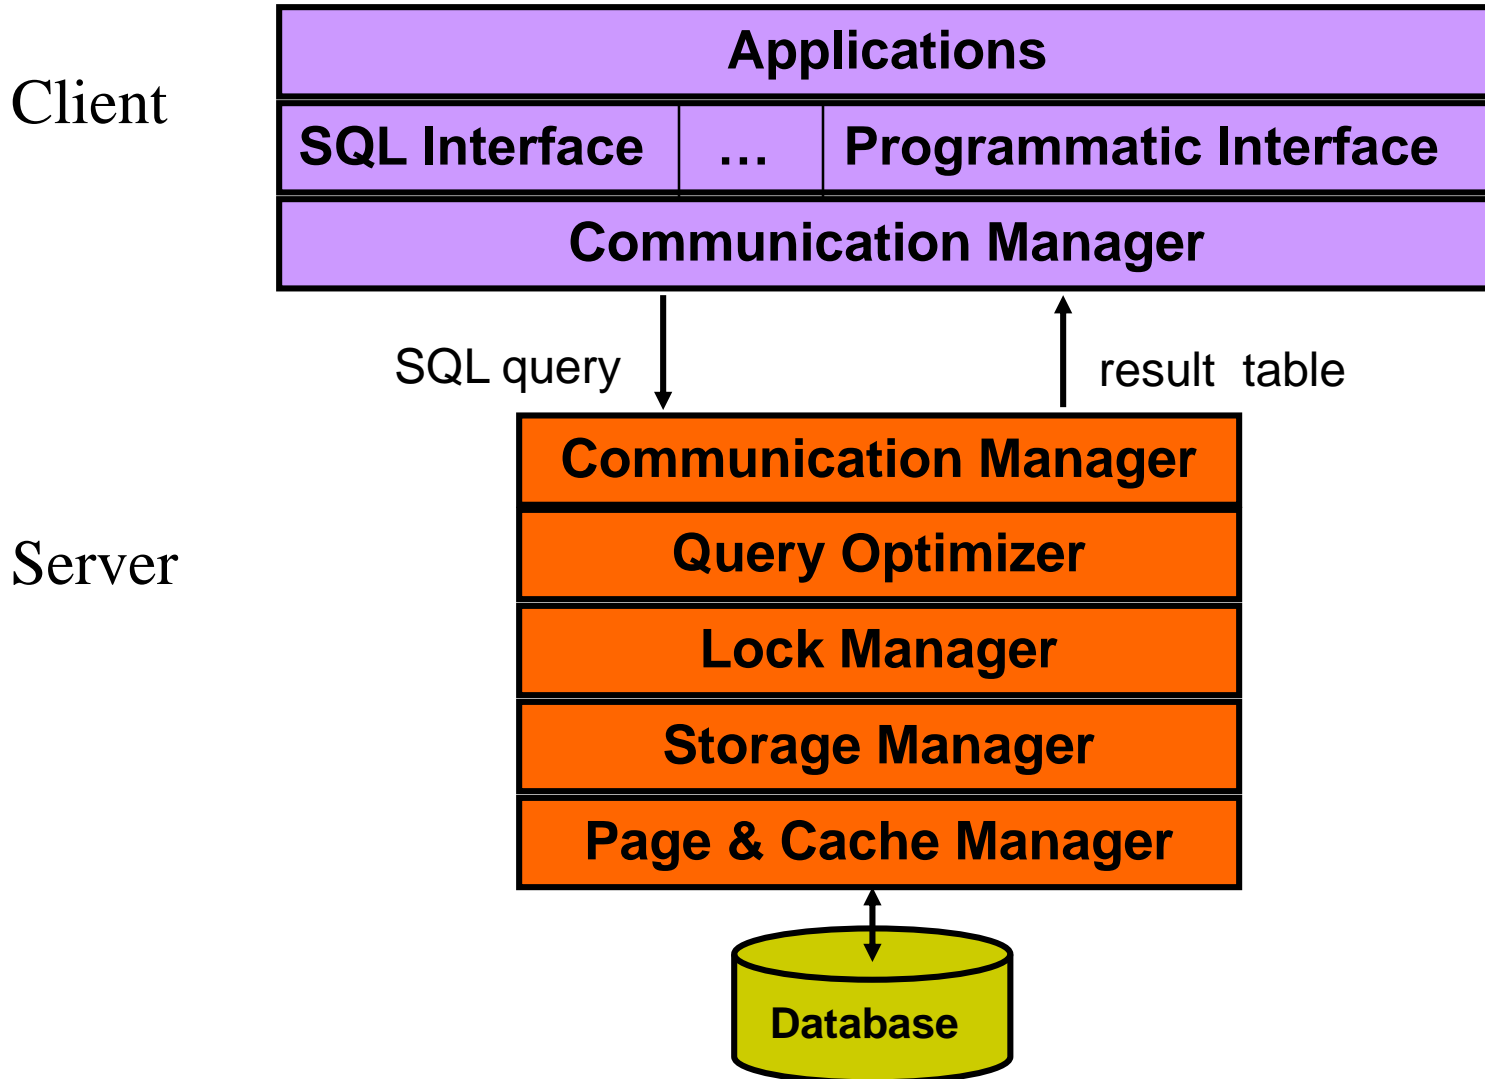
- No data storage
- Host running all software



# 0 – No Distribution (Multiple Clients / Single Server)



# Task Distribution



# Advantages of Client-Server Architectures

---

- ❖ More efficient division of labor
- ❖ Horizontal and vertical scaling of resources
- ❖ Better price/performance on client machines
- ❖ Ability to use familiar tools on client machines
- ❖ Client access to remote data (via standards)
- ❖ Full DBMS functionality provided to client workstations
- ❖ Overall better system price/performance

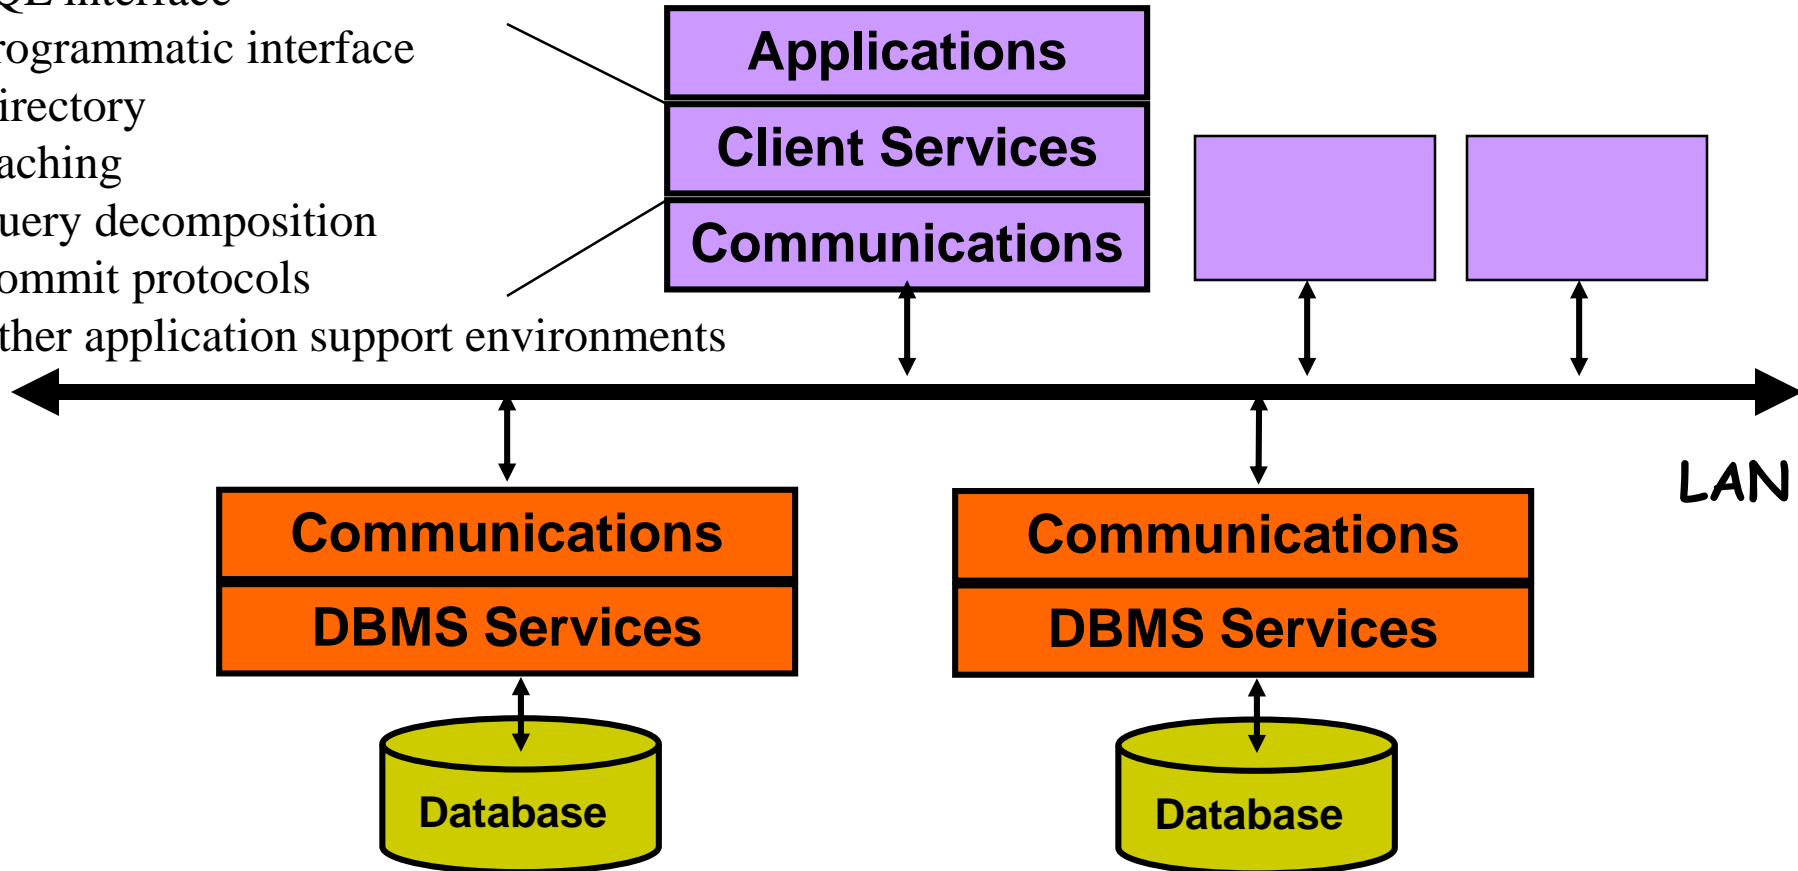
# Problems with Multiple-Clients / Single Server Architectures

---

- ❖ Server forms bottleneck
- ❖ Server forms single point of failure
- ❖ Database scaling difficult

# 1 - Multiple Clients / Multiple Servers

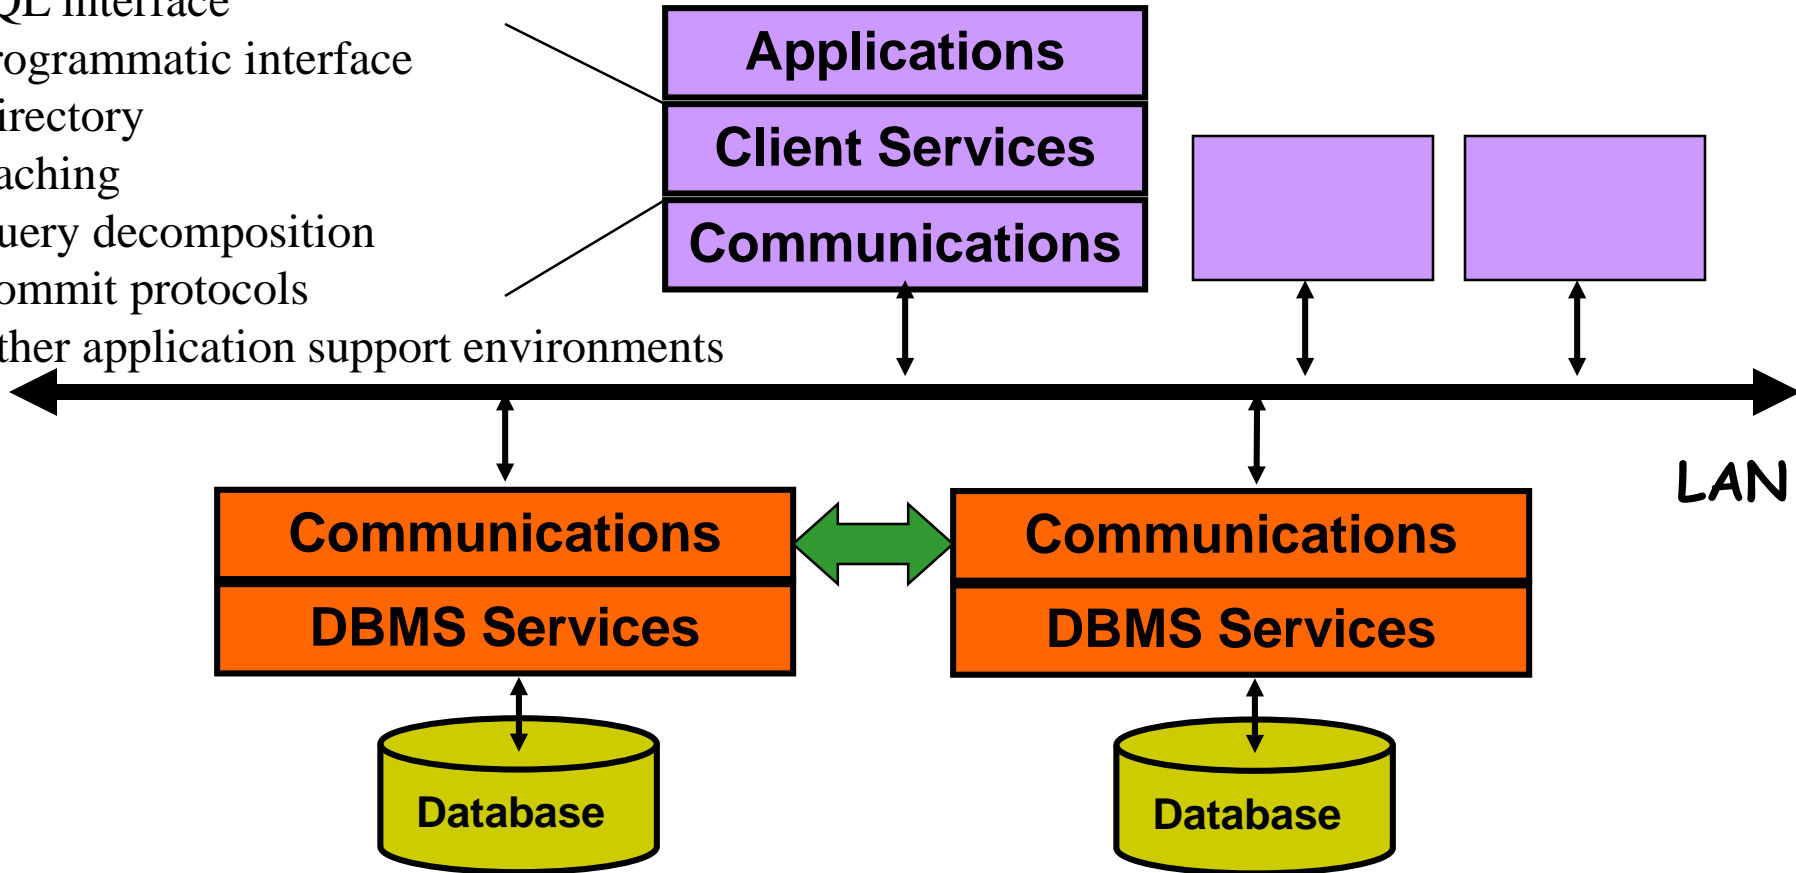
- SQL interface
- Programmatic interface
- Directory
- Caching
- Query decomposition
- Commit protocols
- Other application support environments






# Server to Server

- SQL interface
- Programmatic interface
- Directory
- Caching
- Query decomposition
- Commit protocols
- Other application support environments



# 2 – Master-Slaves Architecture

Apache  
Hadoop

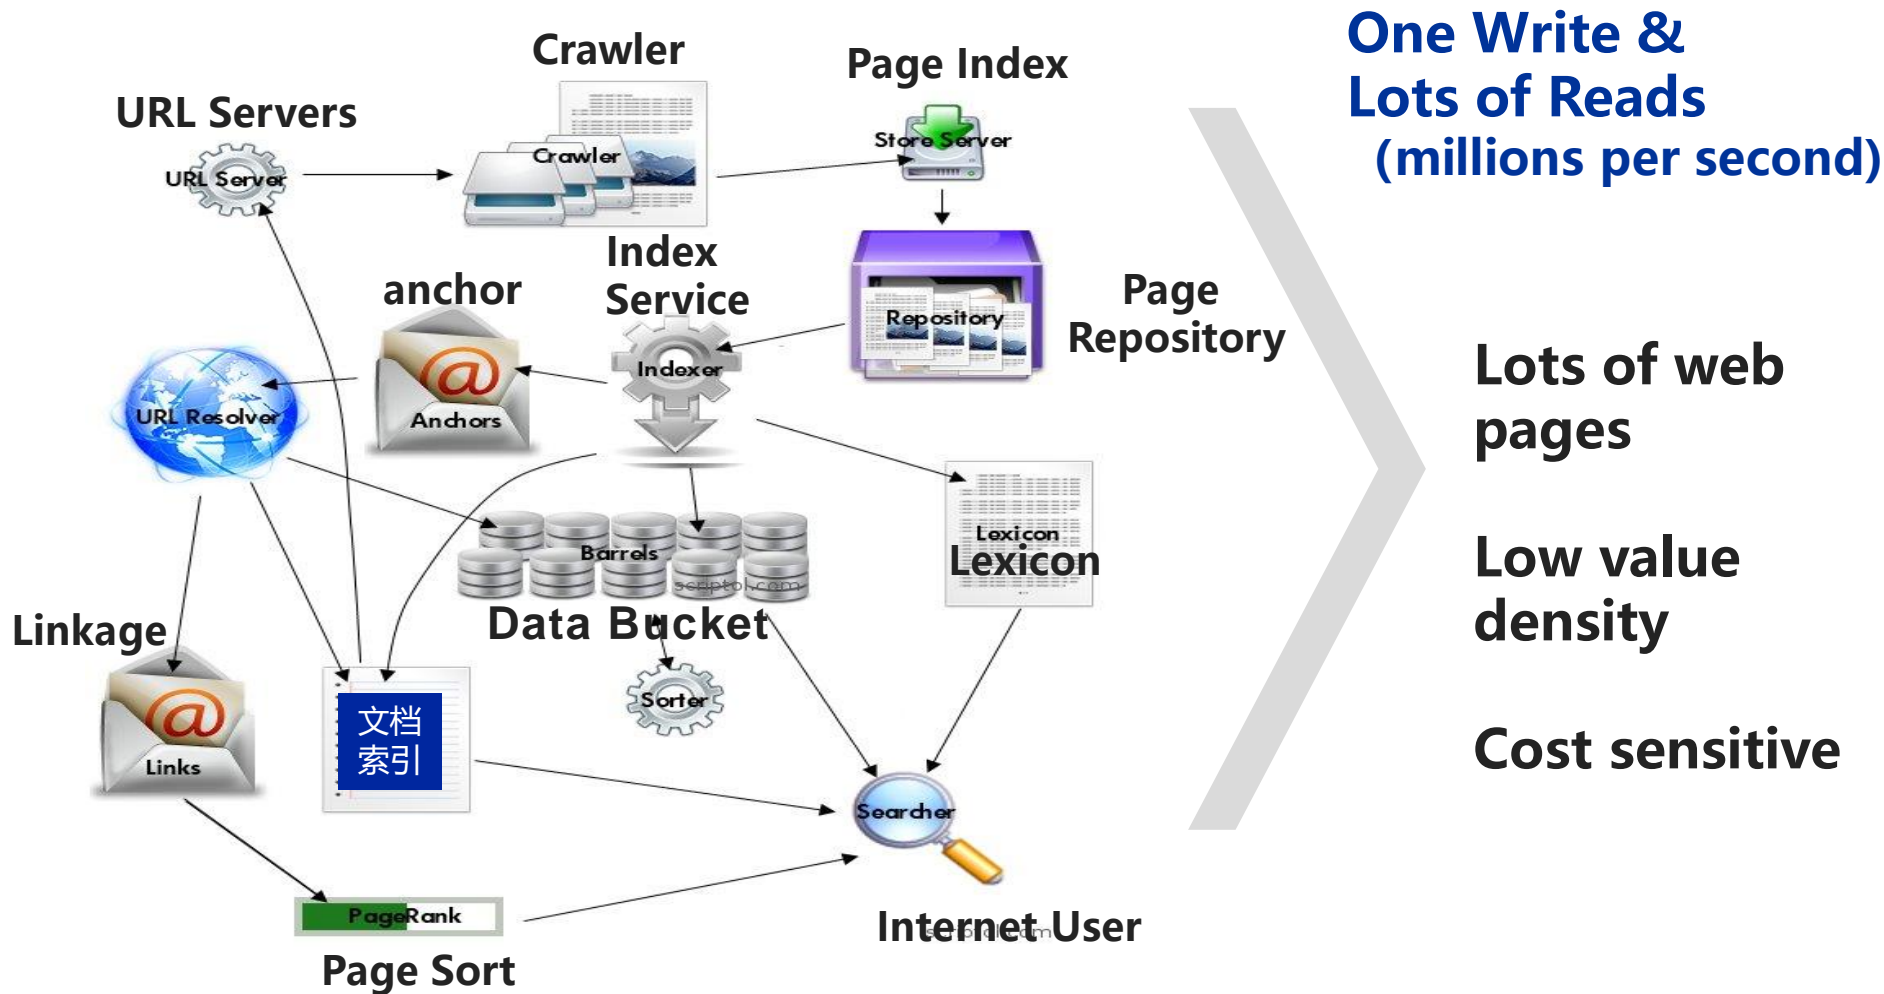


# What is Hadoop?

---


- ❖ Hadoop is a software framework for distributed processing of large datasets across large clusters of computers
  - ◆ **Large datasets** → Terabytes or petabytes of data
  - ◆ **Large clusters** → hundreds or thousands of nodes
- ❖ Hadoop is open-source implementation for **Google MapReduce** (a simple programming model)

# Google's Dream and Challenge





# Google's DIY Hardware Platform



From: Mass Data Processing Technology on Large Scale Clusters

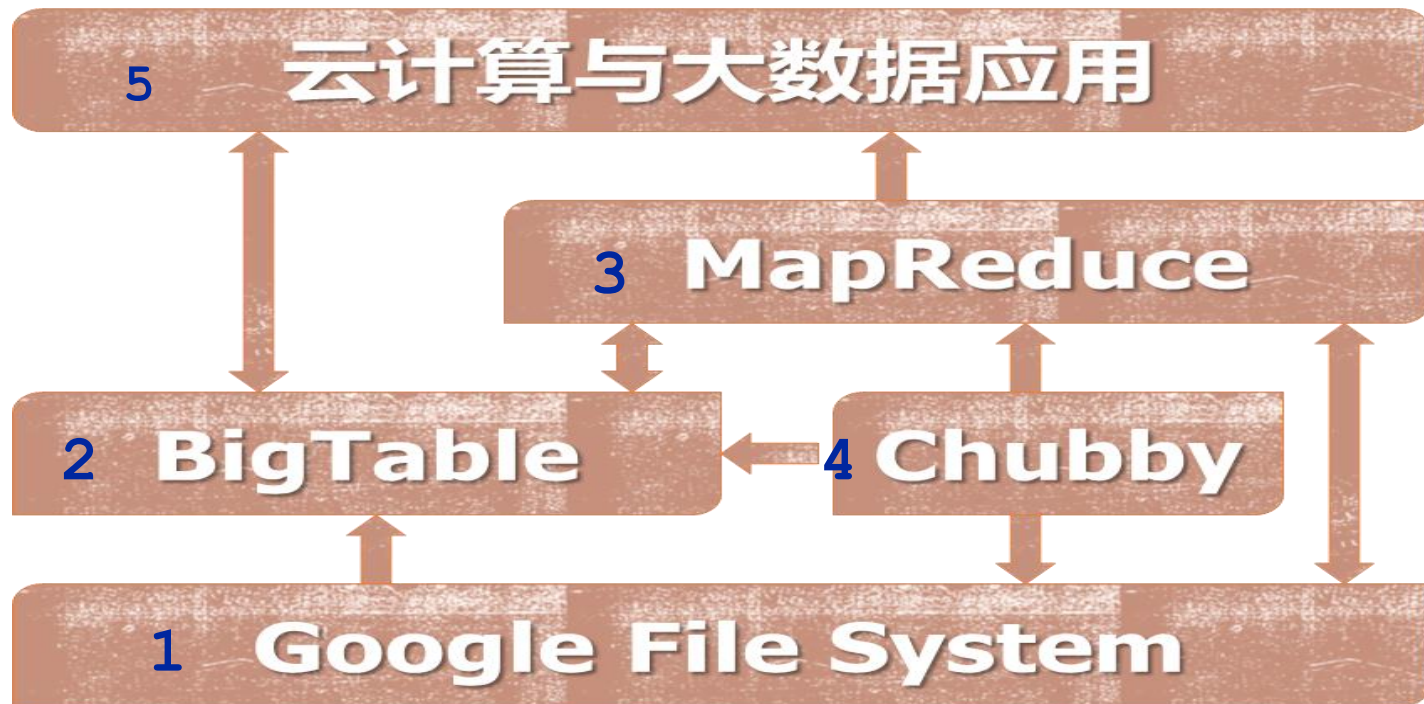
# Google's Choice

## ■ Powerful server or cheap PC?

	Server	PC
Computational Capability	Strong	Weak
Storage	Big	Small
Fault	Seldom	Frequent
Cost	Expensive	Cheap

- High volume of web pages with low value density
- Cost per query < 5 cents

# Google's Solution for Big Data




**Software stack to make up for hardware defects**



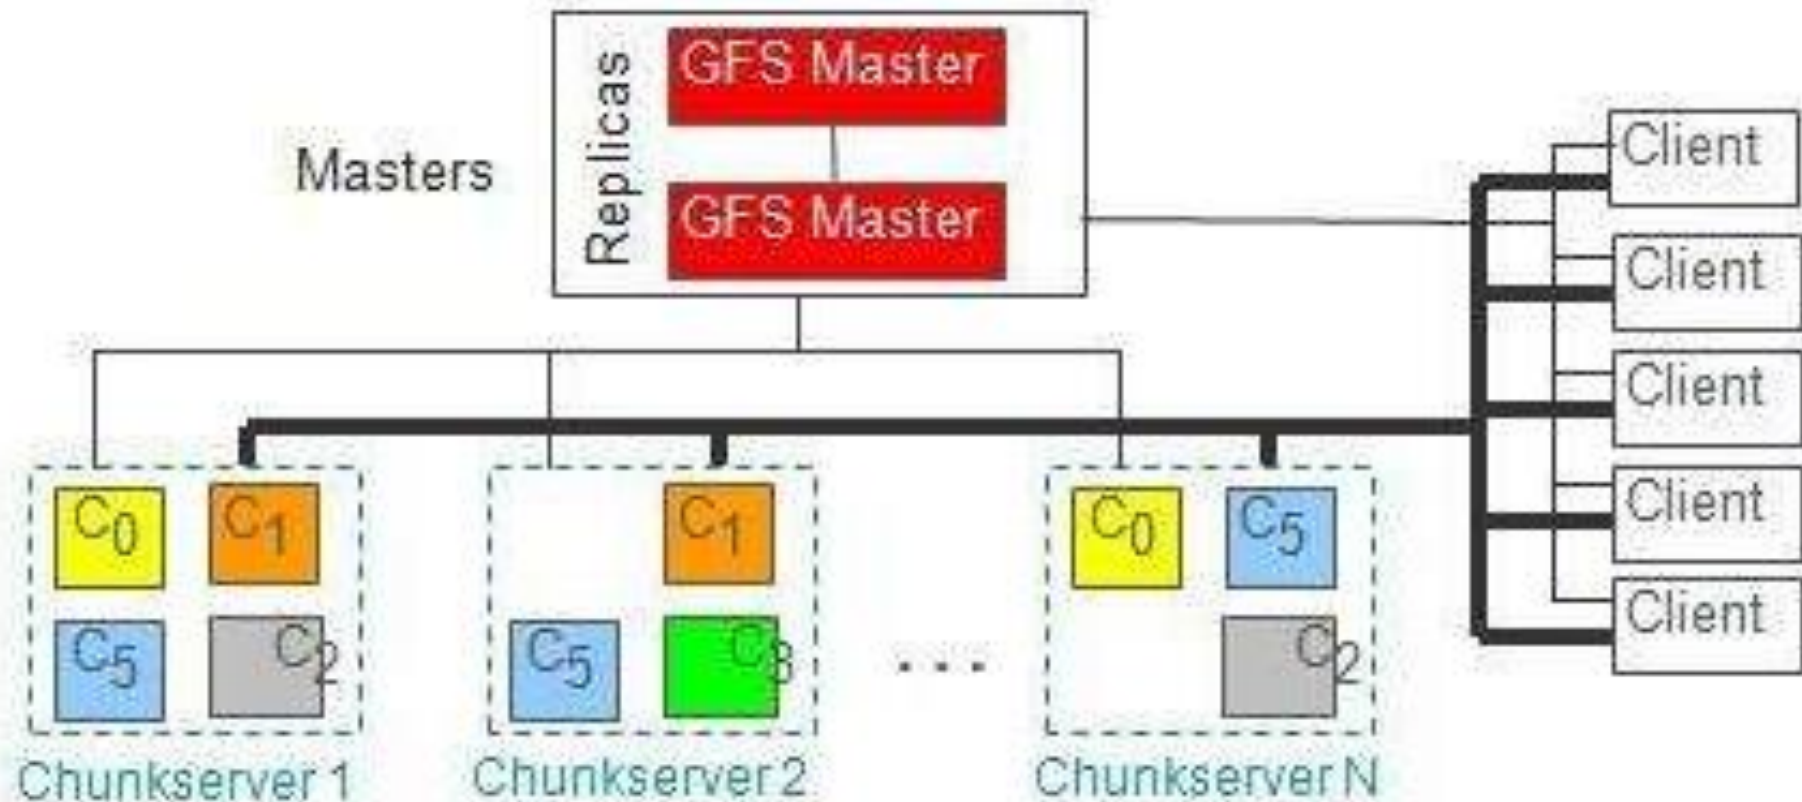
# Google File System (GFS)

- Turn “small machines” into “big system”
- Turn “bad machines” into “good system”
- Prepare for parallel computation





# Google File System (GFS)



- Files broken into chunks (typically 64 MB)
- Master manages metadata
- Data transfers happen directly between clients/chunkservers

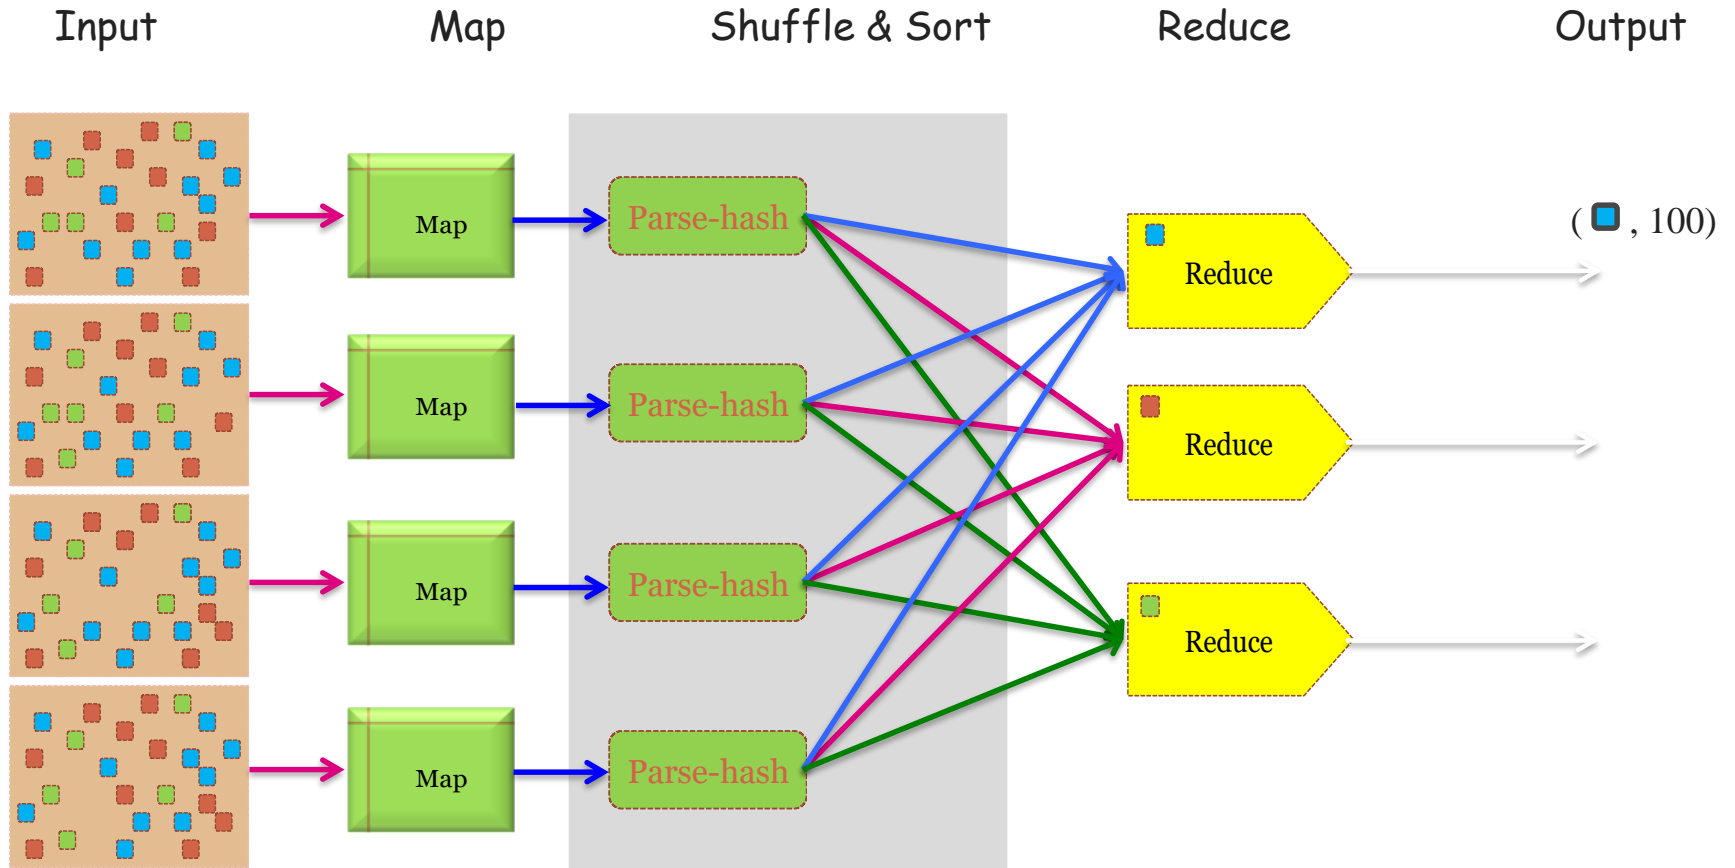
# Google MapReduce

---

- ❖ A reliable, fault-tolerant, parallel computing software framework for large-scale data sets for low-cost hardware clusters
  - ◆ Large volume of data (>1PB)
  - ◆ Large-scale parallel processing (>1 thousand nodes)
  - ◆ Make parallel computation easy

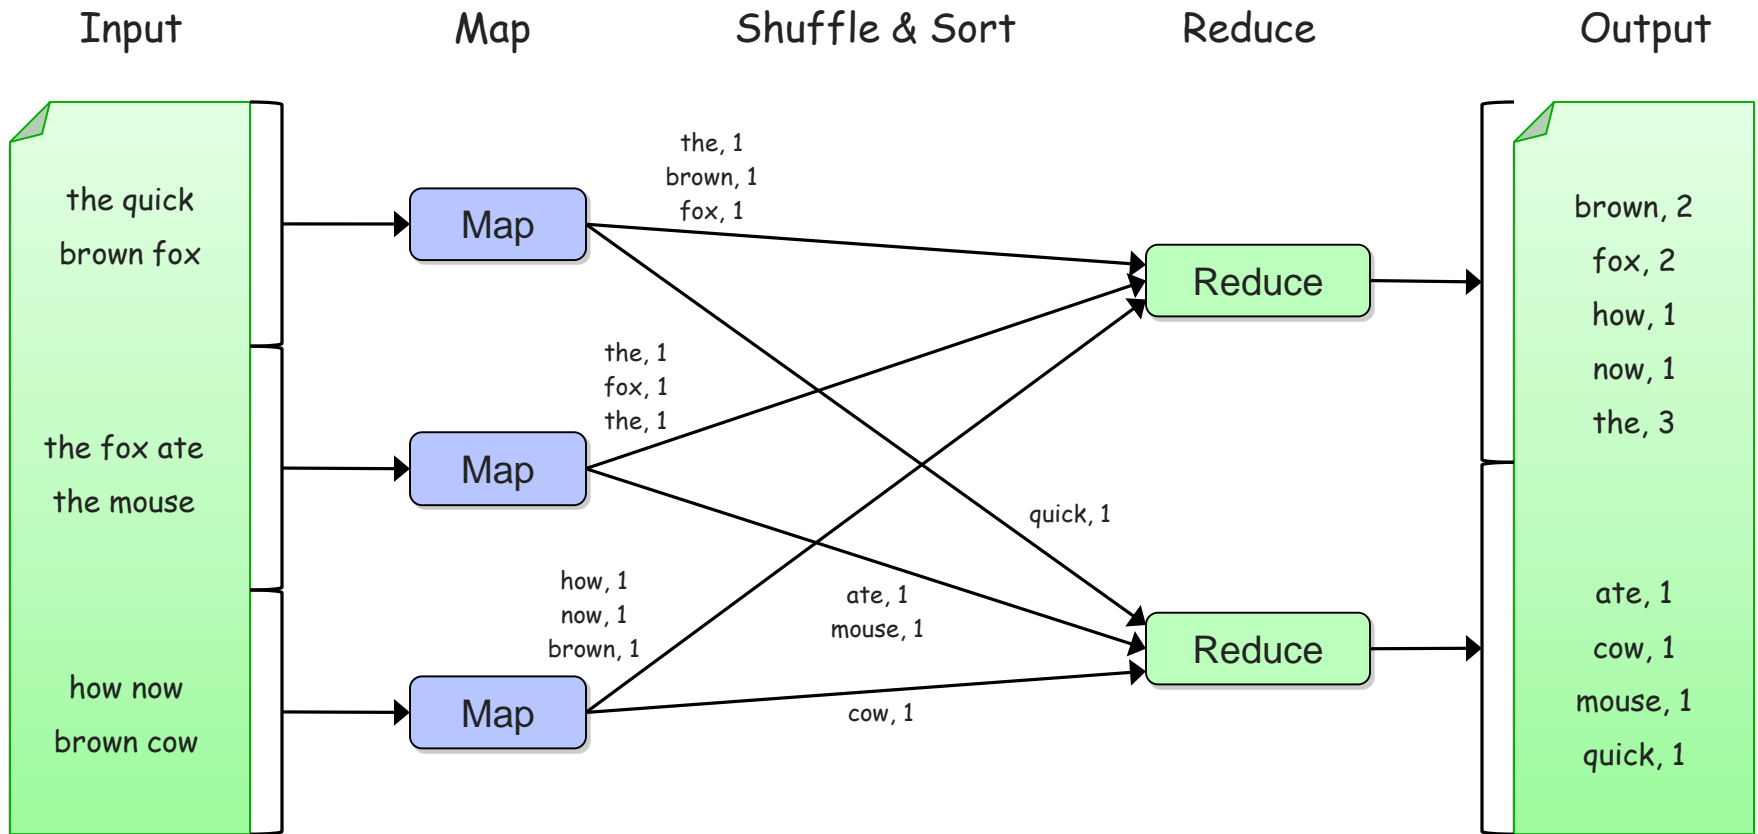
**MapReduce = Map (divide) + Reduce (merge and sort)**

# Example 1: Color Count



*Users only provide the “Map” and “Reduce” functions*

# Example 2: Word Count



# Design Principles of Hadoop

---

- ❖ Need to process big data
- ❖ Need to parallelize computation across thousands of nodes
- ❖ **Commodity hardware**
  - ◆ Large number of low-end cheap machines working in parallel to solve a computing problem
- ❖ This is in contrast to traditional parallel DBs
  - ◆ Small number of high-end expensive machines

# Design Principles of Hadoop (cont.)

---

## ❖ **Automatic parallelization & distribution**

- ◆ Hidden from the end-user

## ❖ **Fault tolerance and automatic recovery**

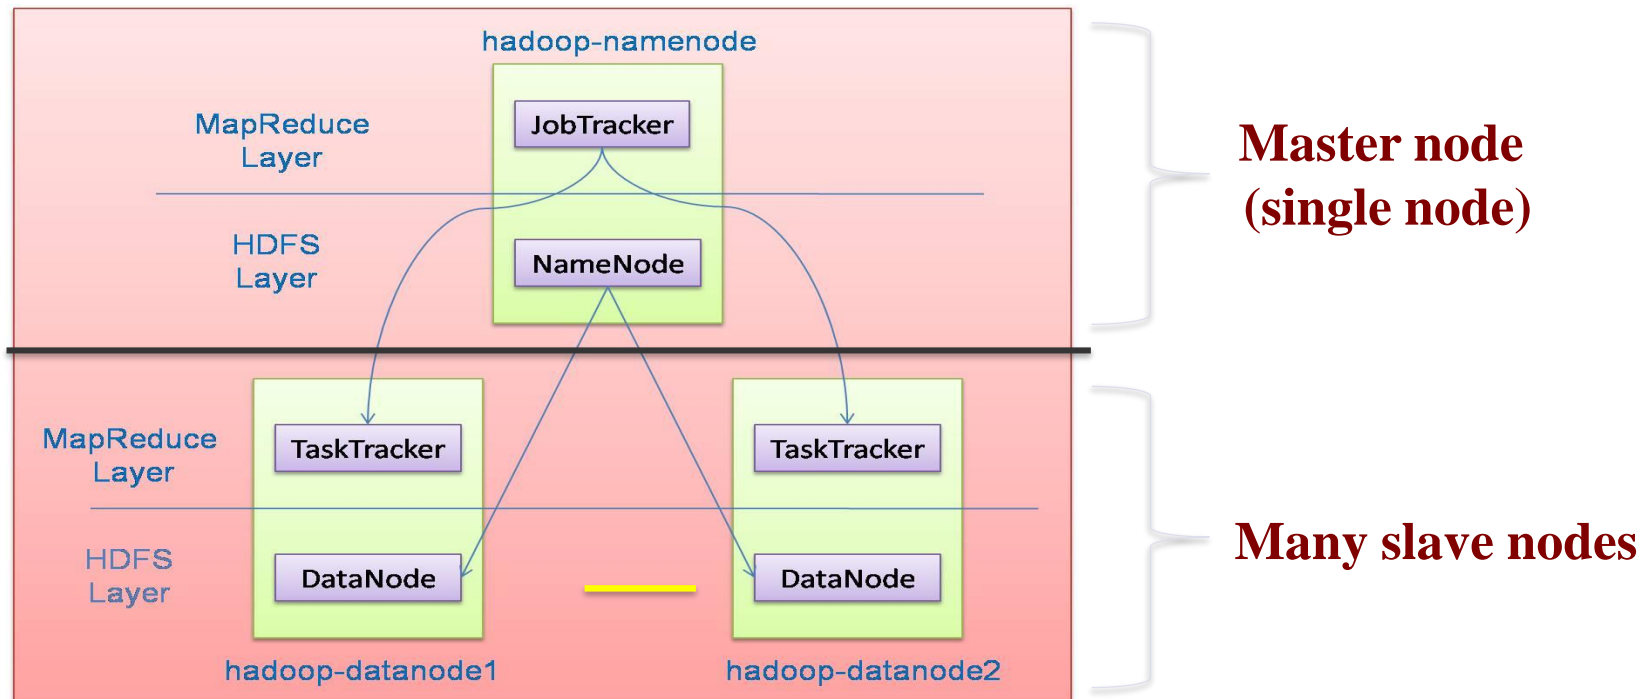
- ◆ Nodes/tasks will fail and will recover automatically

## ❖ **Clean and simple programming abstraction**

- ◆ Users only provide two functions “map” and “reduce”

# Hadoop Architecture

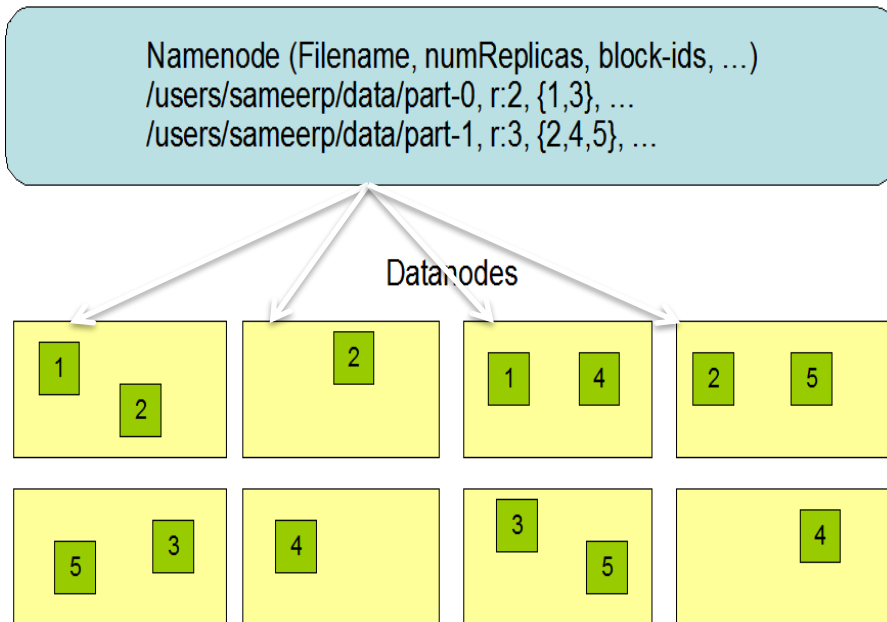
- ❖ Master-slave & shared-nothing
- ❖ Two layers
  - ◆ Distributed file system (HDFS)
  - ◆ Execution engine (MapReduce)



# Hadoop Distributed File System (HDFS)

## Master node

Block Replication



## Slave nodes

## Centralized namenode

- Maintains metadata info about files

File *F*

1	2	3	4	5
---	---	---	---	---

Blocks (64 MB)

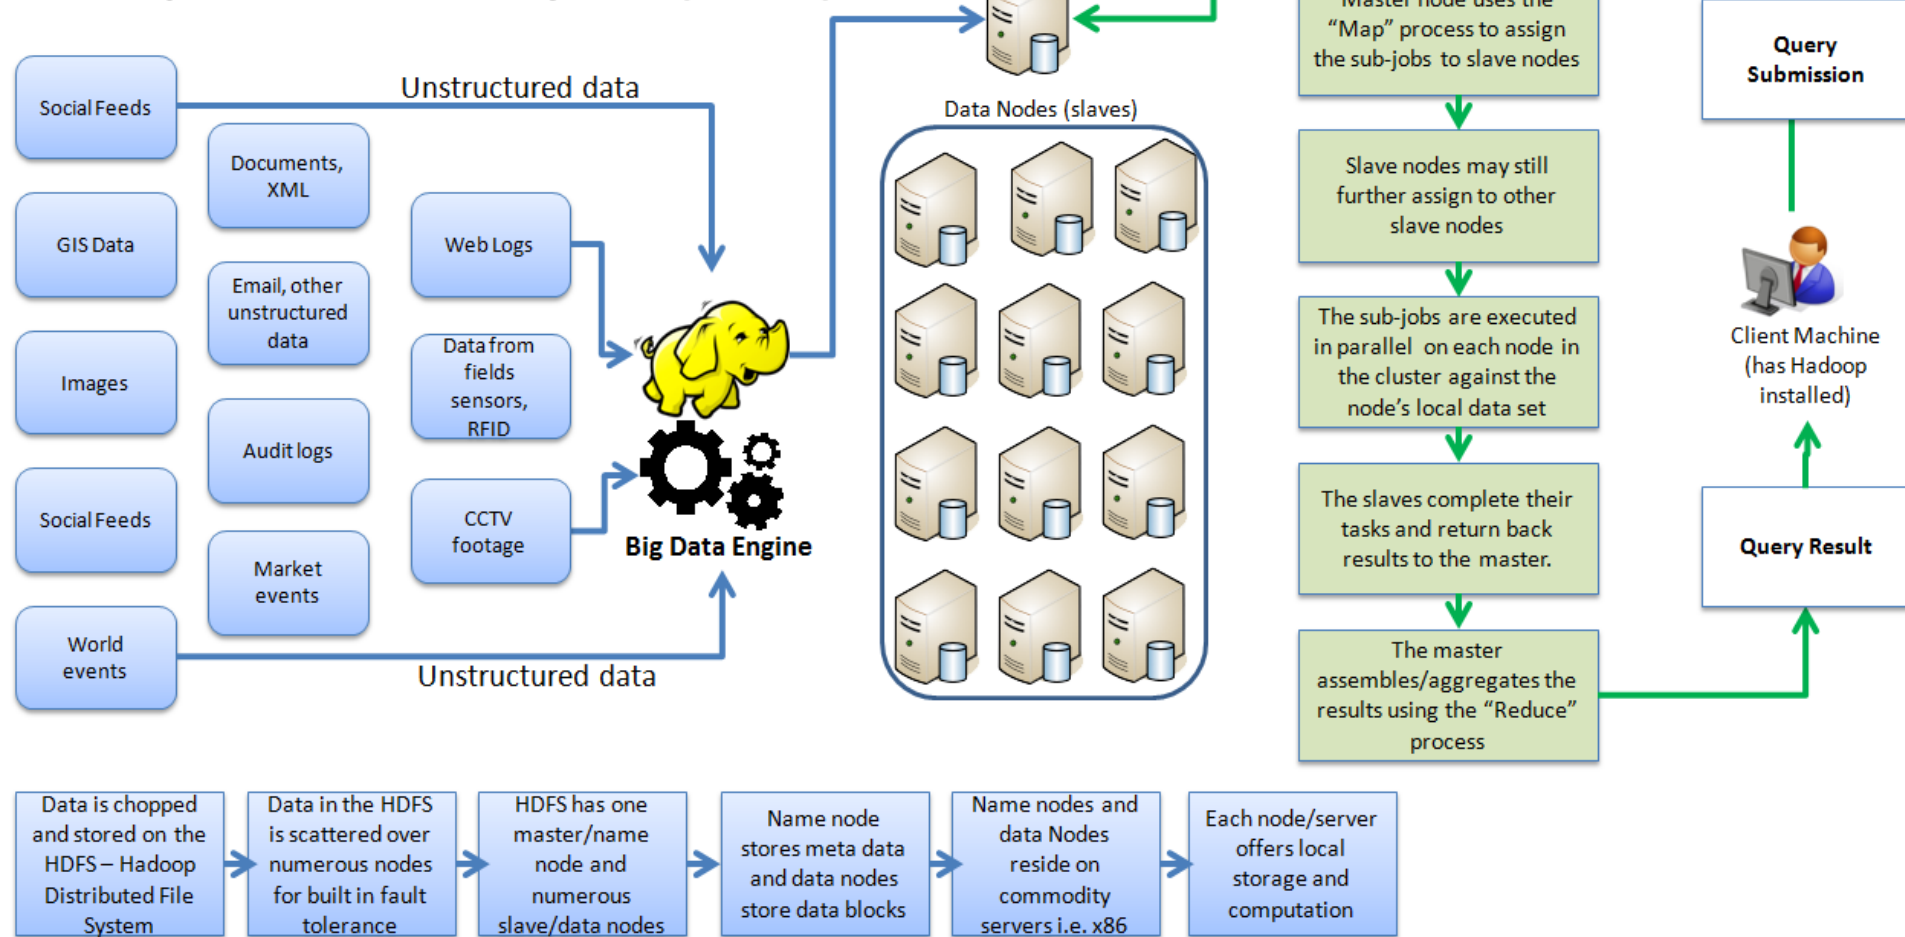
## Many datanode (1000s)

- Store the actual data
- Files are divided into blocks
- Each block is replicated *N* times (Default = 3)



# Master-Slave Architecture

## Storing & Querying Big Data in Hadoop Distributed File System ( HDFS )



Designed by Sri Prakash, November 2012


# Main Properties of HDFS

---

- ❖ **Large:** A HDFS instance may consist of thousands of server machines, each storing part of the file system's data
- ❖ **Replication:** Each data block is replicated many times (default is 3)
- ❖ **Failure:** Failure is the norm rather than exception
- ❖ **Fault Tolerance:** Detection of faults and quick, automatic recovery from them is a core architectural goal of HDFS
  - ◆ Namenode is consistently checking Datanodes

# Properties of MapReduce Engine


- ❖ Job Tracker is the master node (runs with the namenode)
  - ◆ Receives the user's job
  - ◆ Decides on how many tasks will run (number of mappers)
  - ◆ Decides on where to run each mapper (concept of locality)



- This file has 5 blocks → run 5 map tasks
- Where to run the task reading block “1”
  - *Try to run it on Node 1 or Node 3*

# Properties of MapReduce Engine (cont.)

- ❖ Task Tracker is the slave node (runs on each datanode)
  - ◆ Receives the task from Job Tracker
  - ◆ Runs the task until completion (either map or reduce task)
  - ◆ Always in communication with the Job Tracker reporting progress



*In this example, 1 map-reduce job consists of 4 map tasks and 3 reduce tasks*

(Example: Color Count)

# Key-Value Pairs

---

- ❖ Mappers and Reducers are users' code (provided functions)
- ❖ Just need to obey the Key-Value pairs interface
- ❖ **Mappers:**
  - ◆ Consume <key, value> pairs
  - ◆ Produce <key, value> pairs
- ❖ **Reducers:**
  - ◆ Consume <key, <list of values>>
  - ◆ Produce <key, value>
- ❖ **Shuffling and Sorting:**
  - ◆ Hidden phase between mappers and reducers
  - ◆ Groups all similar keys from all mappers, sorts and passes them to a certain reducer in the form of <key, <list of values>>

# Use of MapReduce/Hadoop

---

- ❖ Google: Inventors of MapReduce computing paradigm
- ❖ Yahoo: Developing Hadoop open-source of MapReduce
- ❖ IBM, Microsoft, Oracle
- ❖ Facebook, Amazon, AOL, NetFlex
- ❖ Many others + universities and research labs
- ❖ Applications
  - ❖ Web applications, social networks, scientific applications, applications generating big data

# Large-Scale Data Analytics

## ❖ MapReduce computing paradigm (e.g., Hadoop) vs. Traditional database systems



Scalability (petabytes of data,  
thousands of machines)



Flexibility in accepting all data  
formats (no schema)



Efficient and simple fault-tolerant  
mechanism



Commodity inexpensive hardware



Performance (tons of indexing,  
tuning, data organization tech.)



Features:

- Provenance tracking
- Annotation management
- ....


# Hadoop vs. DDBS

	Distributed Databases	Hadoop
Computing Model	<ul style="list-style-type: none"><li>- Notion of transactions</li><li>- Transaction is the unit of work</li><li>- ACID properties, Concurrency control</li></ul>	<ul style="list-style-type: none"><li>- Notion of jobs</li><li>- Job is the unit of work</li><li>- No concurrency control</li></ul>
Data Model	<ul style="list-style-type: none"><li>- Structured data with known schema</li><li>- Read/Write mode</li></ul>	<ul style="list-style-type: none"><li>- Any data will fit in any format</li><li>- (un)(semi)structured</li><li>- ReadOnly mode</li></ul>
Cost Model	<ul style="list-style-type: none"><li>- Expensive servers</li></ul>	<ul style="list-style-type: none"><li>- Cheap commodity machines</li></ul>
Fault Tolerance	<ul style="list-style-type: none"><li>- Failures are rare</li><li>- Recovery mechanisms</li></ul>	<ul style="list-style-type: none"><li>- Failures are common over thousands of machines</li><li>- Simple yet efficient fault tolerance</li></ul>
Key Characteristics	<ul style="list-style-type: none"><li>- Efficiency, optimizations, fine-tuning</li></ul>	<ul style="list-style-type: none"><li>- Scalability, flexibility, fault tolerance</li></ul>




# Another Master-Slave Fashion

## TAO - Facebook's Distributed Data Store for the Social Graph



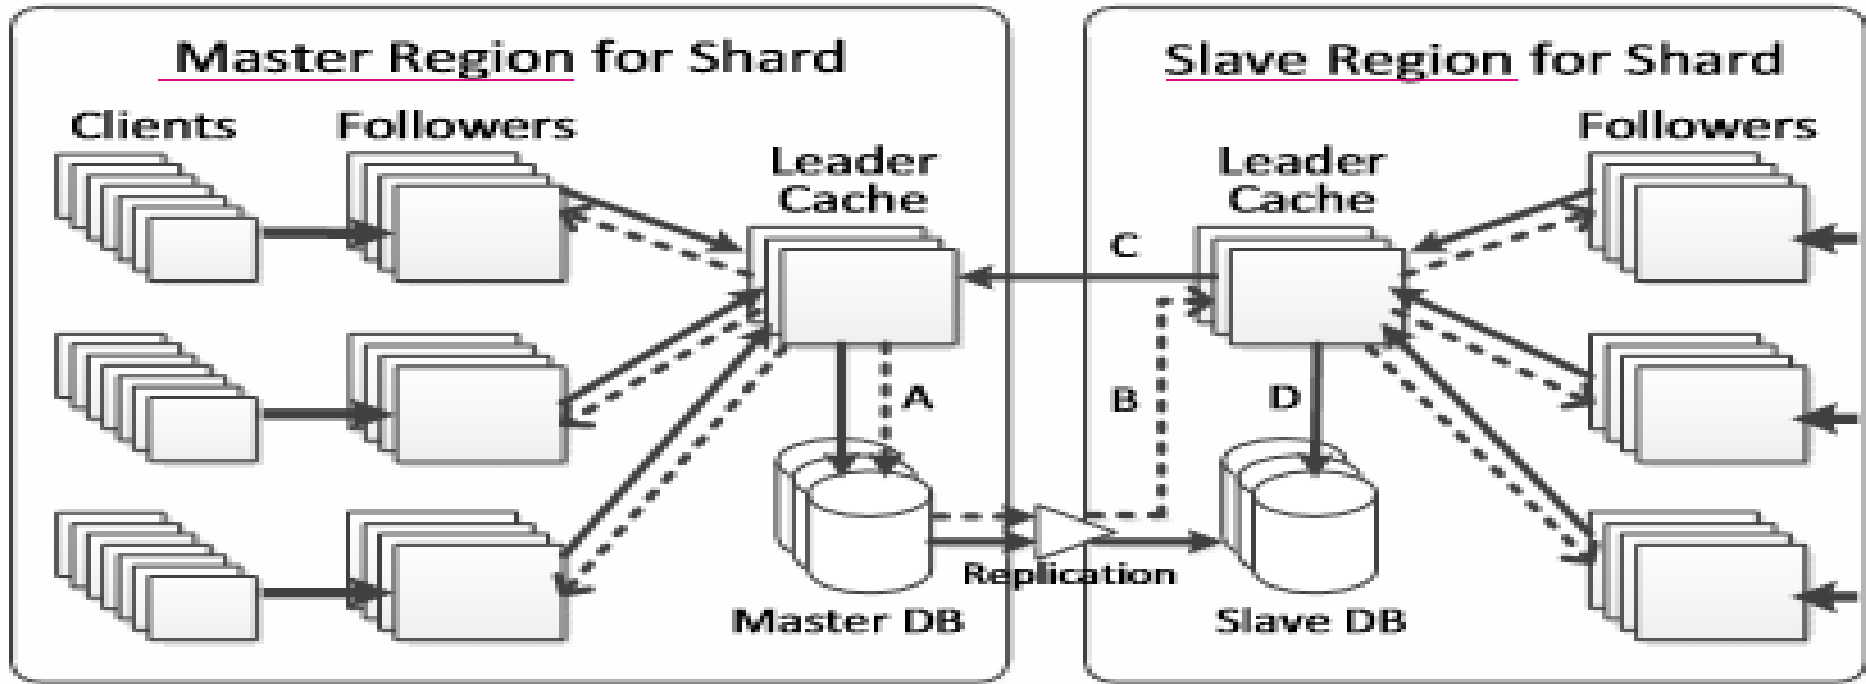
# TAO's Goals/Challenges

Billions of users, billions of data accesses to the social graph  
How to handle this huge workload efficiently?



- Efficiency at scale
- Low read latency
- Timeliness of writes
- High Read Availability

# TAO's Master-Slave Architecture



## Why “Regions” ?

Reason: network latency.


- Each region has **one Leader tier** and **multiple Followers**
- **Master Region** vs. **Slave Region**
- Master/Slave designation is made separately for each shard

## Why Master/Slave ?

Reason: read misses are 25x more frequent than writes. In this way, **Slave Regions** can entirely self-service read misses without bothering **Master Region** (via Slave leader).

- **Writes in a Slave Region go to Master Region**
- Consistency is handled asynchronously by **sending cache maintenance messages from the Master to Slaves**


# TAO's Follower and Leader Caches



1. Client request goes to a single cache server in a local “follower” tier
2. Follower fulfills request. If read miss or a write, go to a Leader
3. If Leader read miss or a write, Leader will go to the DB

# 3 - Peer-to-Peer Distributed Architecture

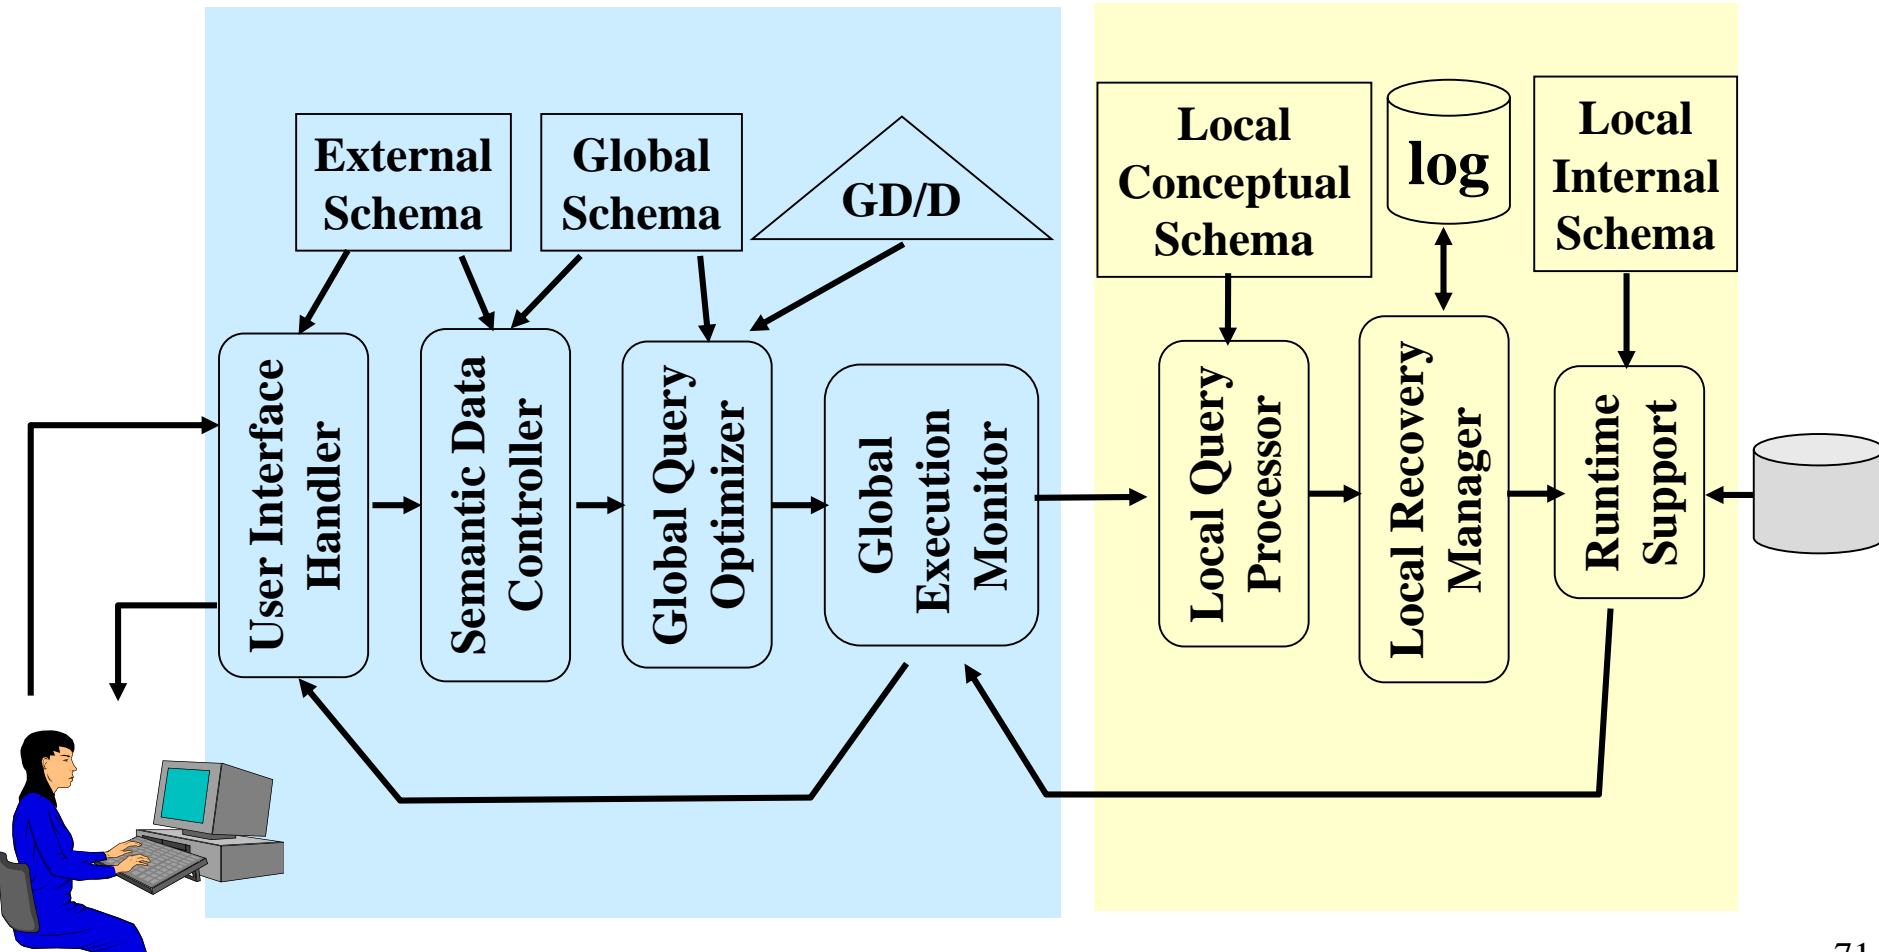
---



# A Peer DBMS

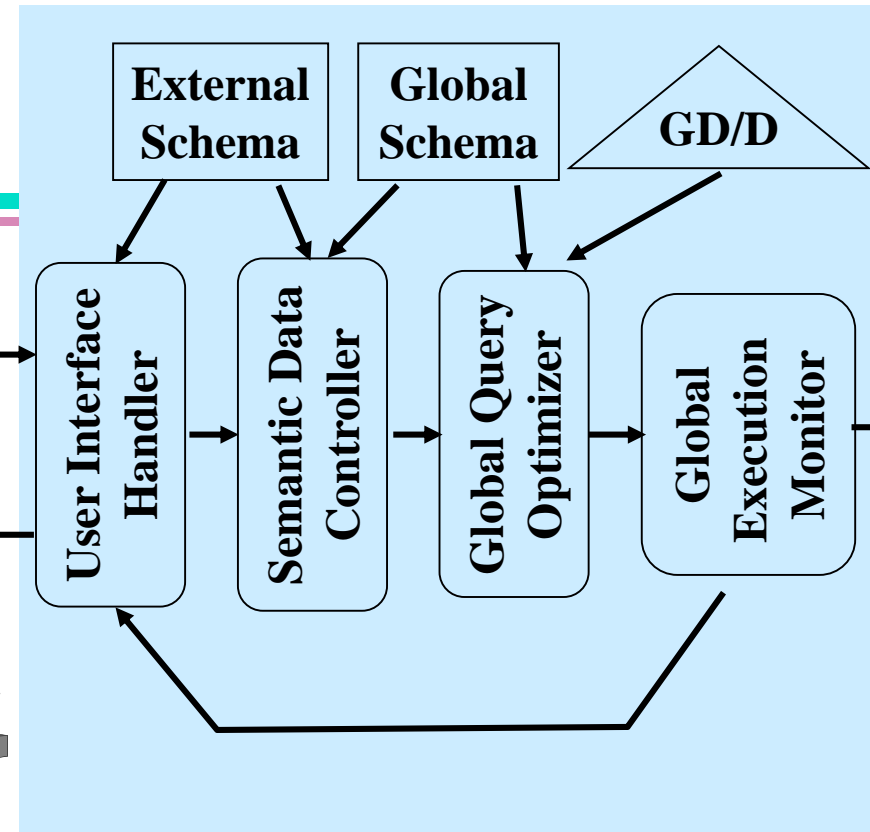
## USER PROCESSOR

## DATA PROCESSOR



# User Processor

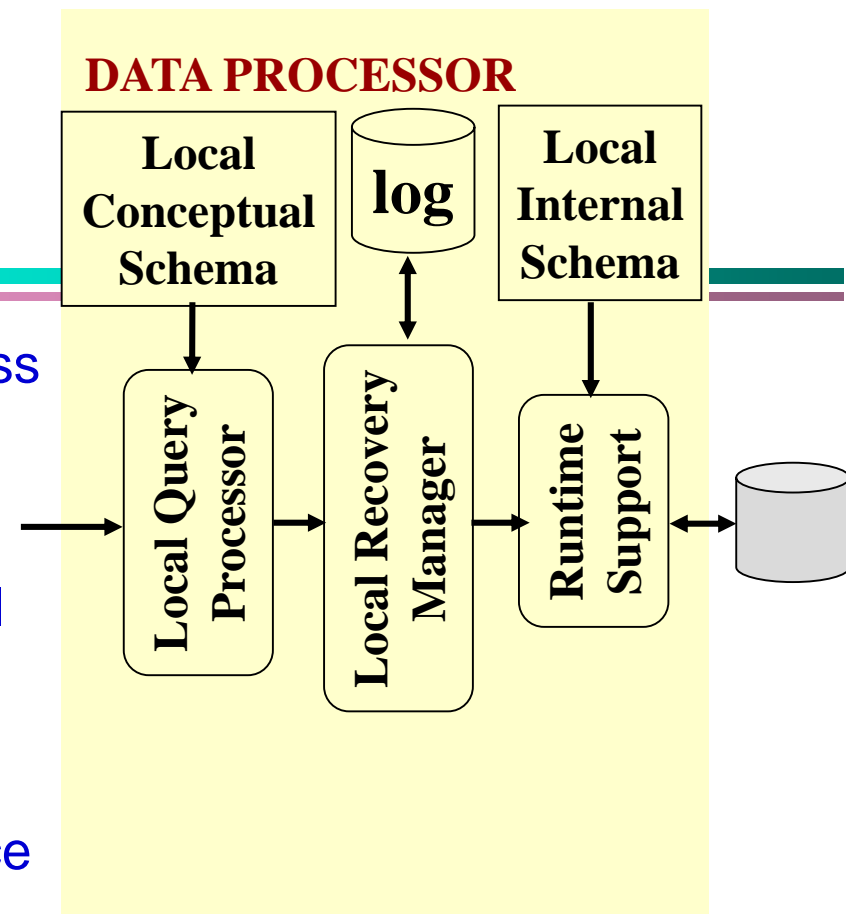
- ❖ **User interface handler** interprets user commands and formats the result data as it is sent to the user.
- ❖ **Semantic data controller** checks the integrity constraints and authorization requirements.
- ❖ **Global query optimizer and decomposer** determines execution strategy, translates global queries to local queries, and generates strategy for distributed join operations.
- ❖ **Global execution monitor** (distributed transaction manager) coordinates the distributed execution of the user request.





# Data Processor

- ❖ **Local query processor** selects the access path and is involved in local query optimization and join operations.
- ❖ **Local recovery manager** maintains local database consistency.
- ❖ **Run-time support processor** physically accesses the database. It is the interface to the OS and contains database buffer manager.






# Apache Cassandra: Peer-to-Peer

---

- ❖ A **peer-to-peer distributed storage system** for managing very large amounts of data spread out across many commodity servers, while providing highly available service with no single point of failure.
- ❖ Originally developed at Facebook, open sourced in 2008, and became a top-level Apache project in 2010.
- ❖ Deployed as the backend storage system for multiple services in Twitter, Reddit, Netflix, Digg, Rackspace, Cisco and more companies that have large, active data sets.



# Application Areas of Cassandra



## Internet of things applications

Perfect for consuming lots of fast incoming data from devices, sensors and similar mechanisms that exist in many different locations.



## Messaging Apps

Serves as the database backbone for numerous mobile phone and messaging providers' applications.




## Social media analytics and recommendation engines

Many online companies, websites, and social media providers use Cassandra to ingest, analyze, and provide analysis and recommendations to their customers.

# Architecture of Cassandra

Masterless “ring” architecture. All nodes play an identical role.



# Gossip

---

- ❖ Peer to peer protocol
- ❖ Gossip every one second, up to 3 others nodes
- ❖ Share state information of the nodes
  
- ❖ Achieve three major functions
  - ◆ Failure detection
  - ◆ Dynamic load balance
  - ◆ Elastic expansion without central control

# Dimension 2: Heterogeneity (异质)

---

- ❖ Various levels (hardware, communication, OS)
- ❖ DBMS important ones (like data model, query language, transaction management algorithms, etc.)
  - ◆ 0 - homogeneous
  - ◆ 1 - heterogeneous

# Dimension 3: Autonomy (自治)

---

- ❖ Refer to the distribution of control, not of data, indicating the degree to which individual DBMSs can operate independently.
- ❖ Requirements of an autonomous system
  - ◆ The **local operations of the individual DBMSs** are not affected by their participation in the DDBS.
  - ◆ The **individual DBMS query processing and optimization** should not be affected by the execution of global queries that access multiple databases.
  - ◆ **System consistency or operation** should not be compromised when individual DBMSs join or leave the distributed database confederation.

# Three Versions of Autonomy

---

## ❖ Design autonomy

- ◆ Ability of a component DBMS to decide on issues related to its own design
- ◆ Freedom for individual DBMSs to use data models and transaction management techniques they prefer

## ❖ Communication autonomy

- ◆ Ability of a component DBMS to decide whether and how to communicate with other DBMSs
- ◆ Freedom for individual DBMSs to decide what information (data & control) is to be exported

## ❖ Execution autonomy

- ◆ Ability of a component DBMS to execute local operations in any manner it wants to.
- ◆ Freedom for individual DBMSs to execute transactions submitted in any way that it wants to

# Dimension 3: Autonomy (*cont.*)

---

- ♦ 0 – Tightly coupled - integrated
- ♦ 1 – Semi-autonomous - federated
- ♦ 2 – Total Isolation – multi-database systems



# Taxonomy of Distributed Databases

---

## ❖ Composite DBMSs - tight integration

- ◆ single image of entire database is available to any user
- ◆ can be single or multiple sites
- ◆ can be homogeneous or heterogeneous

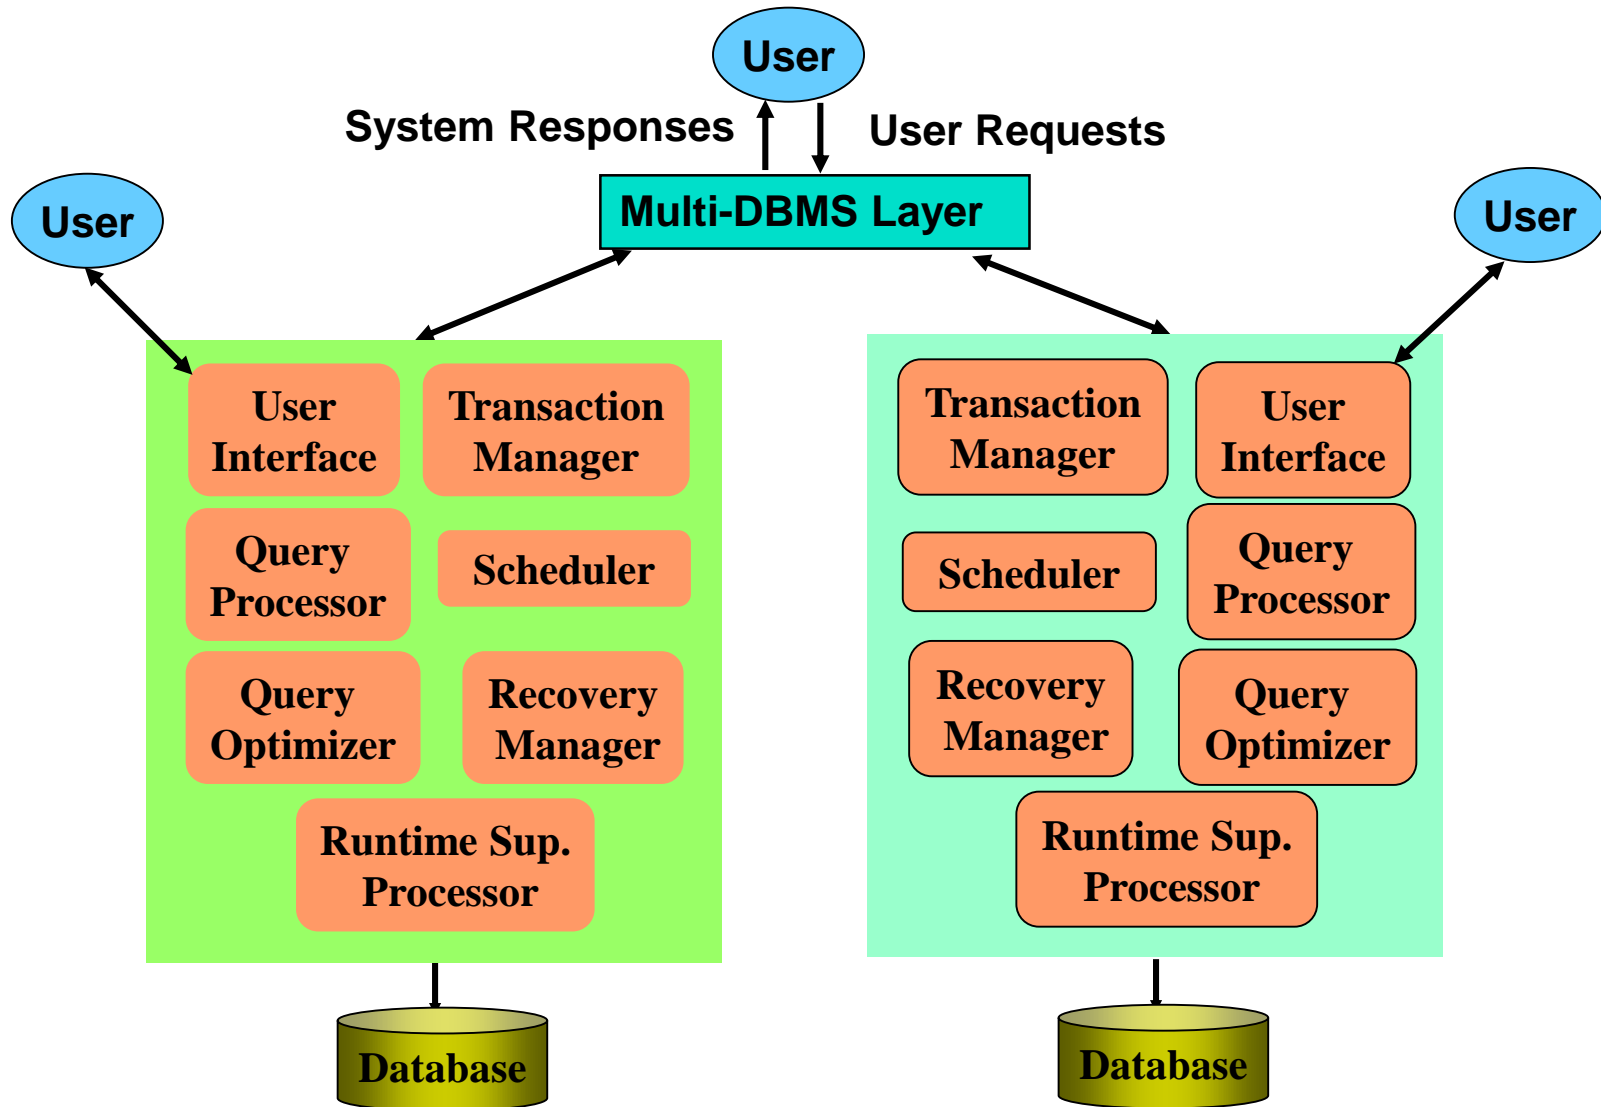
## ❖ Federated DBMSs – semi-autonomous

- ◆ DBMSs that can operate independently, but have decided to make some parts of their local data shareable
- ◆ can be single or multiple sites
- ◆ they need to be modified to enable them to exchange information

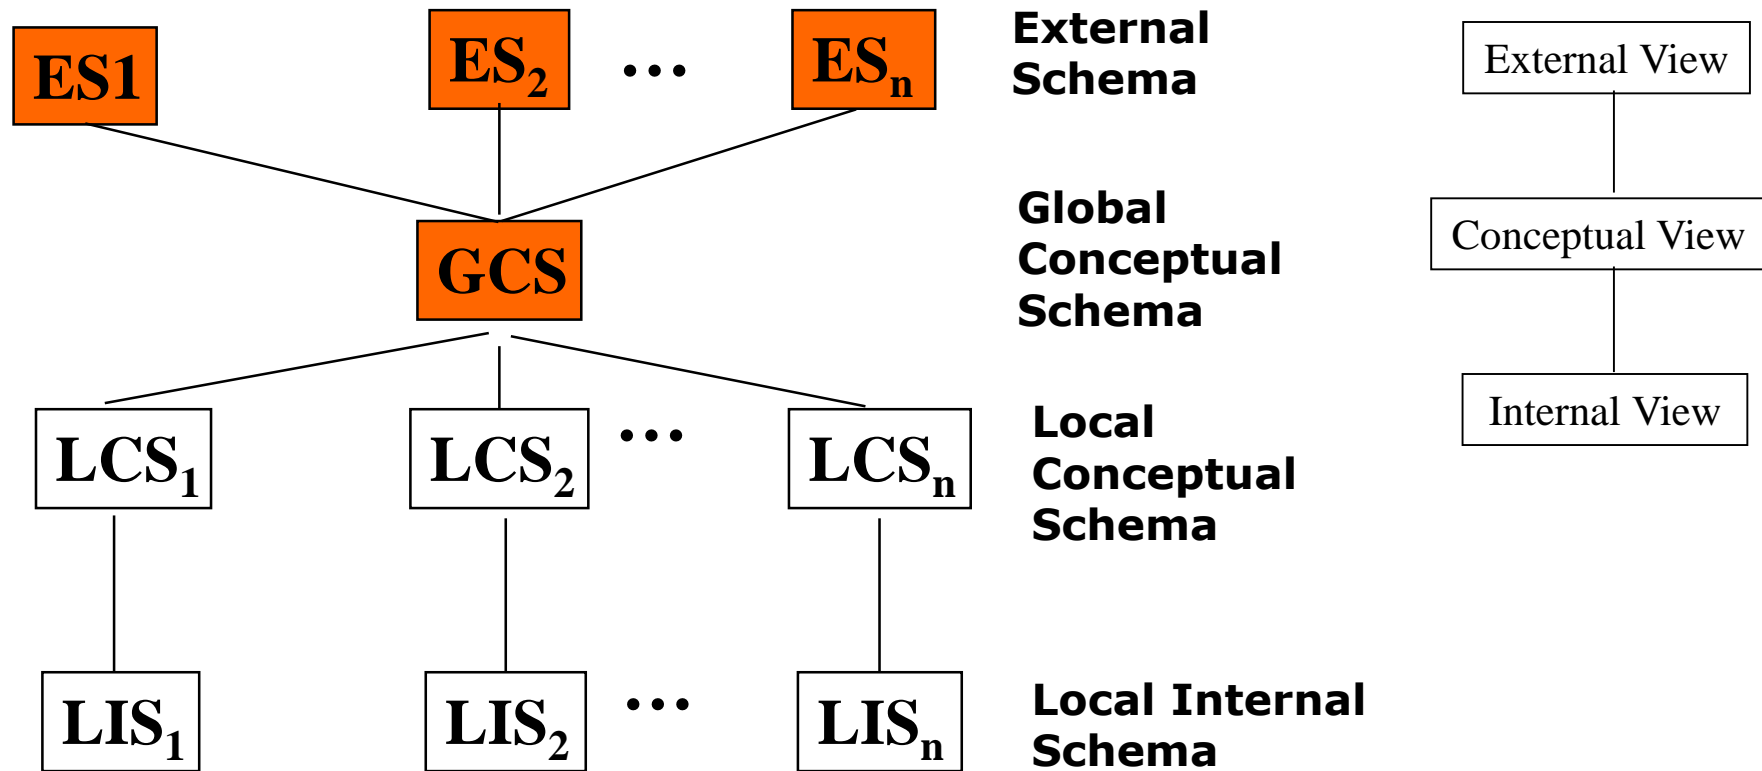
## ❖ Multi-database Systems - total isolation

- ◆ individual systems are stand alone DBMSs, which know neither the existence of other databases or how to communicate with them
- ◆ no global control over the execution of individual DBMSs.
- ◆ can be single or multiple sites
- ◆ homogeneous or heterogeneous

# Components of a Multi-DBMS

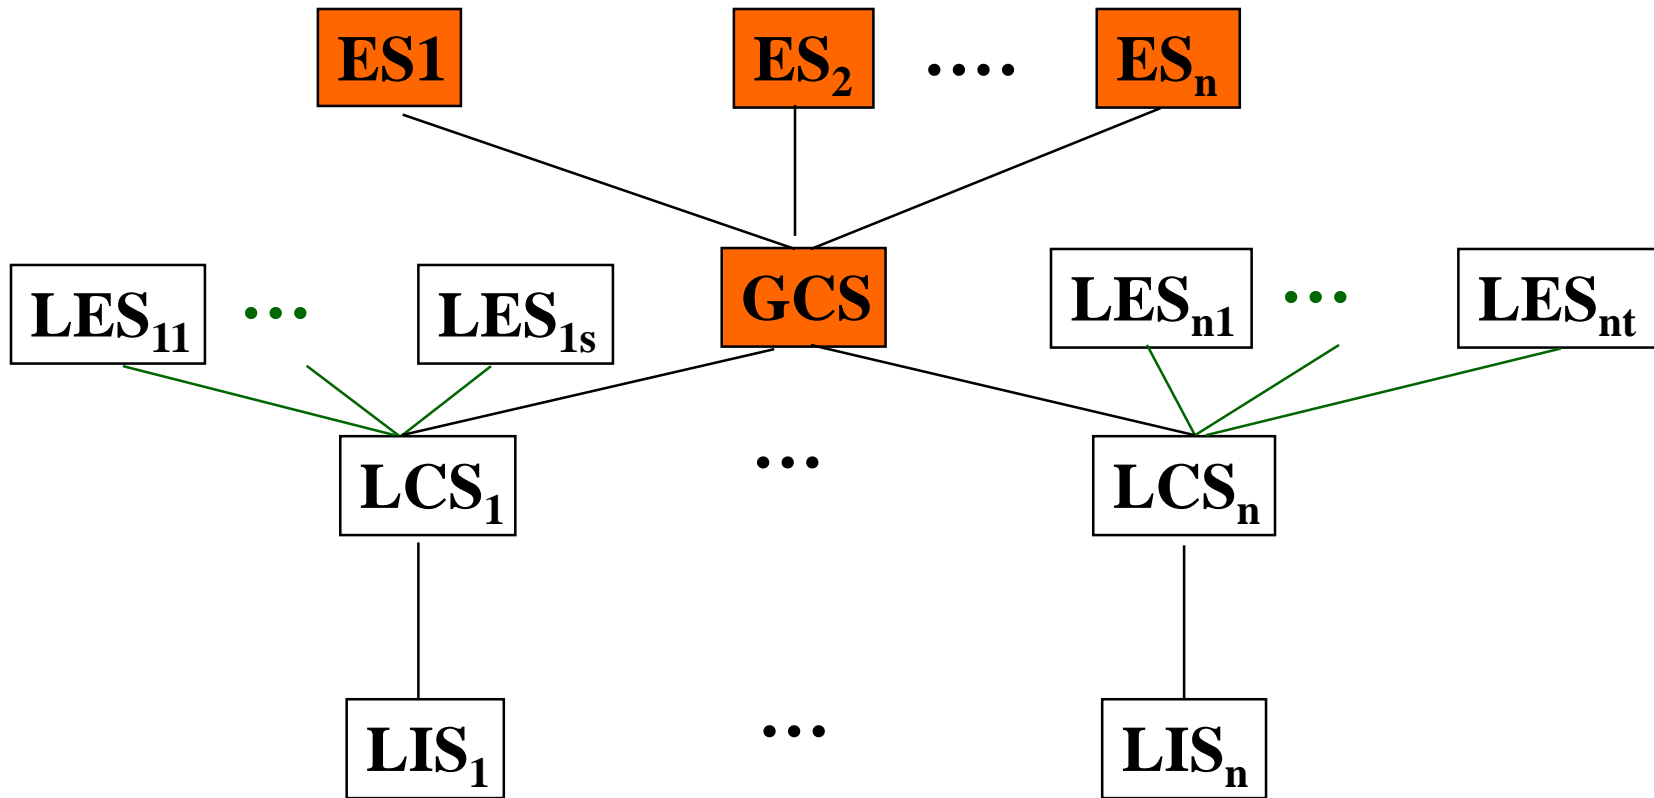


# Distributed Database Reference Architecture




**It is logically integrated. Provide for the levels of transparency**

# Multi-DBMS Architecture with a Global Conceptual Schema



- The GCS is generated by integrating LES's or LCS's
- The users of a local DBMS can maintain their autonomy
- Design of GCS is bottom-up

# Multi-DBMS without a Global Conceptual Schema



- ❖ Local database system layer consists of several DBMSs which present to multidatabase layer part of their databases
- ❖ The shared database has either local conceptual schema or external schema
- ❖ External views on one or more LCSs.
- ❖ Access to multiple databases through application programs

# Multi-DBMS without a Global Conceptual Schema (*cont.*)

---

## ❖ Multi-DBMS components architecture

- ◆ Existence of fully fledged local DBMSs
- ◆ Multi-DBMS is a layer on top of individual DBMSs that support access to different databases
- ◆ The complexity of the layer depends on existence of GCS and heterogeneity

## ❖ Federated Database Systems

- ◆ Do not use global conceptual schema
- ◆ Each local DBMS defines export schema
- ◆ Global database is a union of export schemas
- ◆ Each application accesses global database through import schema (external view)

# Outline

---

- ❖ Introduction
- ❖ Top-Down Design of DDBMS Architecture
  - ◆ Schema and Distribution Transparency
- ❖ Bottom-up Design of DDBMS Architecture
  - ◆ Architectural Alternatives for DDBMSs
  - ◆ Reference Architectures for a DDBMS
- ☞ Global Directory/Dictionary


# Global Directory/Dictionary

---

- ❖ Directory is itself a database that contains **meta-data about the actual data** stored in the database. It includes the support for fragmentation transparency in the classical DDBMS architecture.
- ❖ Directory can be local or distributed.
- ❖ Directory can be replicated and/or partitioned.
- ❖ Directory issues are very important for large multi-database applications, such as digital libraries.



# Alternative Directory Management Strategies



---

# **Question & Answer**