

10. Parallel Database Systems

Chapter 14

Parallel Database Systems

Fact and Database Problem

- ❖ A large volume of data uses disk and large main memory
 - ◆ $\text{Speed}(\text{disk}) \ll \text{speed}(\text{RAM}) \ll \text{speed}(\text{microprocessor})$

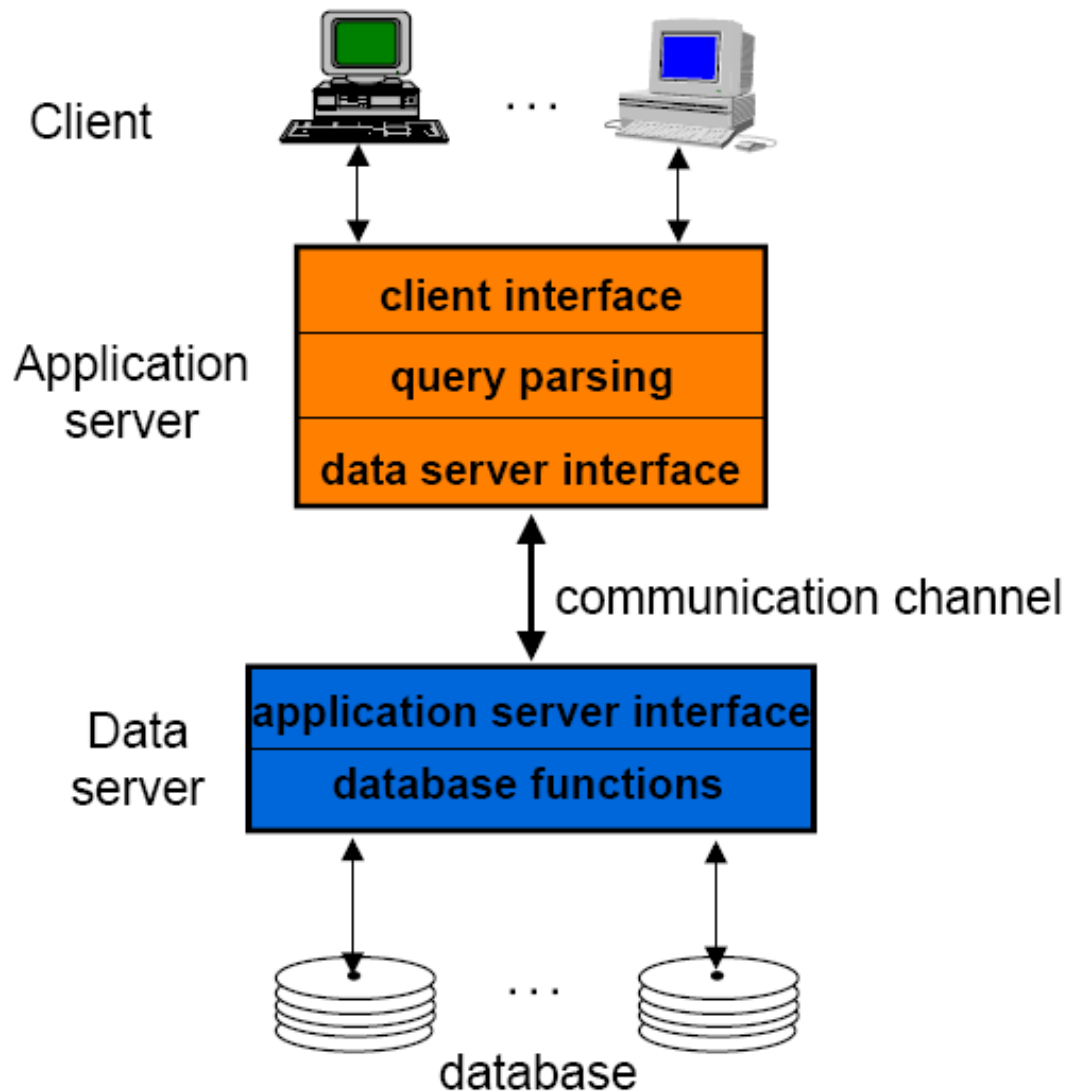
The Solution

- ❖ Increase the I/O bandwidth
 - ◆ Data partitioning
 - ◆ Parallel data access
- ❖ Origins (1980's): database machines
 - ◆ Hardware-oriented: bad cost-performance failure
 - ◆ Notable exception: ICL's CAFS Intelligent Search Processor
- ❖ 1990's: same solution but using standard hardware components integrated in a multiprocessor
 - ◆ Software-oriented
 - ◆ Exploiting continuing technology improvements

Multiprocessor Objectives

- ❖ High-performance with better cost-performance than mainframe or vector super computer
- ❖ Use many nodes, each with good cost-performance, communicating through network
 - ◆ Good cost via high-volume components
 - ◆ Good performance via bandwidth
- ❖ Trends
 - ◆ Microprocessor and memory (DRAM): off-the-shelf
 - ◆ Network (multiprocessor edge): custom
- ❖ The real challenge is to parallelize applications to run with good load balancing

Data Server Architecture



Objectives of Data Servers

- ❖ Avoid the shortcomings of the traditional DBMS approach
 - ◆ Centralization of data and application management
 - ◆ General-purpose OS (not DB-oriented)
- ❖ By separating the functions between
 - ◆ Application server (or host computer)
 - ◆ Data server (or database computer or back-end computer)

Data Server Approach: Assessment

❖ Advantages

- ♦ Integrated data control by the server (black box)
- ♦ Increased performance by dedicated system
- ♦ Can better exploit parallelism
- ♦ Fits well in distributed environments

❖ Potential problems

- ♦ Communication overhead between application and data server
 - High-level interface
- ♦ High cost with mainframe servers

Parallel Data Processing

- ❖ Three ways of exploiting high-performance multiprocessor systems:
 - ◆ Automatically detect parallelism in sequential programs (e.g., Fortran, OPS5)
 - ◆ Augment an existing language with parallel constructs (e.g., C*, Fortran90)
 - ◆ Offer a new language in which parallelism can be expressed or automatically inferred

Parallel Data Processing (*cont.*)

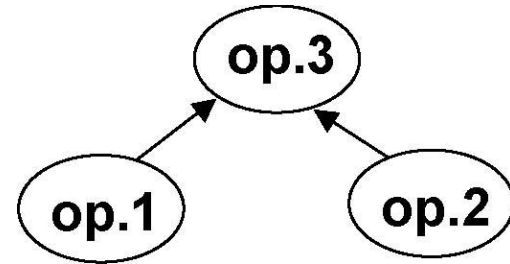
❖ Critique

- ◆ Hard to develop parallelizing compilers, limited resulting speed-up
- ◆ Enables the programmer to express parallel computations but too low-level

Data-based Parallelism


❖ Inter-operation

- ♦ operations of the same query in parallel



❖ Intra-operation

- ♦ the same operation in parallel on different data partitions



Parallel DBMS


- ❖ Loose definition: a DBMS implemented on a tightly coupled multiprocessor
- ❖ Naturally extends to distributed databases with one server per site

Parallel DBMS -Objectives

- ❖ Much better cost / performance than mainframe solution
- ❖ High-performance through parallelism
 - ◆ High throughput with inter-query parallelism
 - ◆ Low response time with intra-operation parallelism
- ❖ High availability and reliability by exploiting data replication
- ❖ Extensibility with the ideal goals
 - ◆ Linear speed-up
 - ◆ Linear scale-up

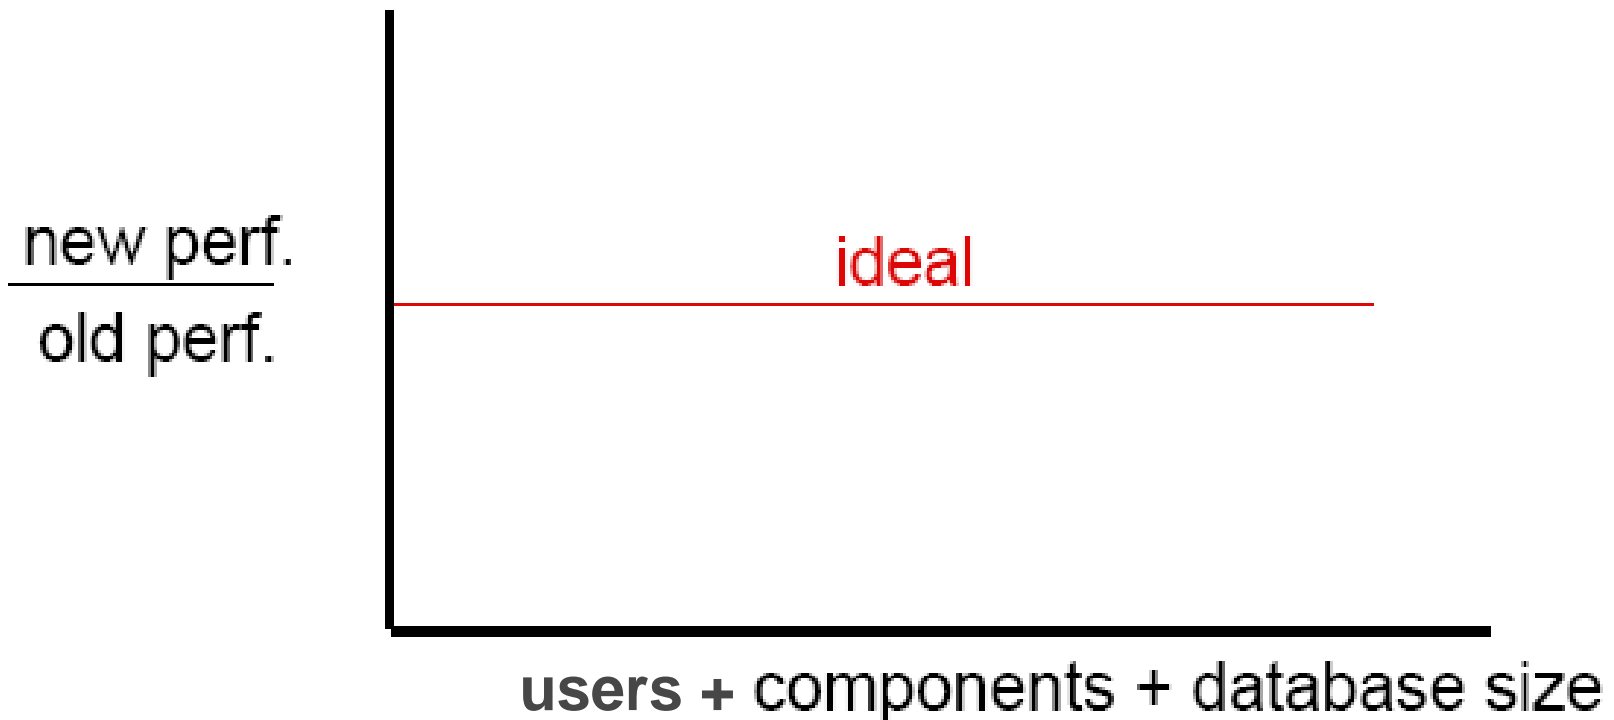
Linear Speed-up

- ❖ Linear increase in performance for a constant DB size and proportional increase of the system components (processor, memory, disk)



Linear Scale-up

- ❖ Sustained performance for a linear increase of database size and proportional increase of the system components.



Barriers to Parallelism

❖ Startup

- ◆ The time needed to start a parallel operation may dominate the actual computation time

❖ Interference

- ◆ When accessing shared resources, each new process slows down the others (hot spot problem)

❖ Skew

- ◆ The response time of a set of parallel processes is the time of the slowest one

❖ Parallel data management techniques intend to overcome these barriers

Parallel DBMS Functions

❖ Session manager

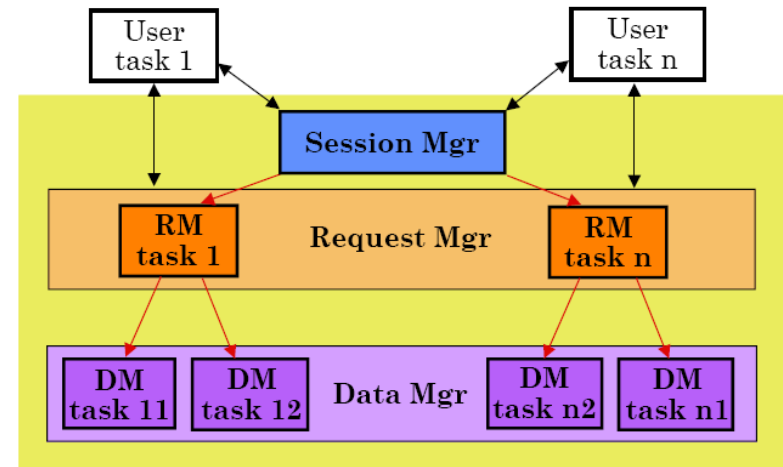
- ◆ Host interface
- ◆ Transaction monitoring

❖ Request manager

- ◆ Compilation and optimization
- ◆ Data directory management
- ◆ Semantic data control
- ◆ Execution control

❖ Data manager

- ◆ Execution of DB operations
- ◆ Transaction management support
- ◆ Data management



Parallel System Architectures

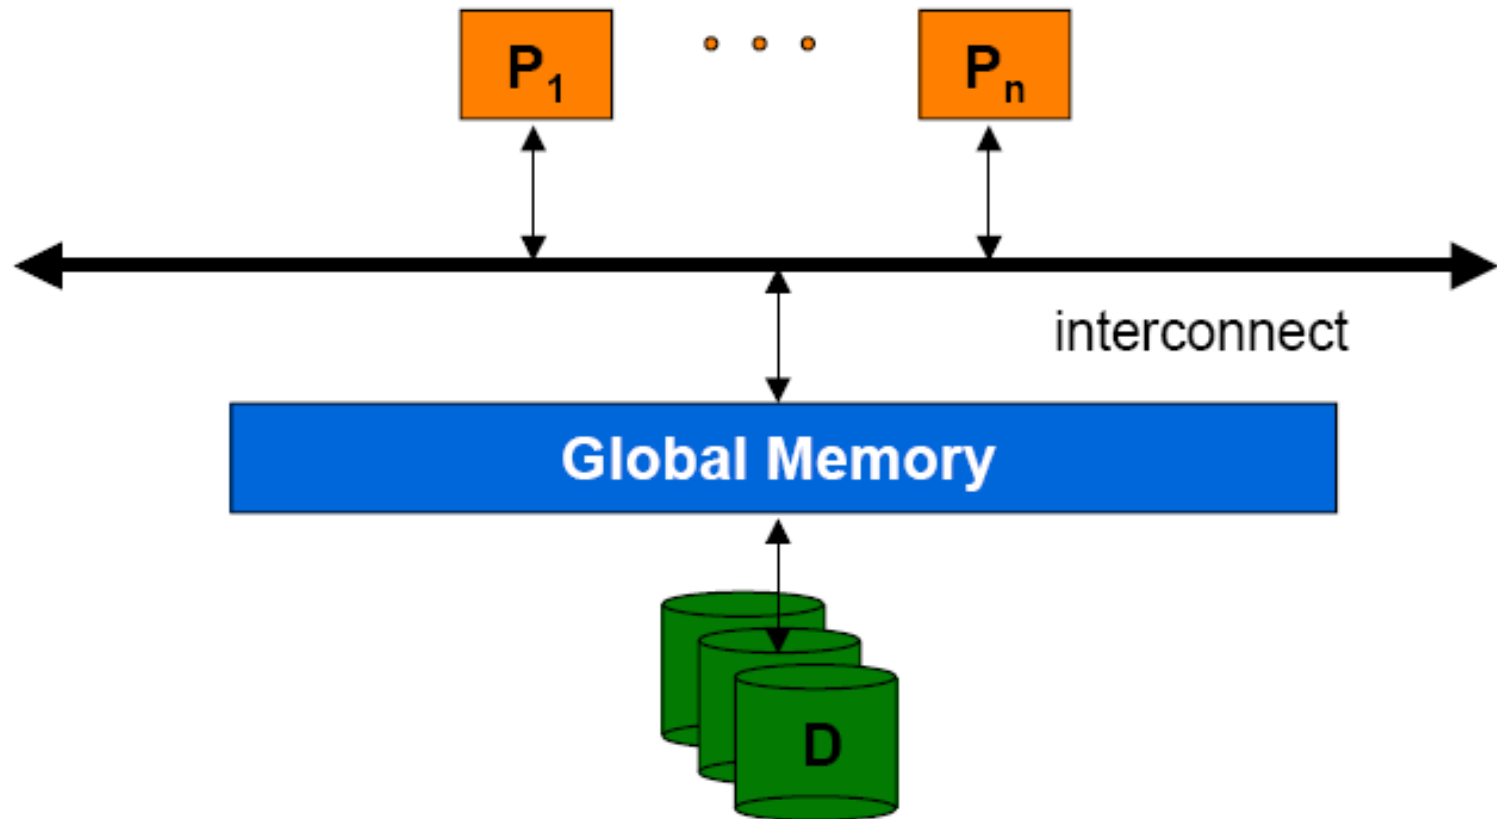
❖ Multiprocessor architecture alternatives

- ◆ Shared memory (shared everything)
- ◆ Shared disk
- ◆ Shared nothing (message-passing)

❖ Hybrid architectures

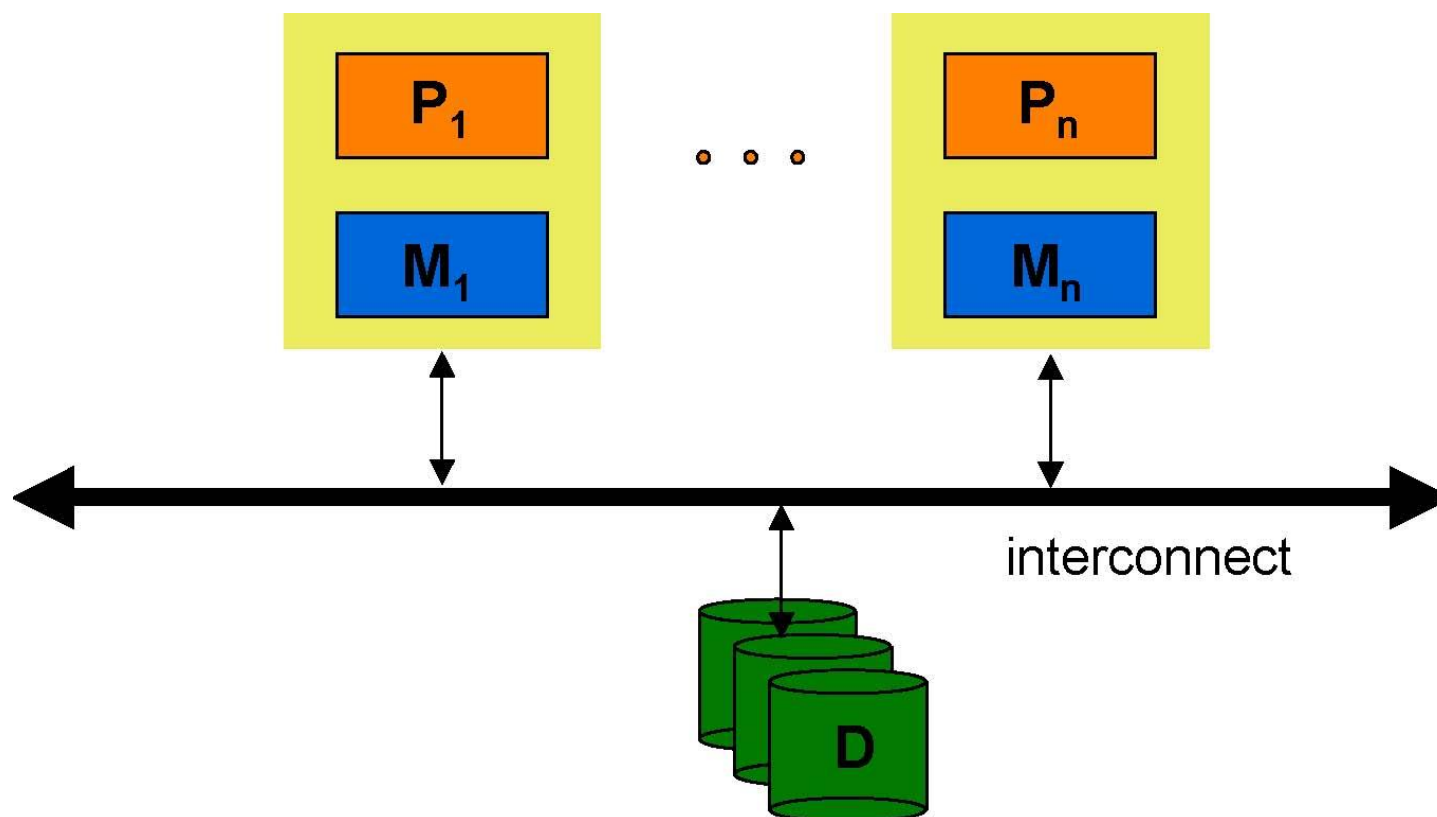
- ◆ Hierarchical (cluster)
- ◆ Non-Uniform Memory Architecture (NUMA)

Shared Memory Architectures



- Examples:
- DBMS on symmetric multiprocessors (Sequent, Encore, Sun, etc.)
 - ▶ Simplicity, load balancing, fast communication
 - ▶ Network cost, low extensibility

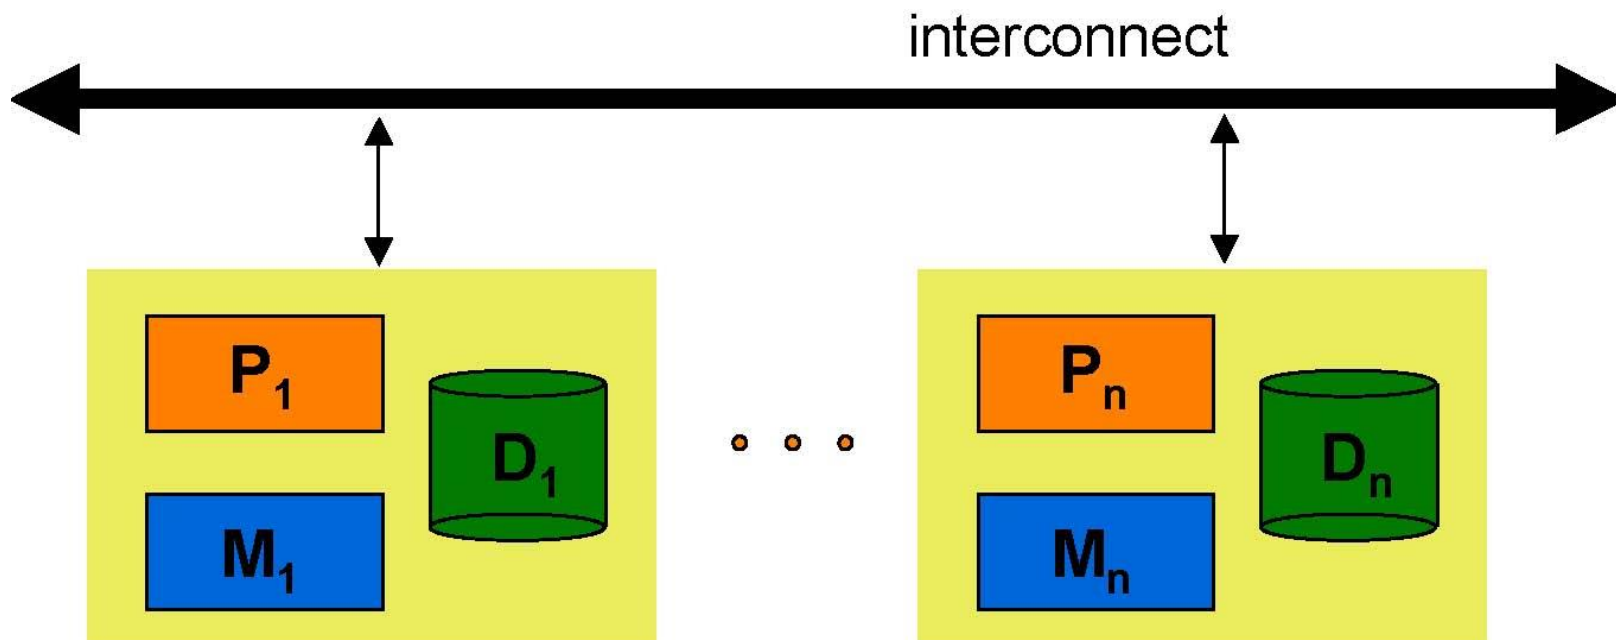
Shared-Disk Architecture



Examples: DEC's VAXcluster, IBM's IMS/VS Data Sharing

- network cost, extensibility, migration from uniprocessor
- complexity, potential performance problem for copy coherency

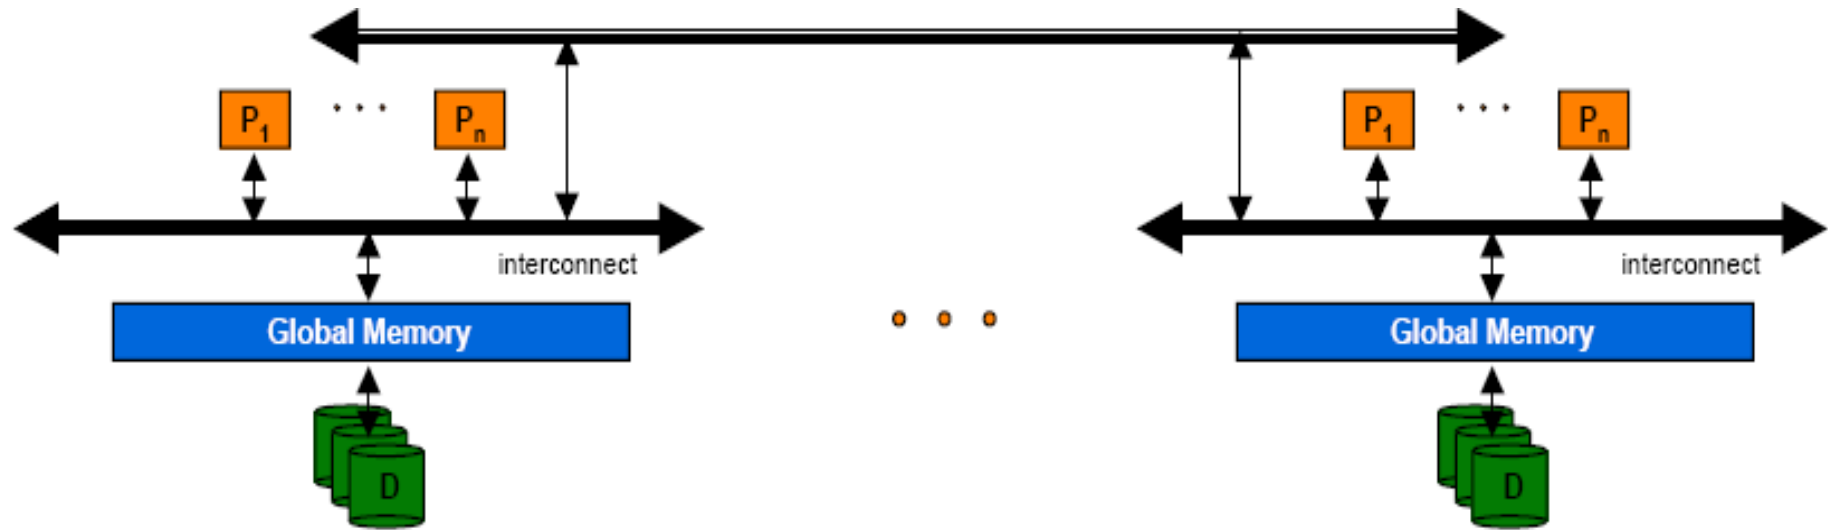
Shared-Nothing Architecture



Examples: Teradata (NCR), NonStopSQL (Tandem-Compaq),
Gamma (U. of Wisconsin), Bubba (MCC)

- Extensibility, availability
- Complexity, difficult load balancing

Hierarchical Architecture



- ❖ Combines good load balancing of SM with extensibility of SN

Shared-Memory vs. Distributed Memory

- ❖ Mix two different aspects : addressing and memory
 - ◆ Addressing
 - Single address space : Sequent, Encore, KSR
 - Multiple address spaces : Intel, Ncube
 - ◆ Physical memory
 - Central : Sequent, Encore
 - Distributed : Intel, Ncube, KSR
- ❖ NUMA (Non-Uniform Memory Architecture):
single address space on distributed physical memory
eases application portability
 - ◆ Extensibility

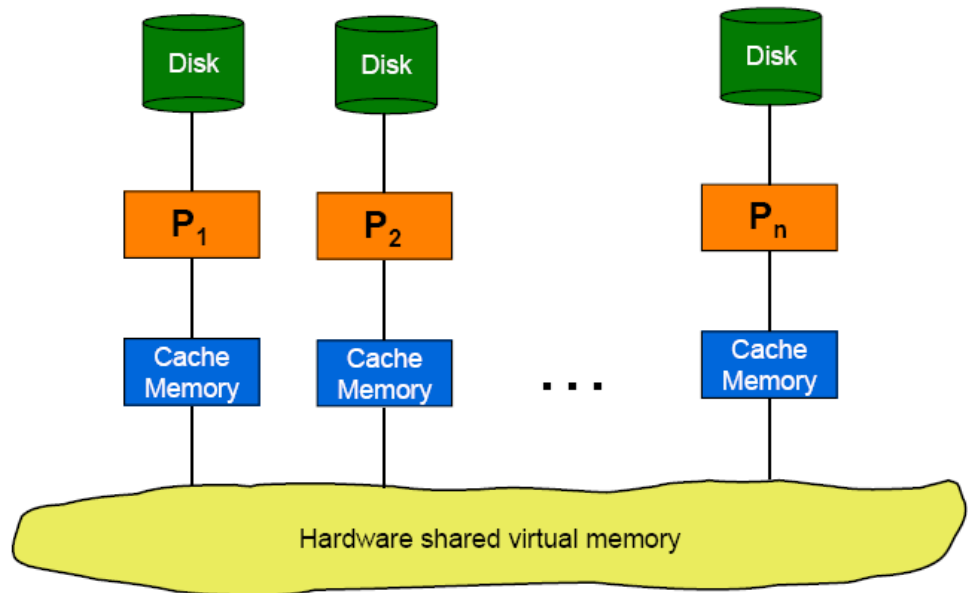
Non-Uniform Memory Architecture (NUMA) Architectures

❖ Cache Coherent NUMA (CC-NUMA)

- ♦ statically divide the main memory among the nodes

❖ Cache-Only Memory Architecture (COMA)

- ♦ convert the per-node memory into a large cache of the shared address space



Parallel DBMS Techniques

❖ Data placement

- ◆ Physical placement of the DB onto multiple nodes
- ◆ Static vs. Dynamic

❖ Parallel data processing

- ◆ Select is easy
- ◆ Join (and all other non-select operations) is more difficult

❖ Parallel query optimization

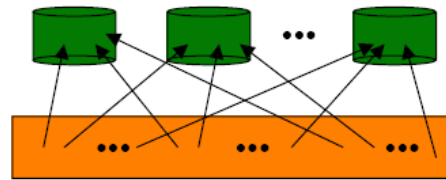
- ◆ Choice of the best parallel execution plans
- ◆ Automatic parallelization of the queries and load balancing

❖ Transaction management


- ◆ Similar to distributed transaction management

Data Partitioning


- ❖ Each relation is divided in n partitions (sub-relations), where n is a function of relation size and access frequency
- ❖ Implementation
 - ◆ Round-robin
 - Map the i -th element to node $(i \bmod n)$
 - Simple but only exact-match queries
 - ◆ B-tree index
 - Support range queries but large index size
 - ◆ Hash function
 - Only exact-match queries but small index size



Round-Robin



Hashing



Interval

Replicated Data Partitioning

❖ High-availability requires data replication

- ♦ interleaved or chained partitioning to achieve load balancing

Node	1	2	3	4
Primary copy	R1	R2	R3	R4
Backup copy		r 1.1	r 1.2	r 1.3
	r 2.3		r 2.1	r 2.2
	r 3.2	r 3.3		r 3.1

Interleaved Partitioning

Node	1	2	3	4
Primary copy	R1	R2	R3	R4
Backup copy	r4	r1	r2	r3

Chained Partitioning

Placement Directory

- ❖ Perform two functions

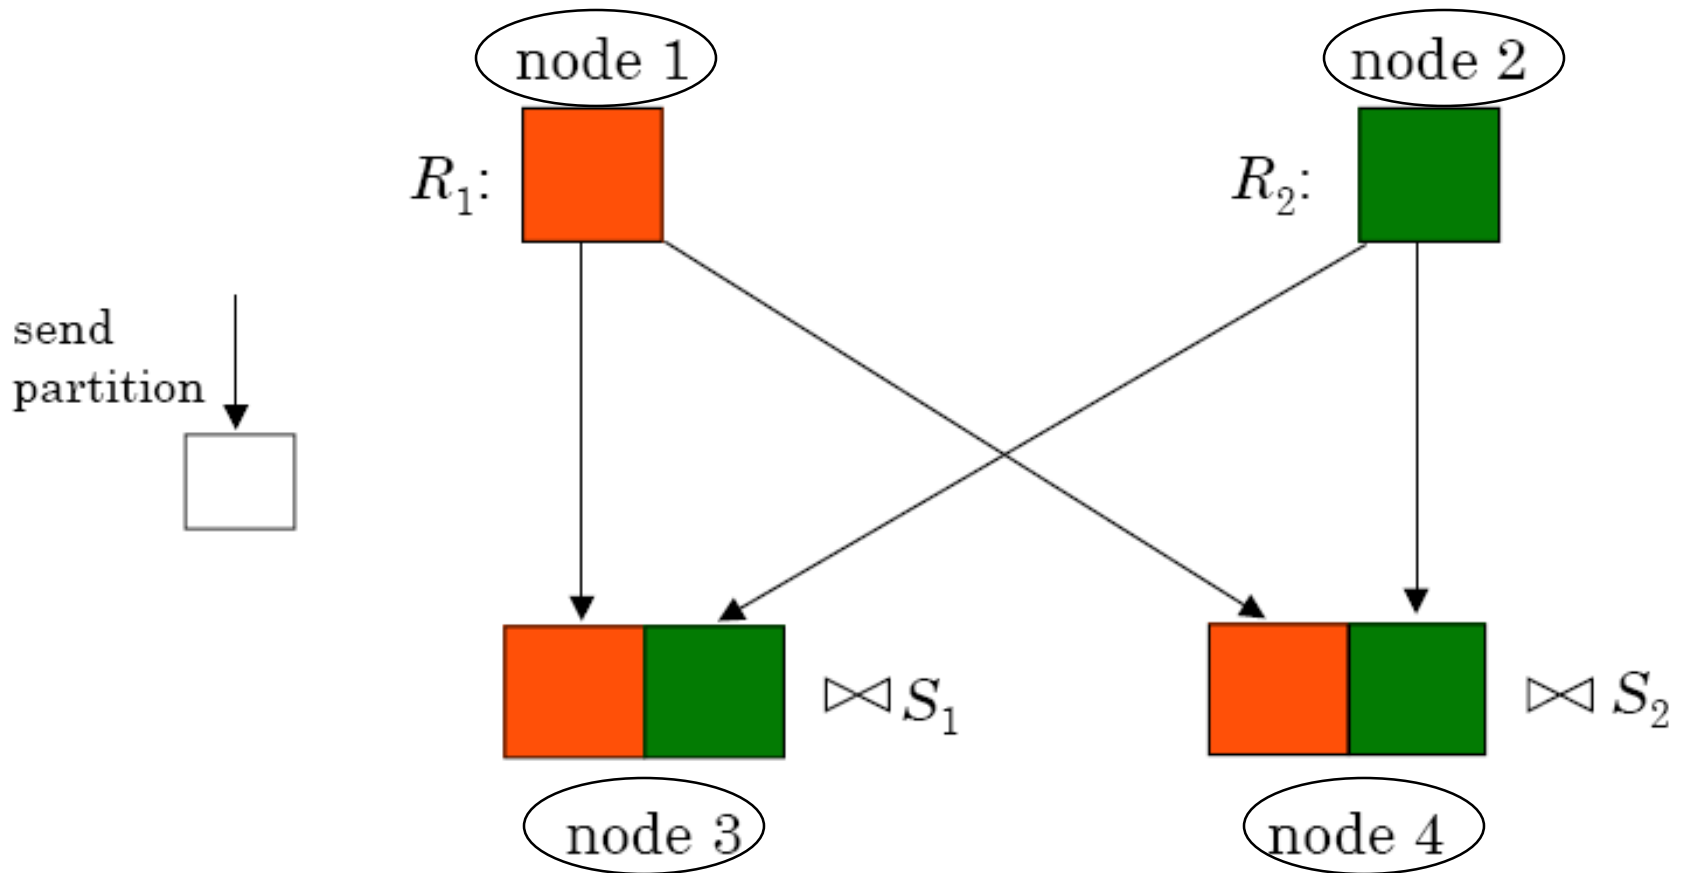
- ♦ $F1$ (relname, placement attval) = log node-id
- ♦ $F2$ (log node-id) = phy node-id

- ❖ In either case, the data structure for $F1$ and $F2$ should be available when needed at each node

Join Processing

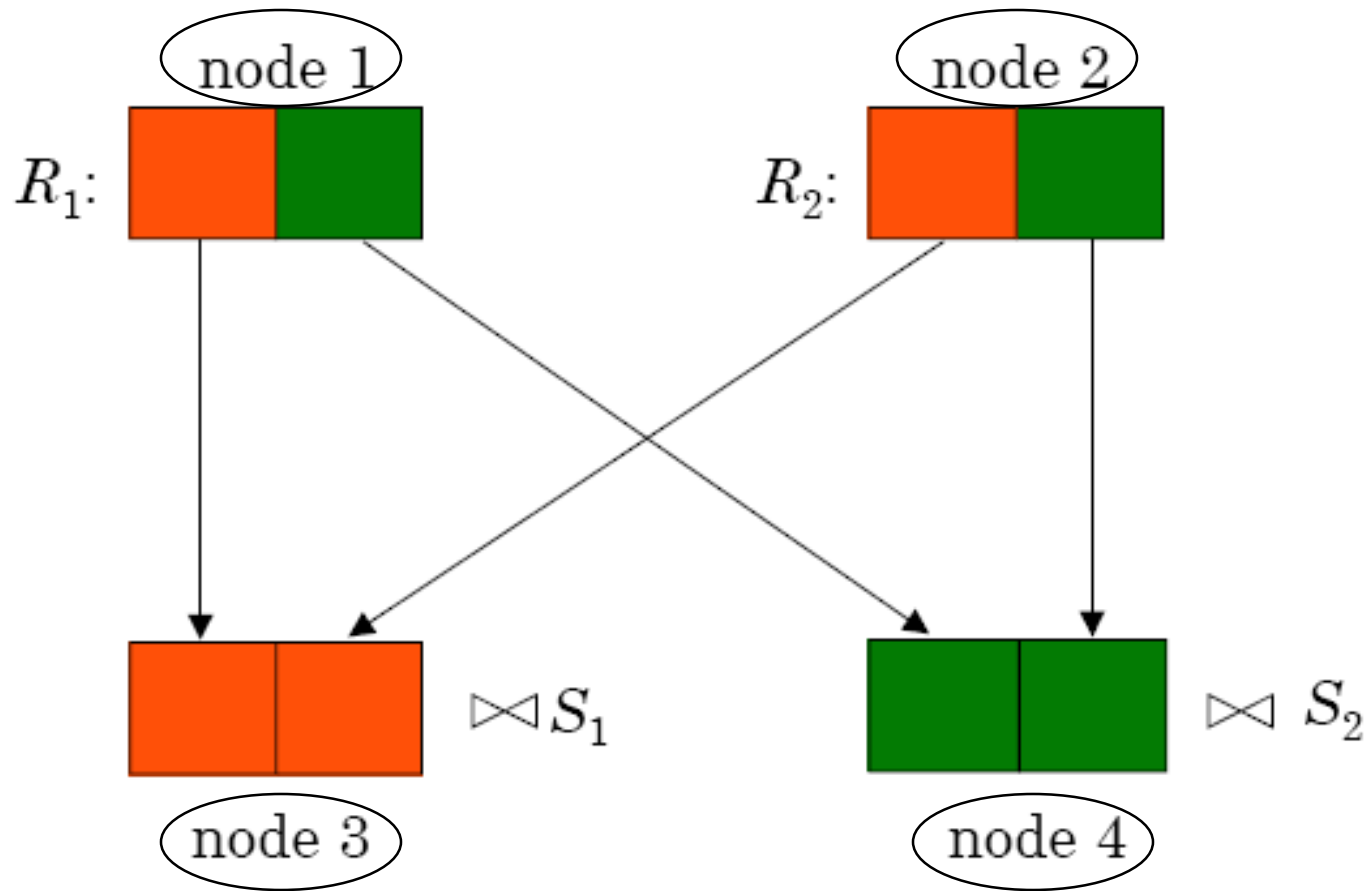
- ❖ Three basic algorithms for intra-operator parallelism
 - ◆ Parallel nested loop join: no special assumption
 - ◆ Parallel associative join: one relation is declustered on join attribute and equi-join
 - ◆ Parallel hash join: equi-join
- ❖ They also apply to other complex operators such as duplicate elimination, union, intersection, etc. with minor adaptation

Parallel Nested Loop Join



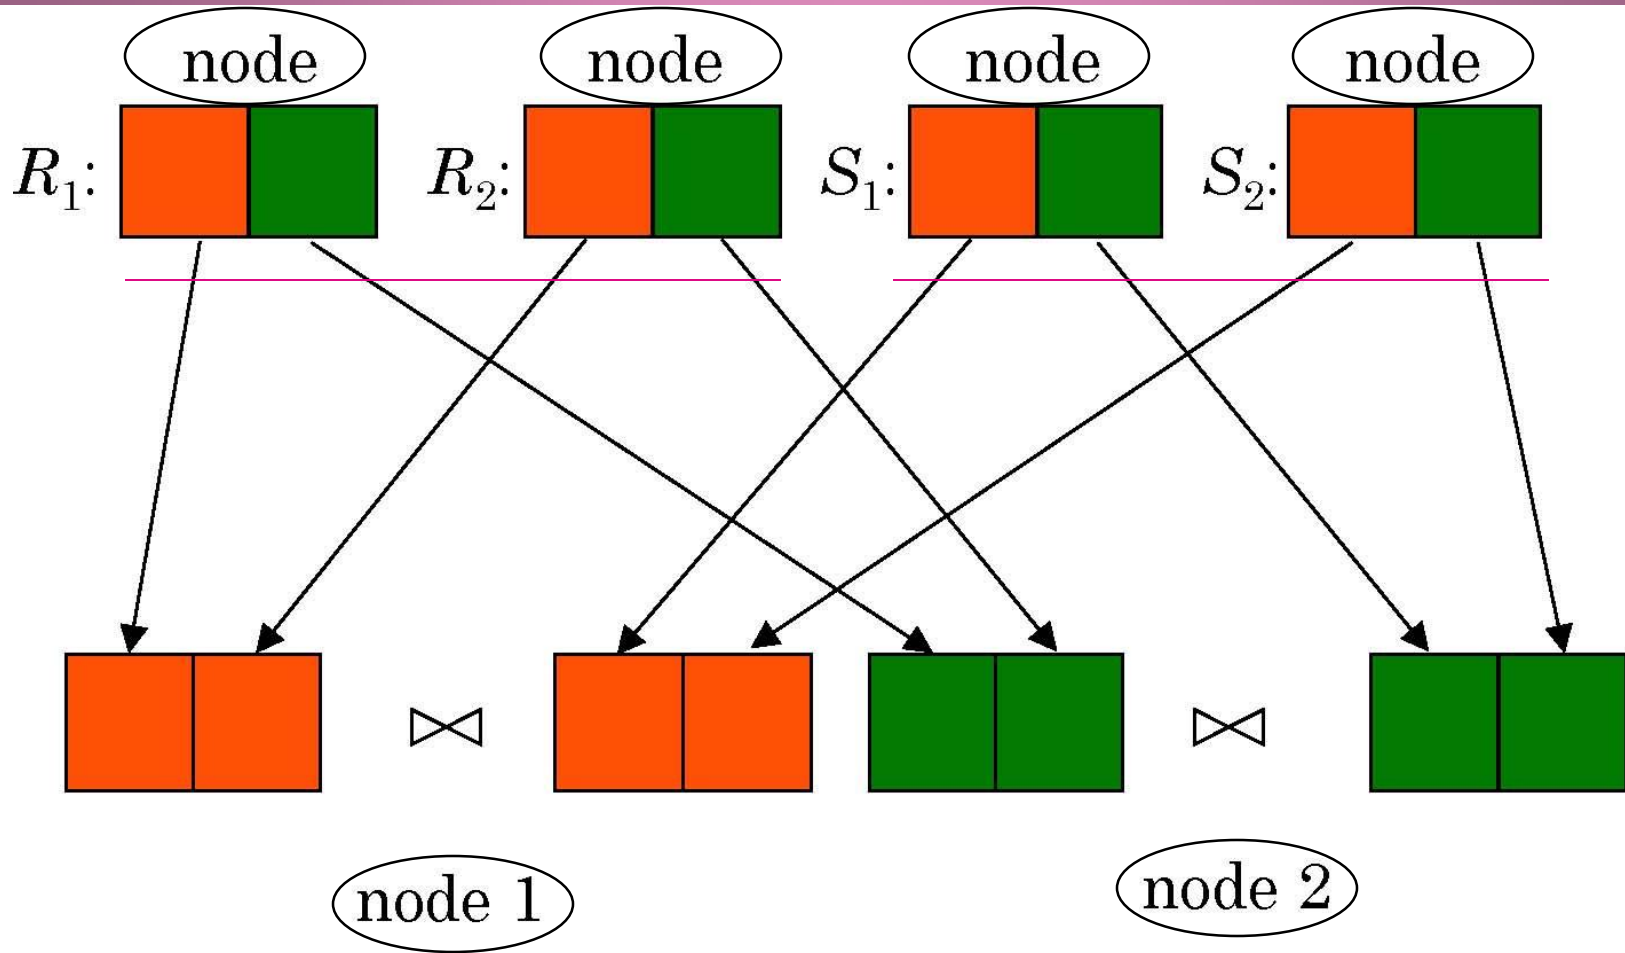
$$R \bowtie S \rightarrow \cup_{i=1,n} (R \bowtie S_i)$$

Parallel Associative Join



$$R \bowtie S \rightarrow \cup_{i=1,n} (R_i \bowtie S_i)$$

Parallel Hash Join

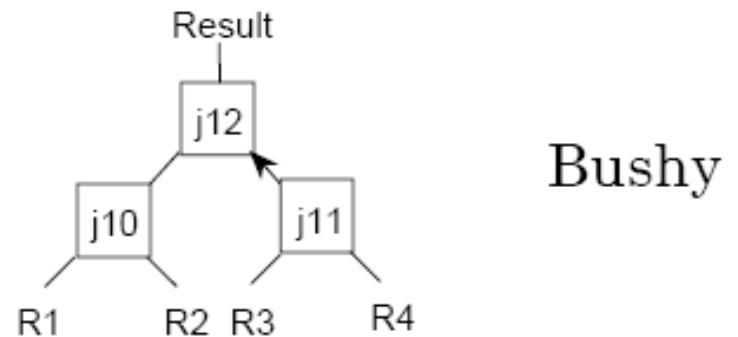
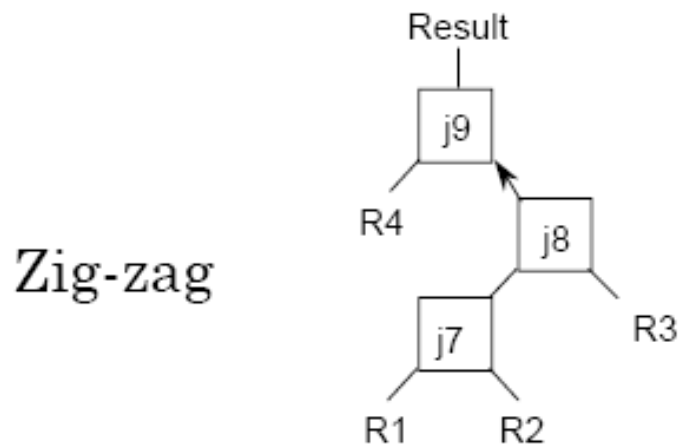
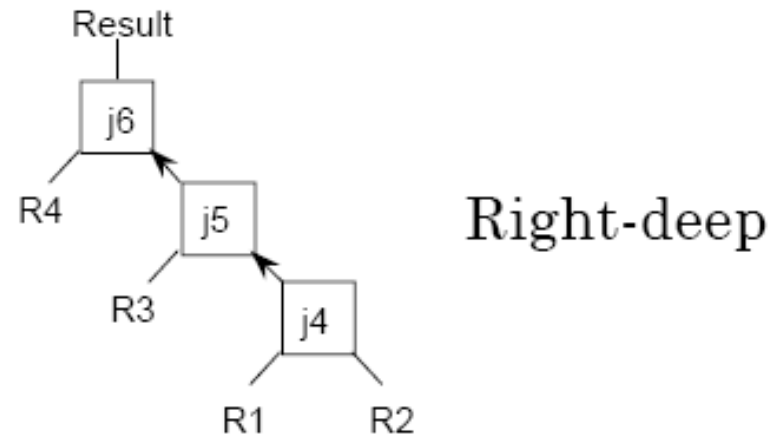
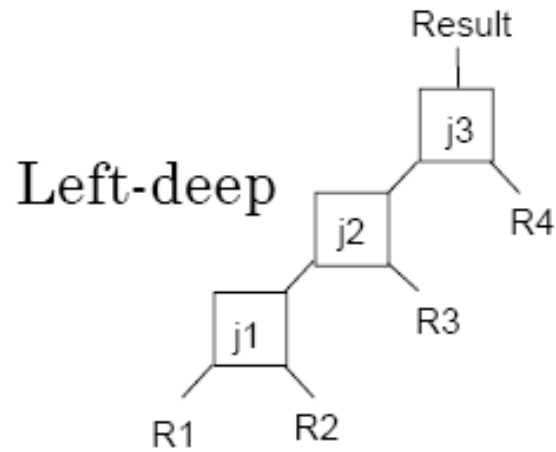


$$R \bowtie S \rightarrow \cup_{i=1,P} (R_i \bowtie S_i)$$

Parallel Query Optimization

- ❖ The objective is to select the "best" parallel execution plan for a query using the following components
- ❖ Search space
 - ◆ Model alternative execution plans as operator trees
 - ◆ Left-deep vs. Right-deep vs. Bushy trees
- ❖ Search strategy
 - ◆ Dynamic programming for small search space
 - ◆ Randomized for large search space
- ❖ Cost model (abstraction of execution system)
 - ◆ Physical schema info. (partitioning, indexes, etc.)
 - ◆ Statistics and cost functions

Execution Plans as Operators Trees



Load Balancing

❖ Problems arise for intra-operator parallelism with skewed data distributions

- ◆ attribute data skew (AVS)
- ◆ tuple placement skew (TPS)
- ◆ selectivity skew (SS)
- ◆ redistribution skew (RS)
- ◆ join product skew (JPS)

❖ Solutions

- ◆ sophisticated parallel algorithms that deal with skew
- ◆ dynamic processor allocation (at execution time)

Some Parallel DBMSs

❖ Prototypes

- ◆ EDS and DBS3 (ESPRIT)
- ◆ Gamma (U. of Wisconsin)
- ◆ Bubba (MCC, Austin, Texas)
- ◆ XPRS (U. of Berkeley)
- ◆ GRACE (U. of Tokyo)

❖ Products

- ◆ Teradata (NCR)
- ◆ NonStopSQL (Tandem-Compac)
- ◆ DB2 (IBM), Oracle, Informix, Ingres, Navigator (Sybase) ...

Research Problems

- ❖ Hybrid architectures
- ❖ OS support: using micro-kernels
- ❖ Benchmarks to stress speedup and scale up under mixed workloads
- ❖ Data placement to deal with skewed data distribution and data replication
- ❖ Parallel data languages to specify independent and pipelined parallelism
- ❖ Parallel query optimization to deal with mix of precompiled queries and complex ad-hoc queries
- ❖ Support of higher functionality such as rules and objects

11. Streaming Data Management

Chapter 18

Streaming Data

Finding a Database Problem

- ❖ Pick a simple but fundamental assumption underlying traditional database systems
 - ◆ Drop it
- ❖ Reconsider all aspects of data management and query processing
 - ◆ Many Ph.D. theses
 - ◆ Prototype from scratch

Facts

❖ Dropped assumptions

- ◆ Data has a fixed schema declared in advance
- ◆ All data is accurate, consistent, and complete
- ◆ First load data, then index it, then run queries
 - Continuous data streams
 - Continuous queries

Streaming Data

- ❖ Continuous, unbounded, rapid, time-varying streams of data elements
- ❖ Occurring in a variety of modern applications
 - ◆ Network monitoring and traffic engineering
 - ◆ Sensor networks, RFID tags
 - ◆ Telecom call records
 - ◆ Financial applications
 - ◆ Web logs and click-streams
 - ◆ Manufacturing processes
- ❖ **DSMS** = Data Stream Management System

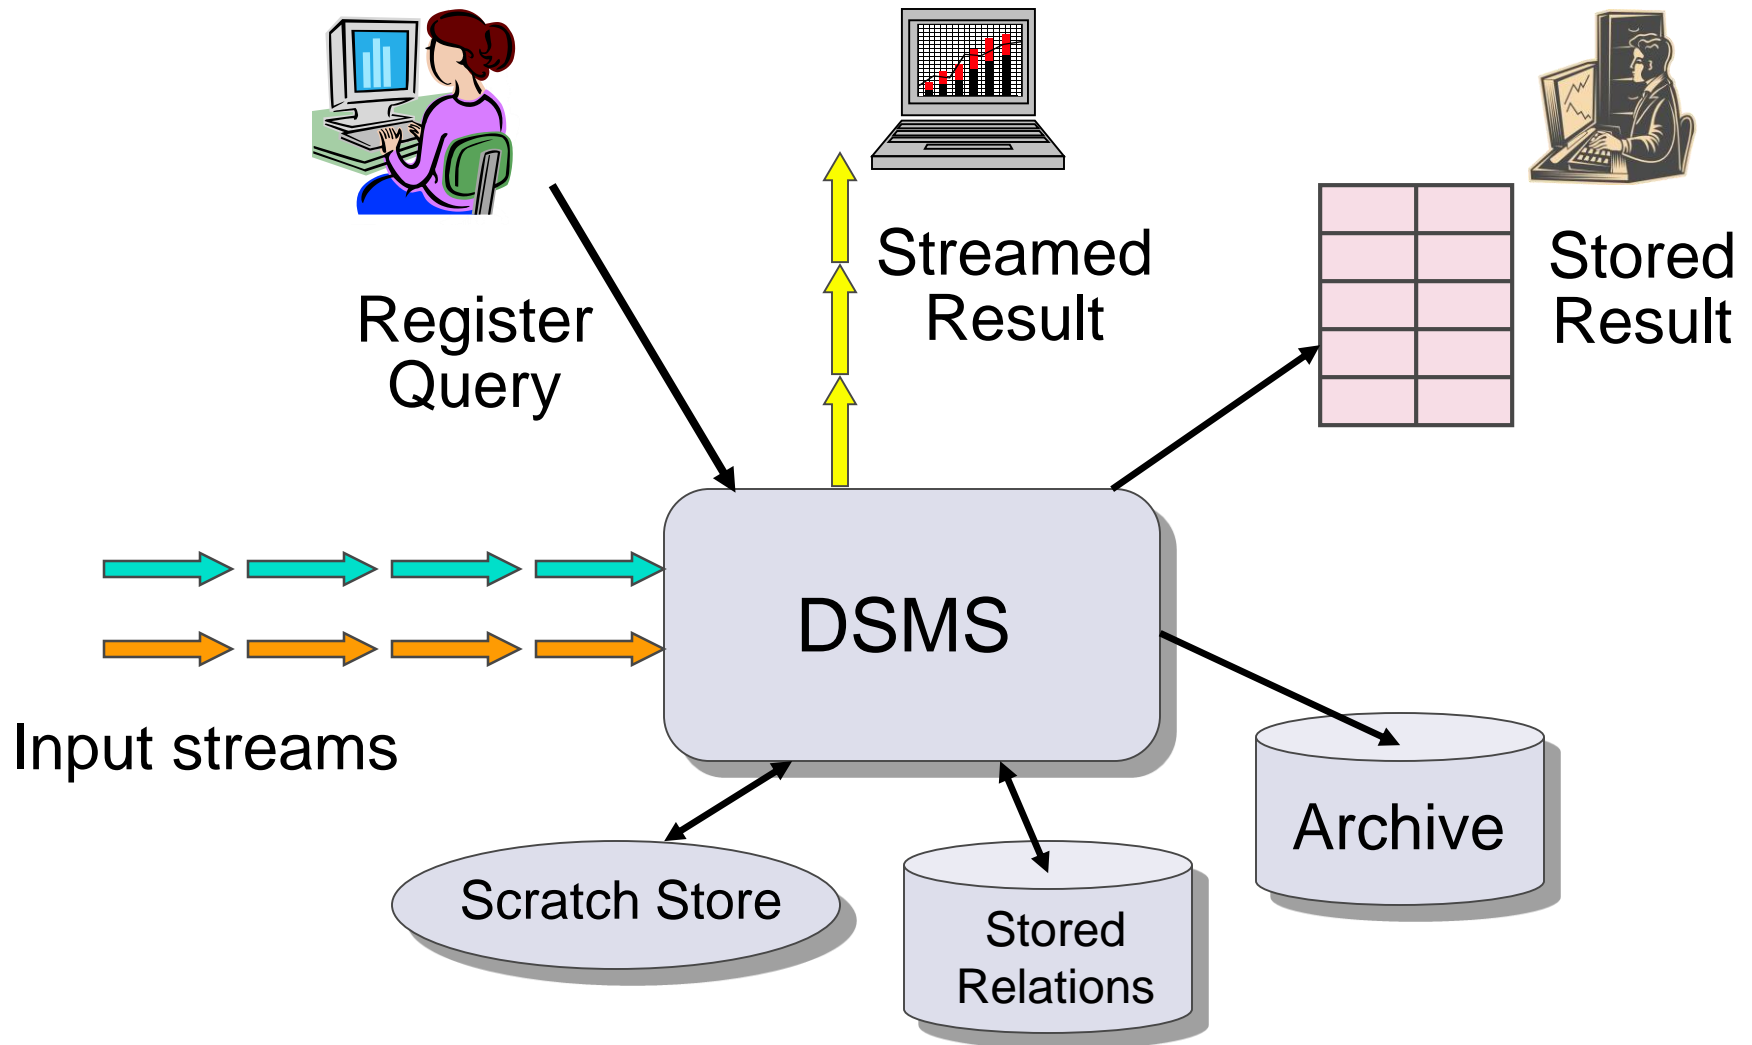
DBMS versus DSMS

- ❖ Persistent relations
- ❖ One-time queries
- ❖ Random access
- ❖ Access plan determined by query processor and physical DB design
- ❖ Transient streams (and persistent relations)
- ❖ Continuous queries
- ❖ Sequential access
- ❖ Unpredictable data characteristics and arrival patterns

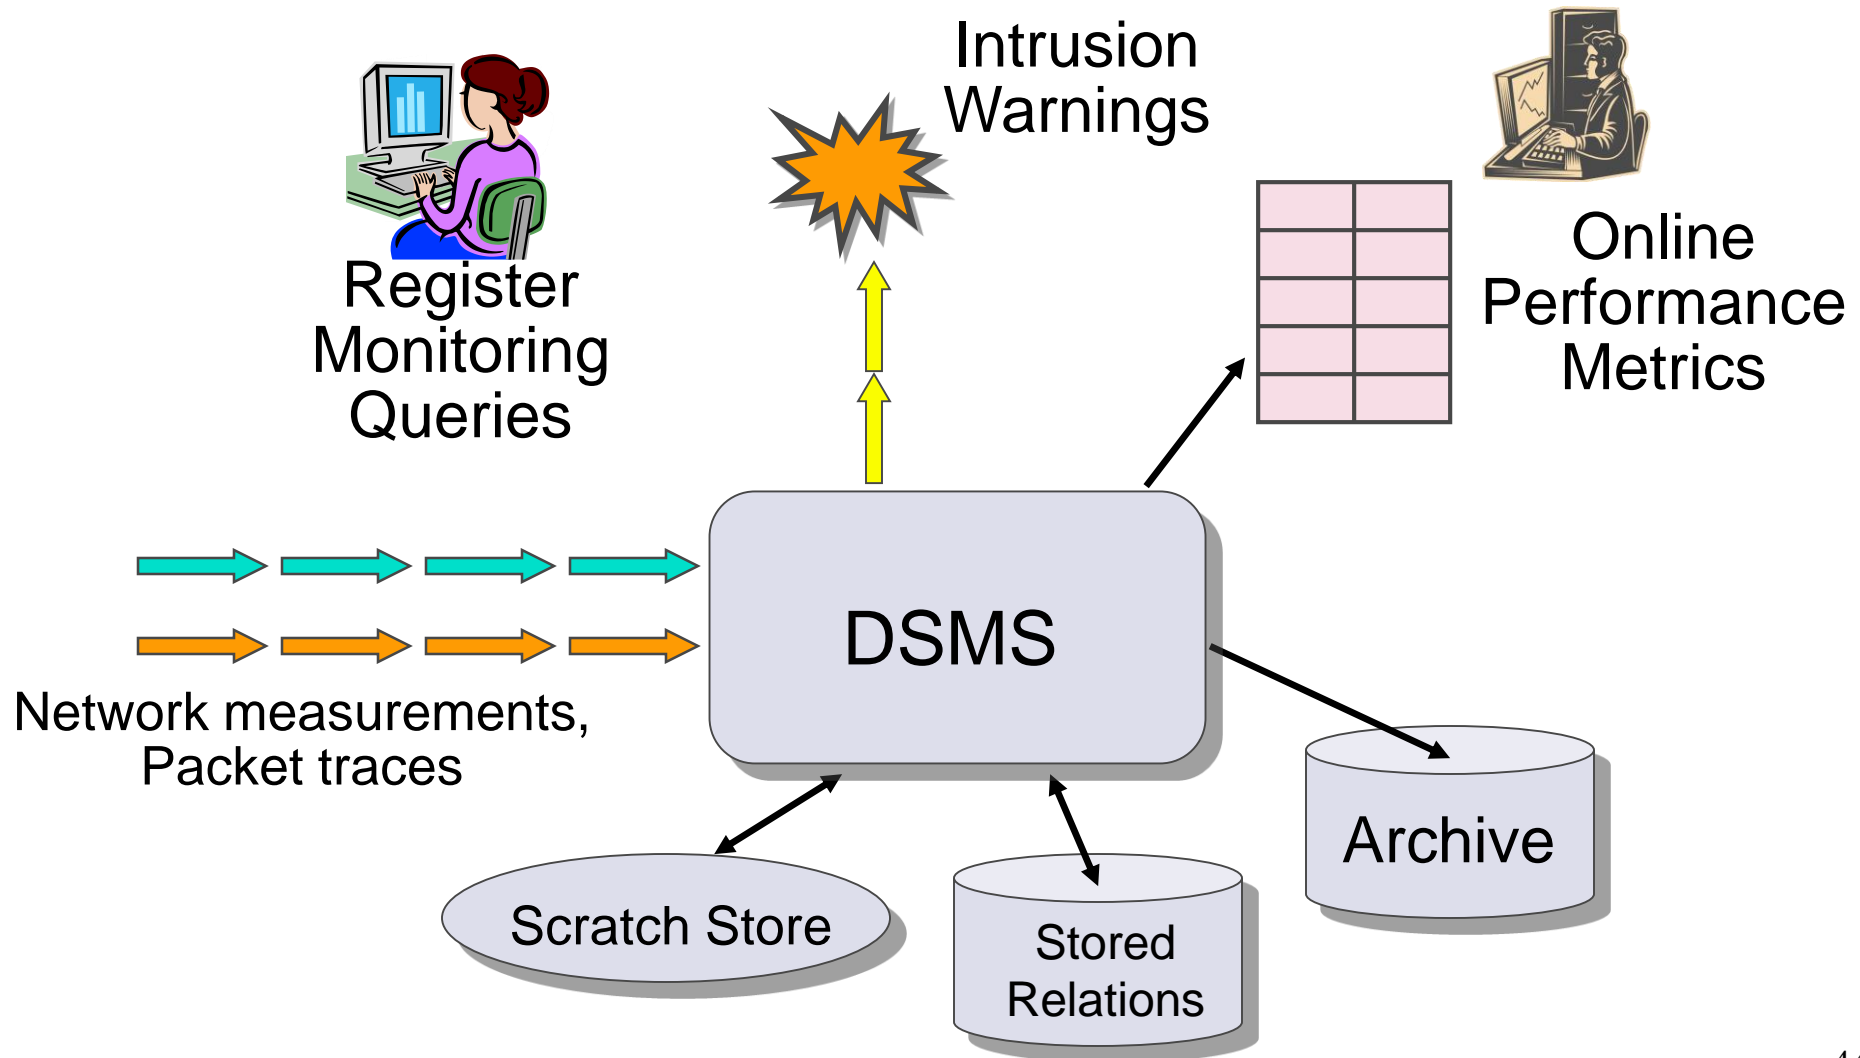
Continuous Queries

- ❖ One time queries – run once to completion over the current data set.
- ❖ Continuous queries – issued once and continuously evaluated over the data, e.g.,
 - ◆ Notify me when the temperature drops below X
 - ◆ Tell me when prices of stock Y > 300

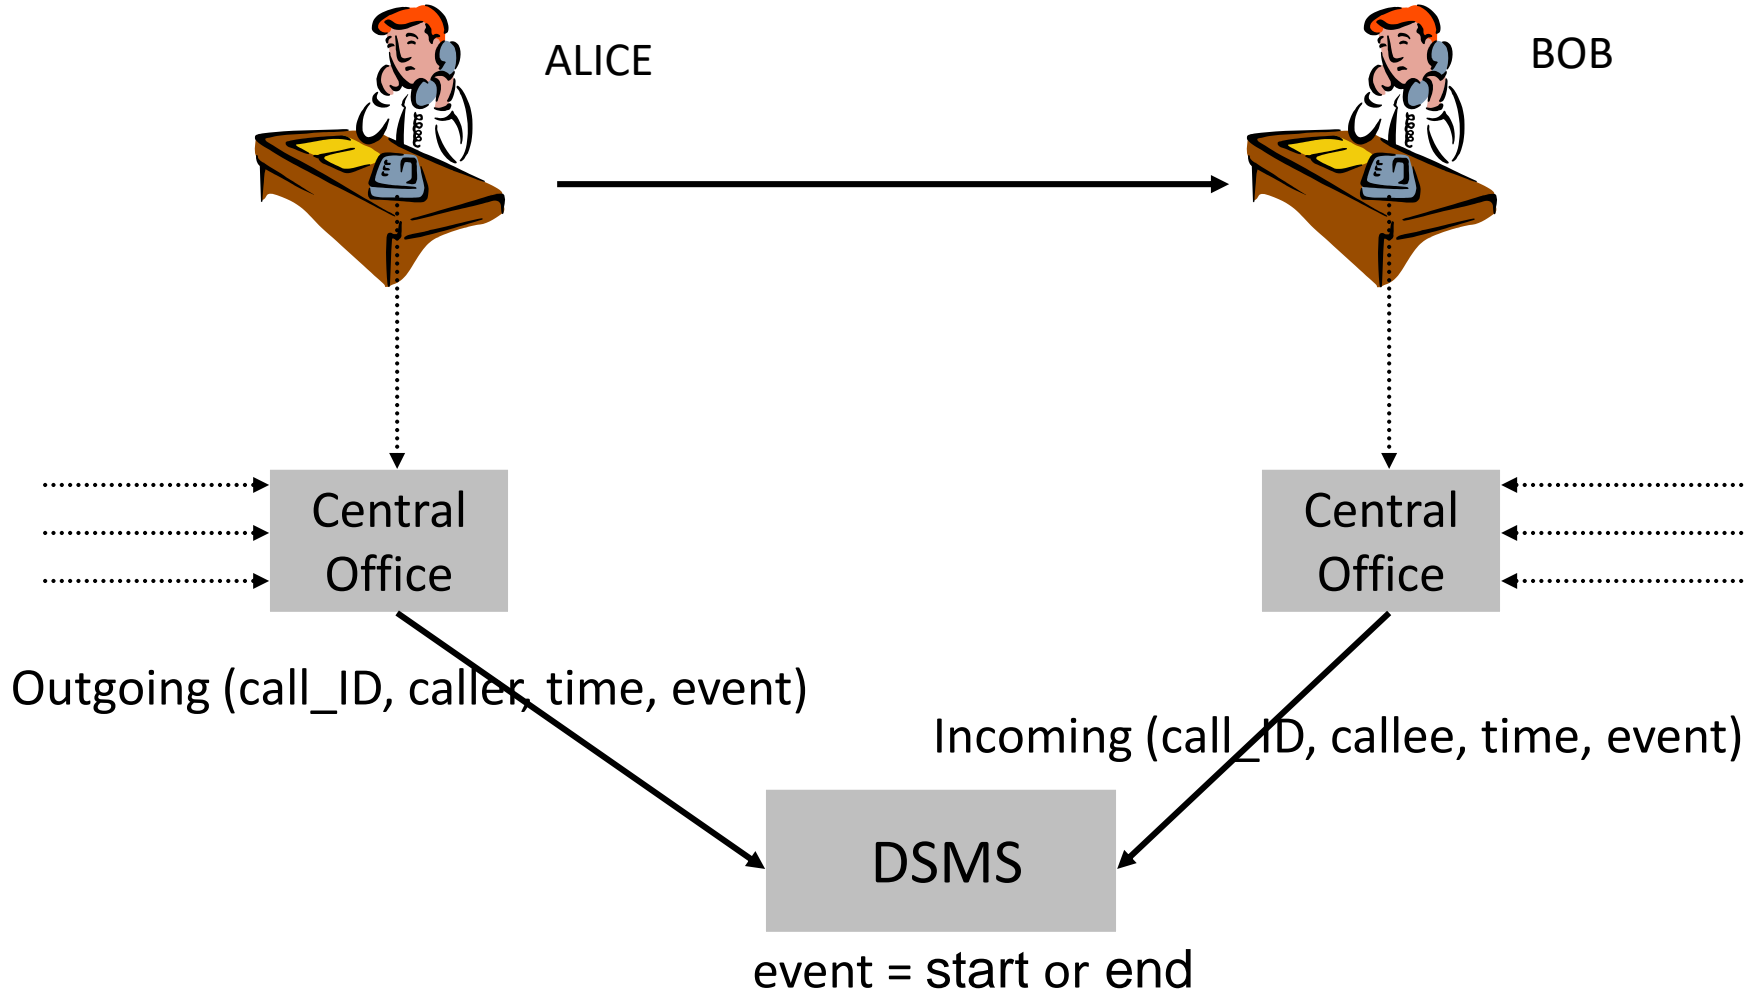
The (Simplified) Big Picture



(Simplified) Network Monitoring



Making Things Concrete



Query 1 (SELF-JOIN)

- ❖ Find all outgoing calls longer than 2 minutes

```
SELECT  O1.call_ID, O1.caller
FROM    Outgoing O1, Outgoing O2
WHERE   (O2.time – O1.time > 2)
        AND  O1.call_ID = O2.call_ID
        AND  O1.event = start
        AND  O2.event = end )
```

- ❖ Result requires unbounded storage
- ❖ Can provide result as data stream
- ❖ Can output after 2 min, without seeing end

Query 2 (JOIN)

- ❖ Pair up callers and callees

```
SELECT  O.caller, I.callee  
FROM    Outgoing O, Incoming I  
WHERE   O.call_ID = I.call_ID
```

- ❖ Can still provide result as data stream
- ❖ Requires unbounded temporary storage ...

Query 3 (group-by aggregation)

❖ Total connection time for each caller

```
SELECT      O1.caller, sum(O2.time – O1.time)
FROM        Outgoing O1, Outgoing O2
WHERE       (O1.call_ID = O2.call_ID
            AND O1.event = start
            AND O2.event = end )

GROUP BY    O1.caller
```

❖ Cannot provide result in (append-only) stream

- ♦ Output **updates?**
- ♦ Provide current value **on demand?**
- ♦ **Memory?**

DSMS – Architecture & Issues

- ❖ Data streams and stored relations – architectural differences
- ❖ Declarative language for registering continuous queries
- ❖ Flexible query plans and execution strategies
- ❖ Centralized ? Distributed ?

DSMS – Options

- ❖ Relation: Tuple Set or Sequence?
- ❖ Update: Modification or Append?
- ❖ Query Answer: Exact or Approximate?
- ❖ Query Evaluation: One or Multiple Pass?
- ❖ Query Plan: Fixed or Adaptive?

Architectural Comparison

DSMS

- ❖ Resource (memory, per-tuple computation) limited
- ❖ Reasonably complex, near real time, query processing
- ❖ Useful to identify what data to populate in database
- ❖ Query evaluation: one pass
- ❖ Query plan: adaptive

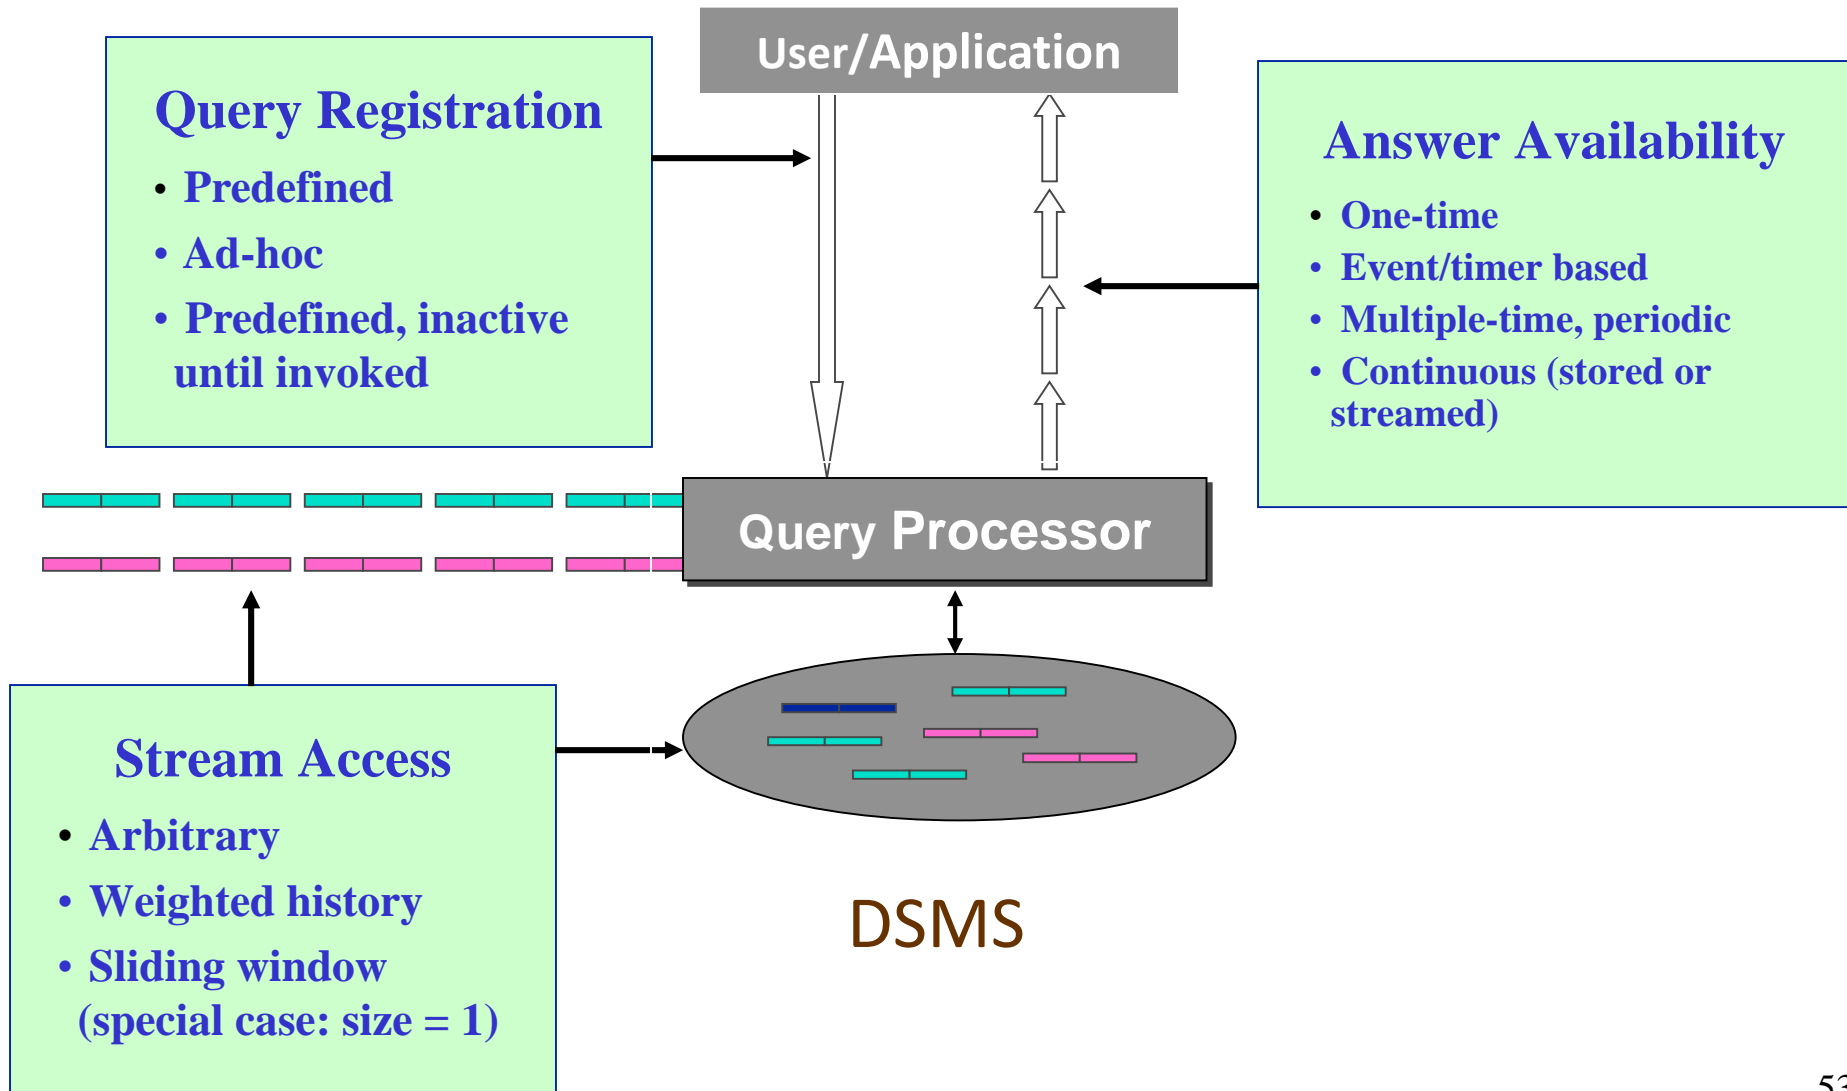
DBMS

- ❖ Resource (memory, disk, per-tuple computation) rich
- ❖ Extremely sophisticated query processing, analysis
- ❖ Useful to audit query results of database systems
- ❖ Query evaluation: arbitrary
- ❖ Query plan: fixed

DSMS Challenges

- ❖ Must cope with:
 - ◆ **Stream rates** that may be high, variable, bursty
 - ◆ **Stream data** that may be unpredictable, variable
 - ◆ **Continuous query loads** that may be high, variable
- ❖ Overload – need to use resources very carefully
- ❖ Changing conditions – adaptive strategy

Query Model




Query Processing

- ❖ Query Language
- ❖ Operators
- ❖ Optimization
- ❖ Multi-Query Optimization

Stream Query Language

- ❖ SQL extension
- ❖ Queries reference/produce relations or streams
- ❖ Examples: GSQL, CQL



Continuous Query Language – CQL

Start with SQL

Then add...

- ❖ Streams as new data type
- ❖ Continuous instead of one-time semantics
- ❖ Windows on streams (derived from SQL-99)
- ❖ Sampling on streams (basic)

Impact of Limited Memory

- ❖ Continuous streams grow unboundedly
- ❖ Queries may require unbounded memory
- ❖ One solution: Approximate query evaluation

Approximate Query Evaluation

❖ Why?

- ◆ Handling load – streams coming too fast
- ◆ Avoid unbounded storage and computation
- ◆ Ad hoc queries need approximate history

❖ How?

- ◆ Sliding windows, synopsis, samples, load-shedding

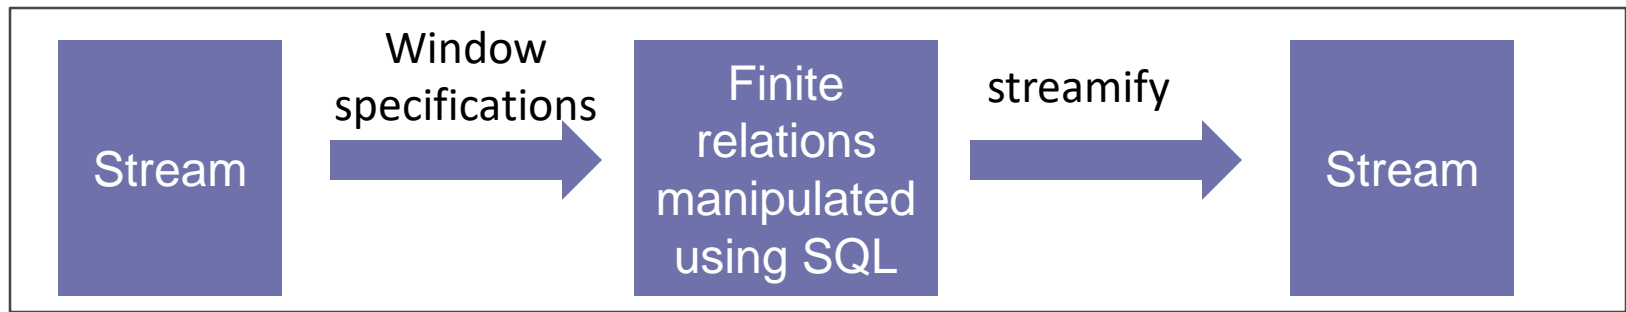
Approximate Query Evaluation (*cont.*)

❖ Major Issues

- ♦ Metric for set-valued queries
- ♦ Composition of approximate operators
- ♦ How is it understood/controlled by user?
- ♦ Integrate into query language
- ♦ Query planning and interaction with resource allocation
- ♦ **Accuracy-efficiency-storage tradeoff and global metric**

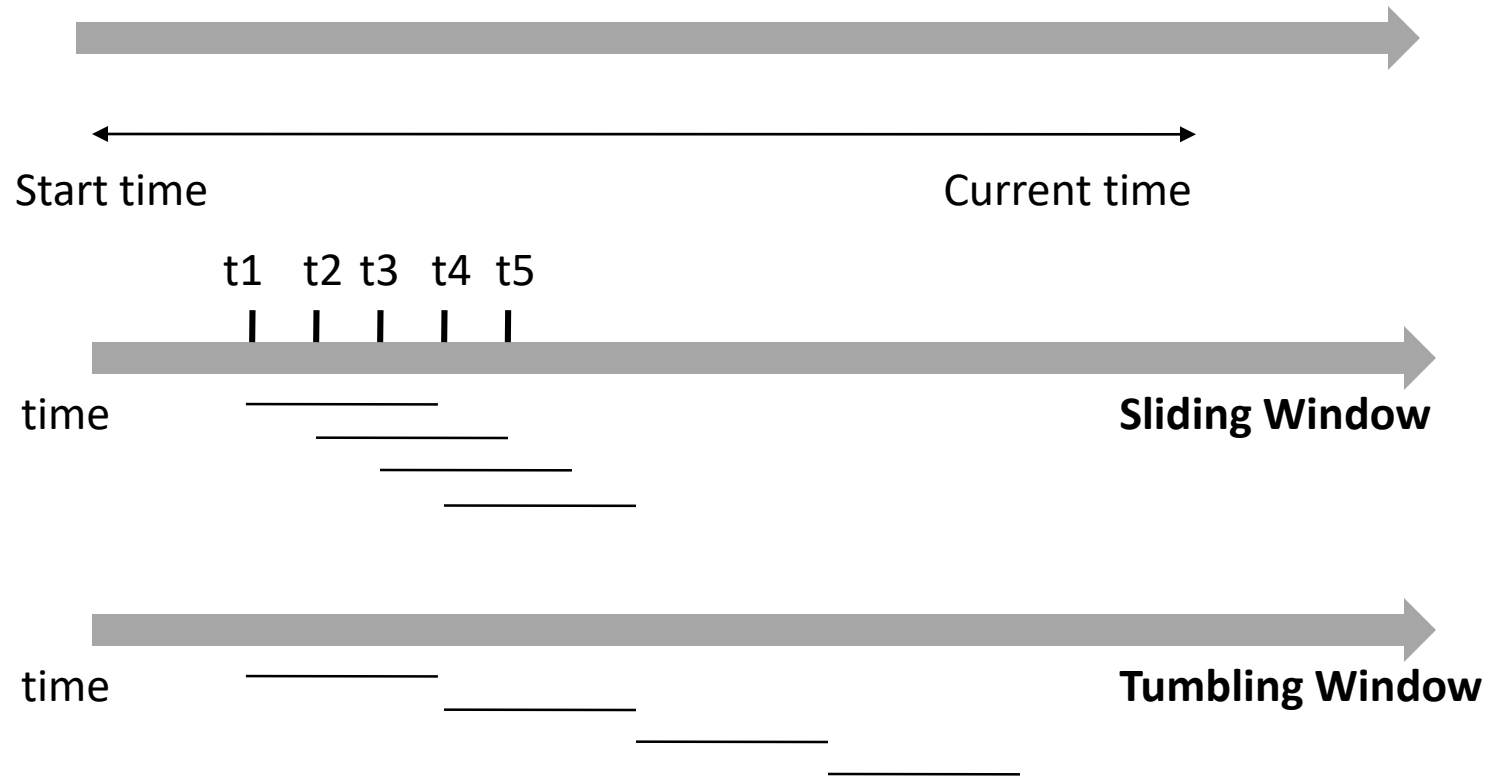
Windows

- ❖ Mechanism for extracting a finite relation from an infinite stream
- ❖ Various window proposals for restricting operator scope
 - ◆ Windows based on ordering attribute (e.g., time)
 - ◆ Windows based on tuple counts
 - ◆ Windows based on explicit markers (e.g. punctuations)
 - ◆ Variants (e.g., partitioning tuples in a window)



Windows (cont.)

❖ Terminology



Query Operators

- ❖ Selection - **Where** clause
- ❖ Projection - **Select** clause
- ❖ Join - **From** clause
- ❖ Group-by (Aggregation) – **Group-by** clause

Query Operators (*cont.*)

- ❖ Selection and projection on streams - straightforward
 - ◆ Local per-element operators
- ❖ Project may need to include ordering attribute
- ❖ Join – Problematic
 - ◆ May need to join tuples that are arbitrarily far apart
 - ◆ Equijoin on stream ordering attributes may be tractable
- ❖ Majority of the work focuses on join using windows.

Blocking Operators

❖ Blocking

- ◆ No output until entire input seen
- ◆ **Streams** – input never ends

❖ Simple Aggregate – output “update” stream

❖ Set Output (sort, group-by)

- ◆ **Root** – could maintain output data structure
- ◆ **Intermediate nodes** – try non-blocking analogs

❖ Join

- ◆ Apply sliding-window restrictions

Optimization in DSMS

- ❖ Traditionally table-based cardinalities used in query optimizer.
 - ◆ Goal of query optimizer: Minimize the size of intermediate results.
- ❖ Problematic in a streaming environment – All streams are unbounded = infinite size!
- ❖ Need novel optimization objectives that are relevant when the input sources are streams.

Query Optimization in DSMS

❖ Novel notions of optimization:

- ◆ Stream rate based [e.g. NiagaraCQ]
- ◆ QoS based [e.g. Aurora]

❖ Continuous adaptive optimization

❖ Possibilities that objectives cannot be met:

- ◆ Resource constraints
- ◆ Bursty arrivals under limited processing capabilities

Typical Stream Projects

- ❖ Amazon/Cougar (Cornell) – sensors
- ❖ Aurora (Brown/MIT) – sensor monitoring, dataflow
- ❖ Hancock (AT&T) – telecom streams
- ❖ Niagara (OGI/Wisconsin) – Internet XML databases
- ❖ OpenCQ (Georgia) – triggers, incr. view maintenance
- ❖ Stream (Stanford) – general-purpose DSMS
- ❖ Tapestry (Xerox) – pub/sub content-based filtering
- ❖ Telegraph (Berkeley) – adaptive engine for sensors
- ❖ Tribeca (Bellcore) – network monitoring
- ❖

Conclusion

- ❖ Conventional DMS technology is inadequate.
- ❖ We need to reconsider all aspects of data management in presence of streaming data.

Question & Answer