

Stat 212b: Topics in Deep Learning

Lecture 2

Joan Bruna
UC Berkeley



Objectives

- Classification, Kernels and metrics
- Representations for recognition
 - curse of dimensionality
 - invariance/covariance
 - discriminability
- Variability models
 - transformation groups and symmetries
 - deformations
 - stationarity
 - clutter and class-specific
- Examples

High-dimensional Recognition Setup

- Input data x lives in a high-dimensional space:

$x \in \Omega, \Omega \subset \mathbb{R}^d$ finite-dimensional (but large $d!$)

$x \in L^2(\mathbb{R}^m), m = 1, 2, 3$. infinite dimensional

High-dimensional Recognition Setup

- Input data x lives in a high-dimensional space:

$x \in \Omega, \Omega \subset \mathbb{R}^d$ finite-dimensional (but large $d!$)

$x \in L^2(\mathbb{R}^m), m = 1, 2, 3$. infinite dimensional

- We observe (x_i, y_i) , $i = 1 \dots n$, where

$y_i \in \mathbb{R}$ (regression)

$y_i \in \{1, K\}$. (classification)

High-dimensional Recognition Setup

- Input data x lives in a high-dimensional space:

$$x \in \Omega, \Omega \subset \mathbb{R}^d \quad \text{finite-dimensional (but large } d\text{!)}$$

$$x \in L^2(\mathbb{R}^m), m = 1, 2, 3 . \text{ infinite dimensional}$$

- We observe (x_i, y_i) , $i = 1 \dots n$, where


$$y_i \in \mathbb{R} \quad (\text{regression})$$

$$y_i \in \{1, K\} . \text{ (classification)}$$

- We can reduce the former to “interpolating” a function
 $f : \Omega \rightarrow \mathbb{R}^K \quad (f(x) = p(y | x) \text{ in the classification case})$

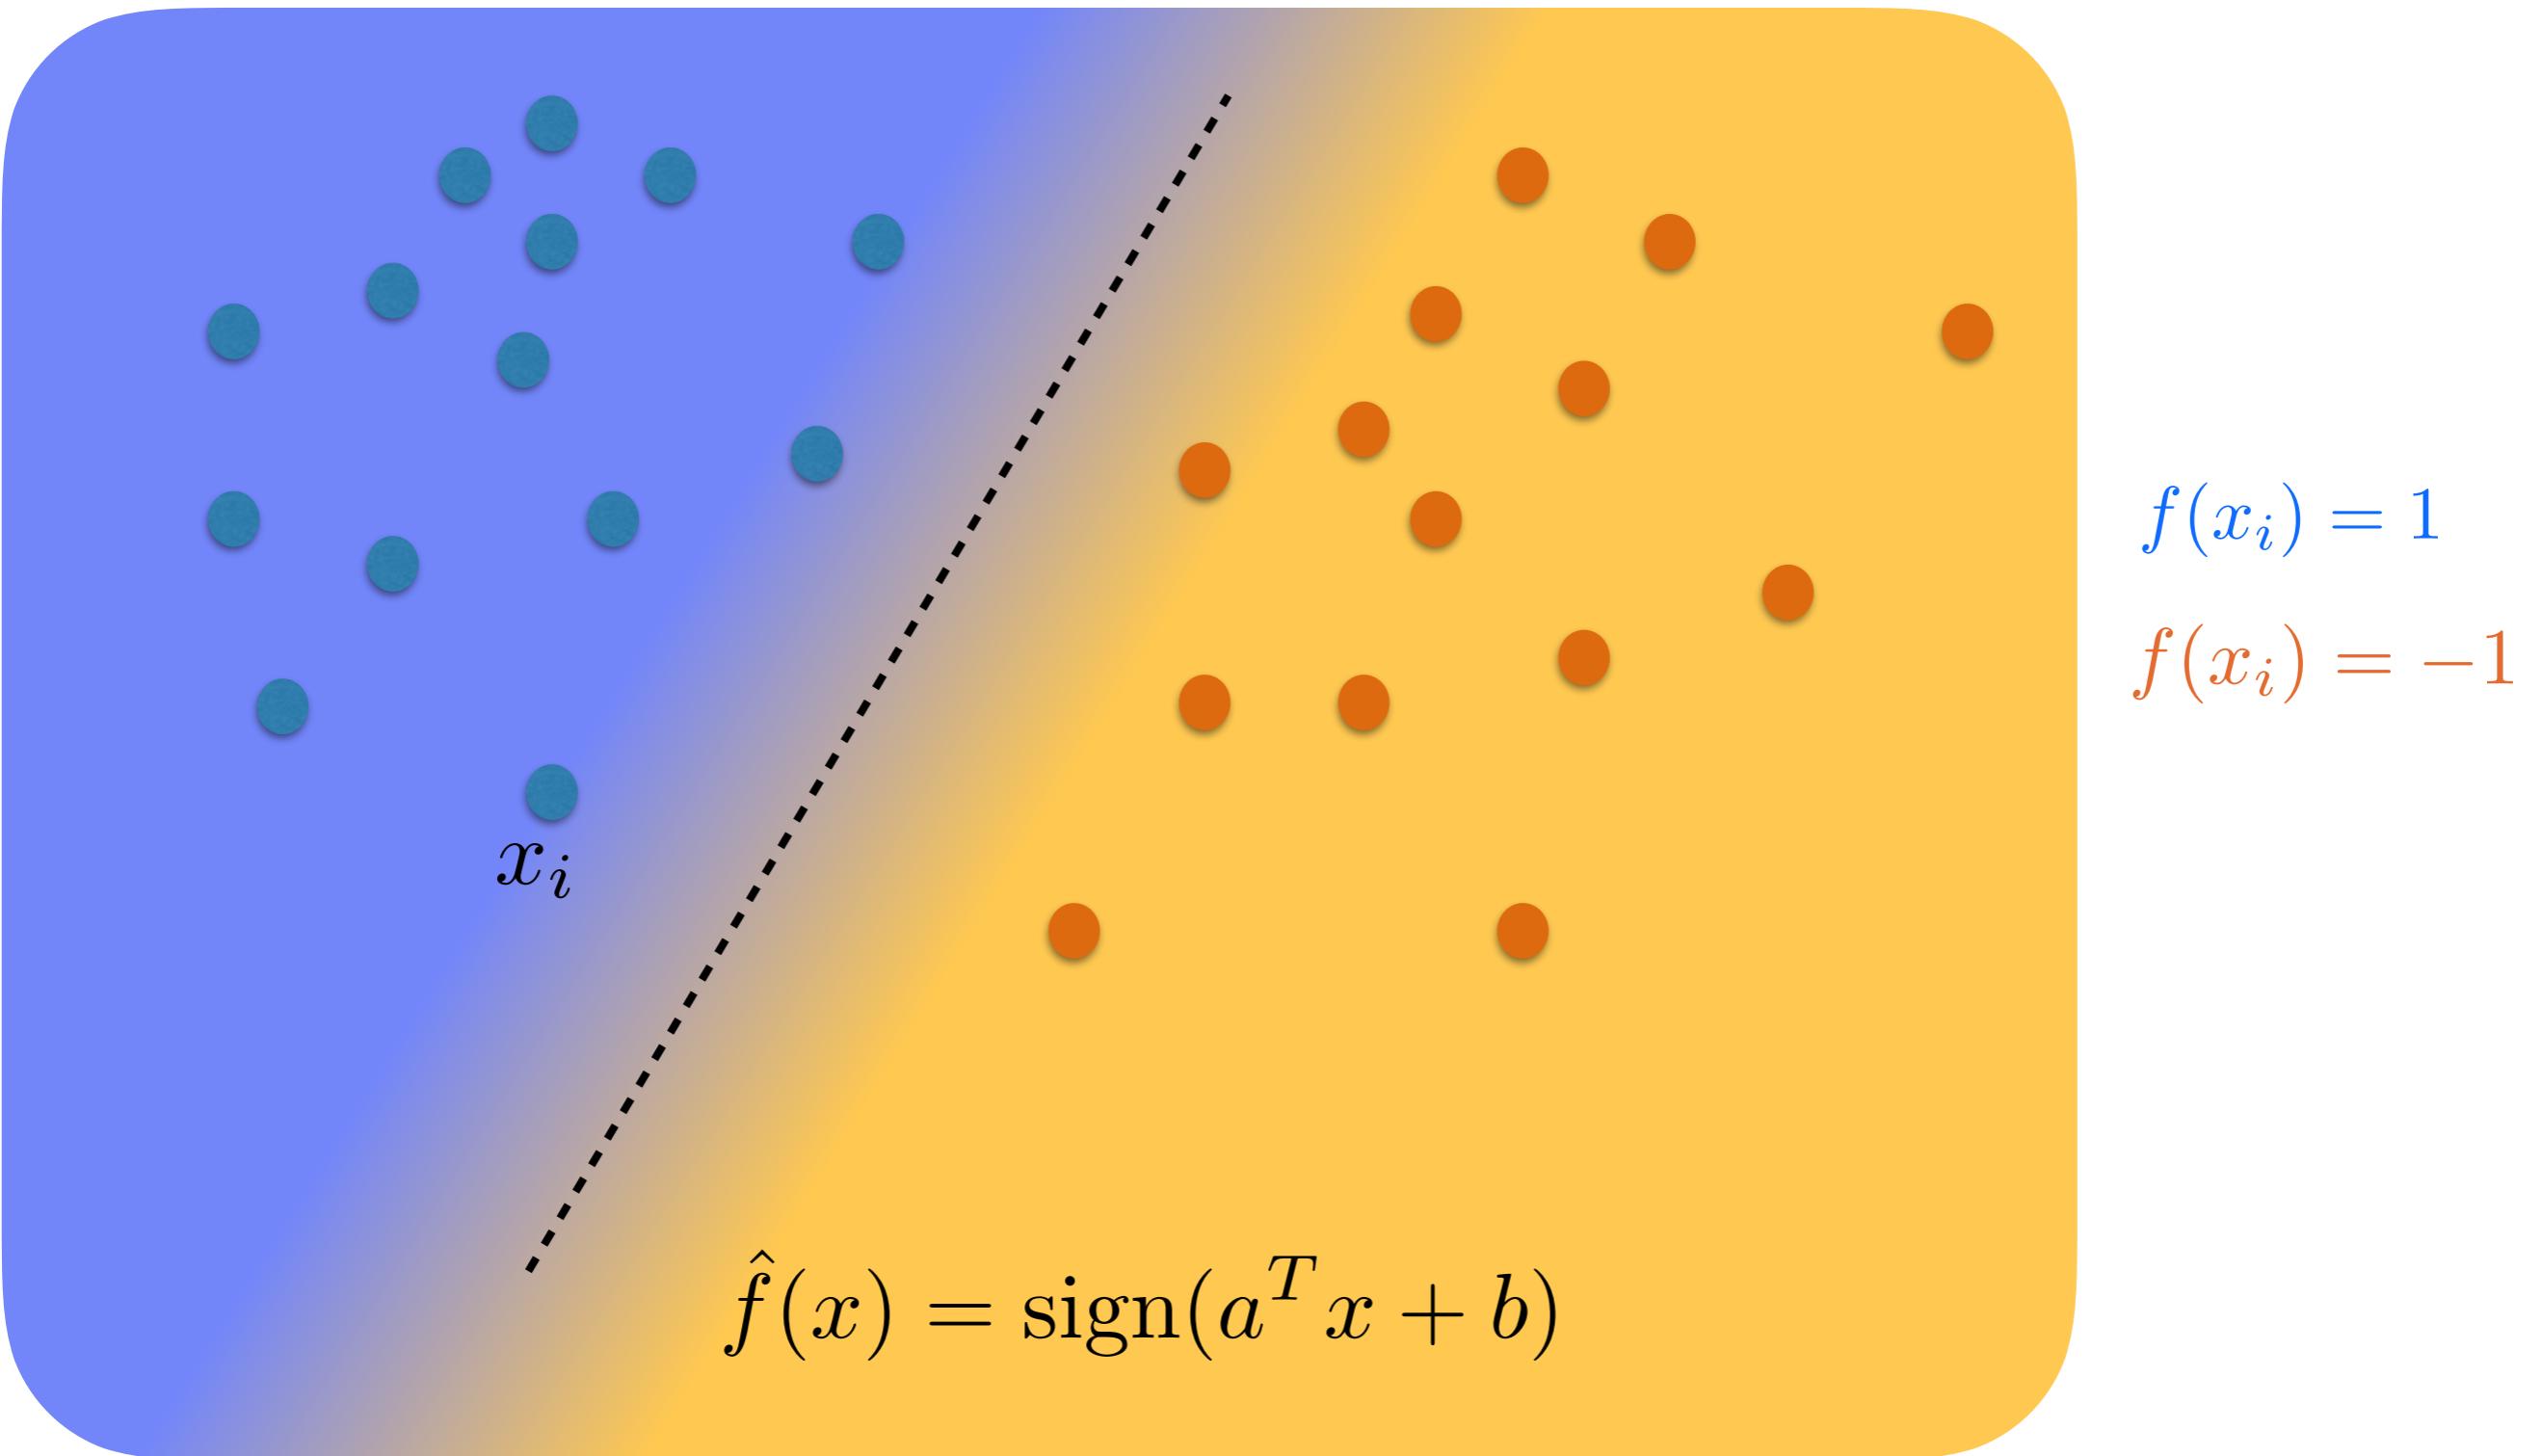
How to “interpolate” in high-dimensions?

- Let's start with a (very) simple low-dimensional setting:




How to “interpolate” in high-dimensions?

- Let's start with a (very) simple low-dimensional setting:



How to interpolate in high dimensions?




- We have found (linear) features $\Phi(x) = a^T x$ such that

$$|f(x) - f(x')| \leq C \|\Phi(x) - \Phi(x')\|$$

-
- The previous example corresponds to a binary classification problem that is **linearly separable**: there exists a hyperplane that separates the classes.

- The previous example corresponds to a binary classification problem that is **linearly separable**: there exists a hyperplane that separates the classes.
- By projecting $\Phi(x) = a^T x$ we transform the high-dimensional problem into a simple low-dimensional interpolation problem:



Support Vector Machines

Support Vector Machines

- The previous example is formalized by Support Vector Machines [Vapnik et al, '90s]: given a binary classification problem with data (x_i, y_i) , we consider an estimator for $f(x)$ of the form

$$\hat{f}(x) = \text{sign} (a^T x + b)$$

Support Vector Machines

- The previous example is formalized by Support Vector Machines [Vapnik et al, '90s]: given a binary classification problem with data (x_i, y_i) , we consider an estimator for $f(x)$ of the form

$$\hat{f}(x) = \text{sign} (a^T x + b)$$

- Empirical Risk Minimization:

$$\min_{a,b} \frac{1}{n} \sum_{i=1}^n \ell(y_i, \hat{f}(x_i)) + \lambda \|a\|^2 ,$$

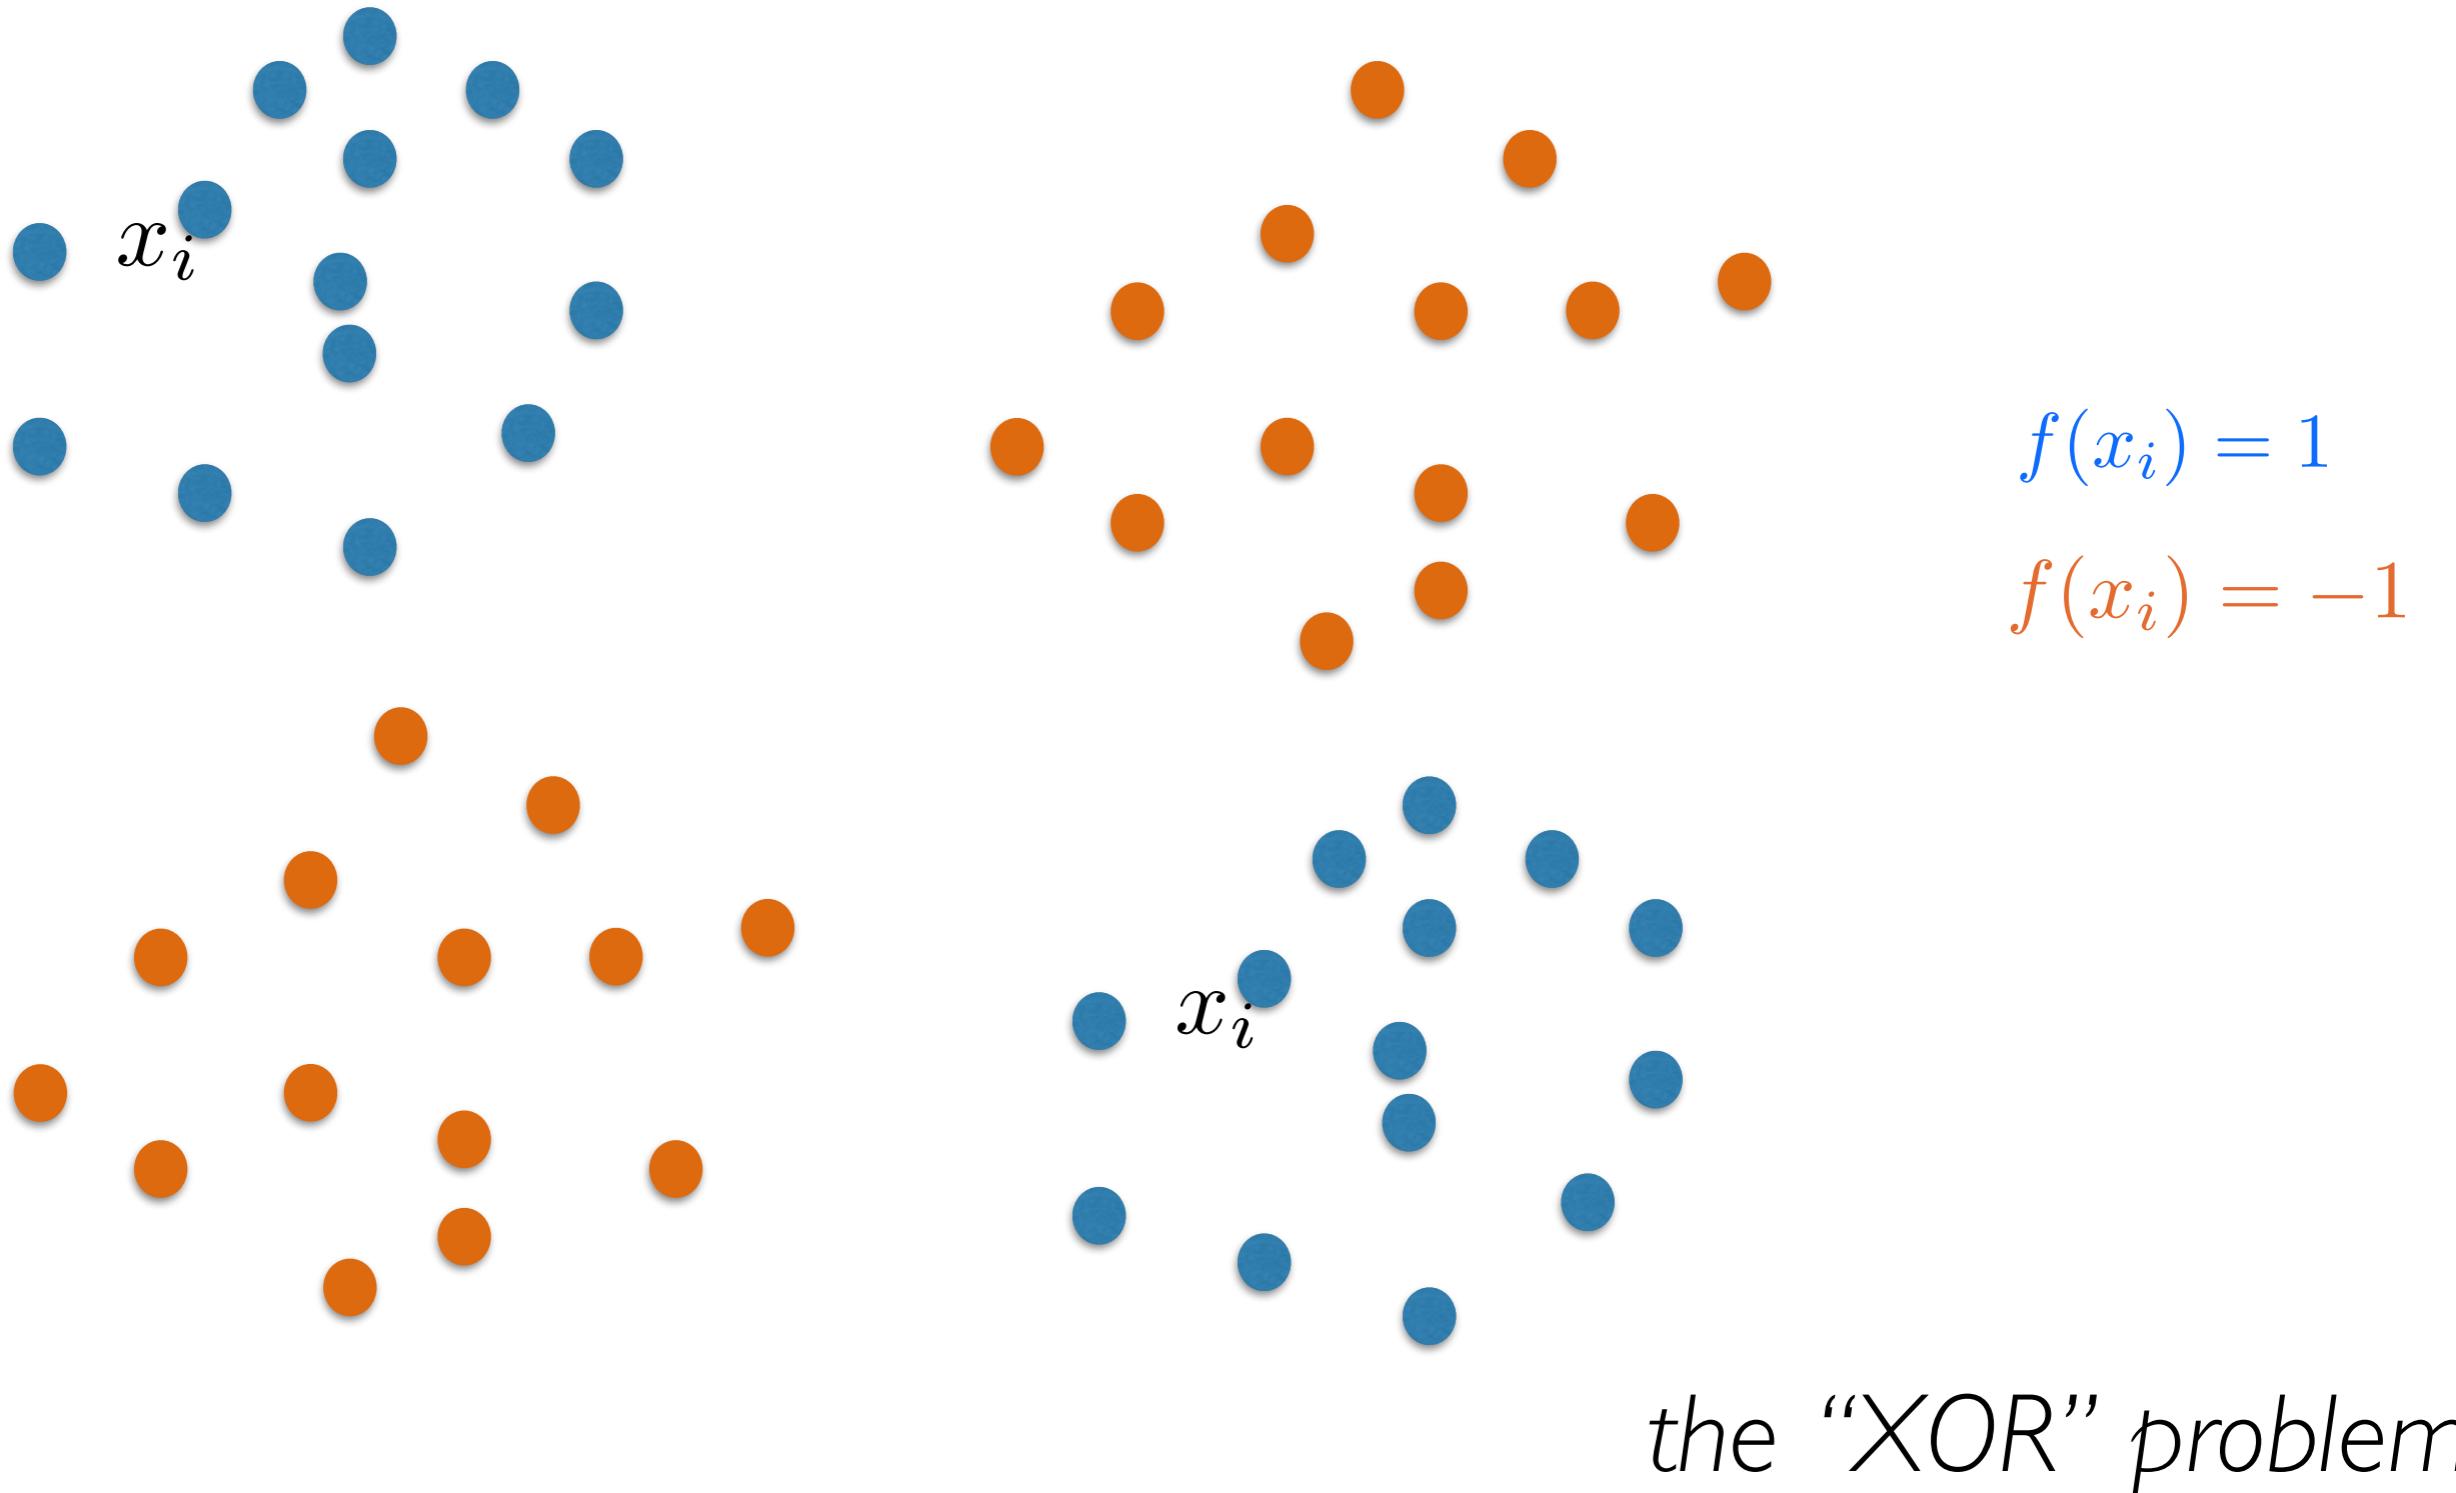
enforces training examples
to fall in the right side of the hyperplane

enforces large margin

$$\ell(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y}) \quad : \text{hinge loss} .$$

SVMs and Kernels

- Not all problems are linearly separable:



- By using the Lagrangian dual of the previous program, we can rewrite our previous solution as

$$\hat{f}(x) = \text{sign} \left(\sum_i \alpha_i y_i K(x_i, x) \right),$$

where $K(x_i, x) = \langle x_i, x \rangle$ is the Euclidean dot product.

- By using the Lagrangian dual of the previous program, we can rewrite our previous solution as

$$\hat{f}(x) = \text{sign} \left(\sum_i \alpha_i y_i K(x_i, x) \right),$$

where $K(x_i, x) = \langle x_i, x \rangle$ is the Euclidean dot product.

- We can replace the linear kernel by a non-linear one, eg
 - polynomial: $K(x, y) = \langle x, y \rangle^d$.
 - Gaussian radial basis function: $K(x, y) = \exp(-\|x - y\|^2/\sigma^2)$.

The Kernel “trick”

- For a wide class of psd kernels (Mercer Kernels), we have a representation in terms of an inner product:

$$\forall x, x' \in \Omega , \quad K(x, x') = \langle \Phi(x), \Phi(x') \rangle , \quad \Phi : \Omega \rightarrow \Omega'$$

The Kernel “trick”

- For a wide class of psd kernels (Mercer Kernels), we have a representation in terms of an inner product:

$$\forall x, x' \in \Omega , \quad K(x, x') = \langle \Phi(x), \Phi(x') \rangle , \quad \Phi : \Omega \rightarrow \Omega'$$

- It results that our estimate is **linear** in the features $\Phi(x)$:

$$\hat{f}(x) = \text{sign}(\langle w, \Phi(x) \rangle + b) , \quad w = \sum_i \alpha_i y_i \Phi(x_i).$$

The Kernel “trick”

- For a wide class of psd kernels (Mercer Kernels), we have a representation in terms of an inner product:

$$\forall x, x' \in \Omega , \quad K(x, x') = \langle \Phi(x), \Phi(x') \rangle , \quad \Phi : \Omega \rightarrow \Omega'$$

- It results that our estimate is **linear** in the features $\Phi(x)$

$$\hat{f}(x) = \text{sign}(\langle w, \Phi(x) \rangle + b) , \quad w = \sum_i \alpha_i y_i \Phi(x_i).$$

- Features need to be *discriminative*:

$$|f(x) - f(x')| \leq C \|\Phi(x) - \Phi(x')\|$$

The Kernel “trick”

- For a wide class of psd kernels (Mercer Kernels), we have a representation in terms of an inner product:

$$\forall x, x' \in \Omega , \quad K(x, x') = \langle \Phi(x), \Phi(x') \rangle , \quad \Phi : \Omega \rightarrow \Omega'$$

- It results that our estimate is **linear** in the features $\Phi(x)$:

$$\hat{f}(x) = \text{sign}(\langle w, \Phi(x) \rangle + b) , \quad w = \sum_i \alpha_i y_i \Phi(x_i).$$

- Features need to be *discriminative*:

$$|f(x) - f(x')| \leq C \|\Phi(x) - \Phi(x')\|$$

- Is this enough to characterize good features/kernels?

Generalization Error

- It is easy to construct discriminative features:
 - Using a Gaussian RBF, it suffices to let $\sigma^2 \rightarrow 0$.
 - The estimator converges to the *nearest neighbor* classifier:

$$\hat{f}(x) = f(x_{i(x)}) , \quad i(x) = \arg \min_i \|x - x_i\|$$

Generalization Error

- It is easy to construct discriminative features:
 - Using a Gaussian RBF, it suffices to let $\sigma^2 \rightarrow 0$.
 - The estimator converges to the *nearest neighbor* classifier:

$$\hat{f}(x) = f(x_{i(x)}) , \quad i(x) = \arg \min_i \|x - x_i\|$$

- While it may be easy to correctly classify our *training* examples, we do not necessarily improve our generalization error:

$$\mathbf{E}_{(x,y)}(\ell(\hat{f}(x), y))$$

Generalization Error

- It is easy to construct discriminative features:
 - Using a Gaussian RBF, it suffices to let $\sigma^2 \rightarrow 0$.
 - The estimator converges to the *nearest neighbor* classifier:

$$\hat{f}(x) = f(x_{i(x)}) , \quad i(x) = \arg \min_i \|x - x_i\|$$

- While it may be easy to correctly classify our *training* examples, we do not necessarily improve our generalization error:

$$\mathbf{E}_{(x,y)}(\ell(\hat{f}(x), y))$$

- The larger the embedding dimension, the higher is the risk of overfitting.

Generalization Error

- It is easy to construct discriminative features:
 - Using a Gaussian RBF, it suffices to let $\sigma^2 \rightarrow 0$.
 - The estimator converges to the *nearest neighbor* classifier:

$$\hat{f}(x) = f(x_{i(x)}) , \quad i(x) = \arg \min_i \|x - x_i\|$$

- While it may be easy to correctly classify our *training* examples, we do not necessarily improve our generalization error:

$$\mathbf{E}_{(x,y)}(\ell(\hat{f}(x), y))$$

- The larger the embedding dimension, the higher is the risk of overfitting.
- Underlying question: how to compare signals in high-dim?

Curse of Dimensionality

- In a finite-dimensional, bounded space, all metrics are equivalent:

for each $x \in \Omega$, exists constants c, C such that
 $\forall x' \in \Omega, cd(x, x') \leq \tilde{d}(x, x') \leq Cd(x, x')$.

Curse of Dimensionality

- In a finite-dimensional, bounded space, all metrics are equivalent:

for each $x \in \Omega$, exists constants c, C such that
 $\forall x' \in \Omega, cd(x, x') \leq \tilde{d}(x, x') \leq Cd(x, x')$.

- But as the dimension increases, metrics start to “diverge”.
 - In particular, the Euclidean distance in high-dimensional spaces is typically a poor measure of similarity for practical purposes.

Curse of Dimensionality

- In a finite-dimensional, bounded space, all metrics are equivalent:

for each $x \in \Omega$, exists constants c, C such that
 $\forall x' \in \Omega, cd(x, x') \leq \tilde{d}(x, x') \leq Cd(x, x')$.

- But as the dimension increases, metrics start to “diverge”.
 - In particular, the Euclidean distance in high-dimensional spaces is typically a poor measure of similarity for practical purposes.
- Local decisions around training do not extend to the whole space.
- So, we need a guiding principle that plays well with our data (images, sounds, etc.)

2-dimensional
embedding of
CIFAR-10 using
Euclidean similarity

from A. Karpathy



Linearization

- We want to obtain a representation $\Phi(x)$ such that

$$\hat{f}(x) = \text{sign}(a^T \Phi(x) + b)$$

is a good approximation of $f(\mathbf{x})$.

Linearization

- We want to obtain a representation $\Phi(x)$ such that

$$\hat{f}(x) = \text{sign}(a^T \Phi(x) + b)$$

is a good approximation of $f(\mathbf{x})$. Thus $f(\mathbf{x})$ is approximately linearized by $\Phi(x)$:

$$f(x) \approx \text{sign}(a^T \Phi(x) + b)$$

Linearization

- We want to obtain a representation $\Phi(x)$ such that

$$\hat{f}(x) = \text{sign}(a^T \Phi(x) + b)$$

is a good approximation of $f(x)$. Thus $f(x)$ is approximately linearized by $\Phi(x)$:

$$f(x) \approx \text{sign}(a^T \Phi(x) + b)$$

- In particular, we should have

$$a^T(\Phi(x) - \Phi(x')) = 0 \implies f(x) = f(x') .$$

Linearization

- We want to obtain a representation $\Phi(x)$ such that

$$\hat{f}(x) = \text{sign}(a^T \Phi(x) + b)$$

is a good approximation of $f(\mathbf{x})$. Thus $f(\mathbf{x})$ is approximately linearized by $\Phi(x)$:


$$f(x) \approx \text{sign}(a^T \Phi(x) + b)$$

- In particular, we should have

$$a^T(\Phi(x) - \Phi(x')) = 0 \implies f(x) = f(x') .$$

Thus the level sets of f should be mapped to parallel hyperplanes by Φ

Linearization



In order to beat the curse of dimensionality, we need features that **linearize intra-class variability** and **preserve inter-class variability**.

Invariance and Symmetry

- A global symmetry is an operator $\varphi \in Aut(\Omega)$ that leaves f invariant:

$$\forall x \in \Omega , f(\varphi(x)) = f(x) .$$

Invariance and Symmetry

- A global symmetry is an operator $\varphi \in Aut(\Omega)$ that leaves f invariant:

$$\forall x \in \Omega , f(\varphi(x)) = f(x) .$$

- They can be absorbed by Φ to varying degrees:

Invariants: $\Phi(\varphi(x)) = \Phi(x)$ for each x .

Covariants: $\Phi(\varphi(x)) = A_\varphi \Phi(x)$ for each x ,
where A_φ is “simpler” than φ

Invariance and Symmetry

- A global symmetry is an operator $\varphi \in Aut(\Omega)$ that leaves f invariant:

$$\forall x \in \Omega , f(\varphi(x)) = f(x) .$$

- They can be absorbed by Φ to varying degrees:

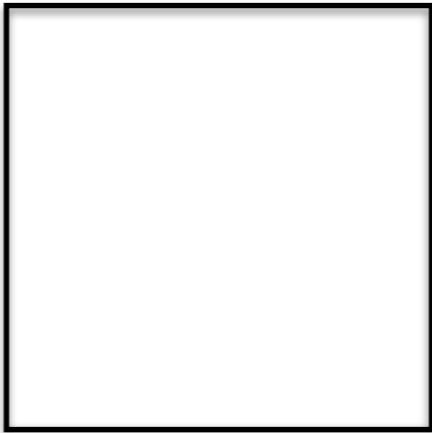
Invariants: $\Phi(\varphi(x)) = \Phi(x)$ for each x .

Covariants: $\Phi(\varphi(x)) = A_\varphi \Phi(x)$ for each x ,
where A_φ is “simpler” than φ

- What are those symmetries? How to impose them on Φ without breaking discriminability?

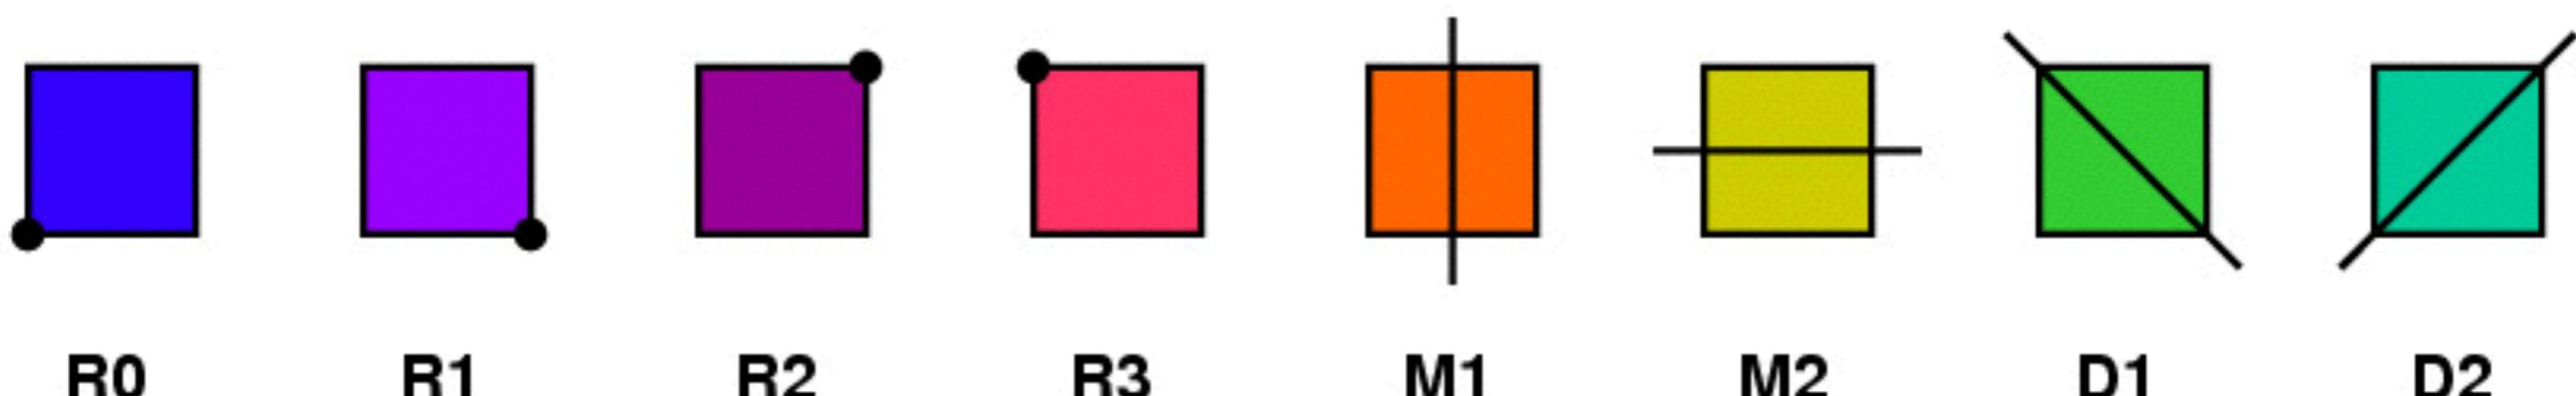

Discrete symmetries

- Which transformations leave this square unchanged?



Discrete symmetries

- Which transformations leave this square unchanged?

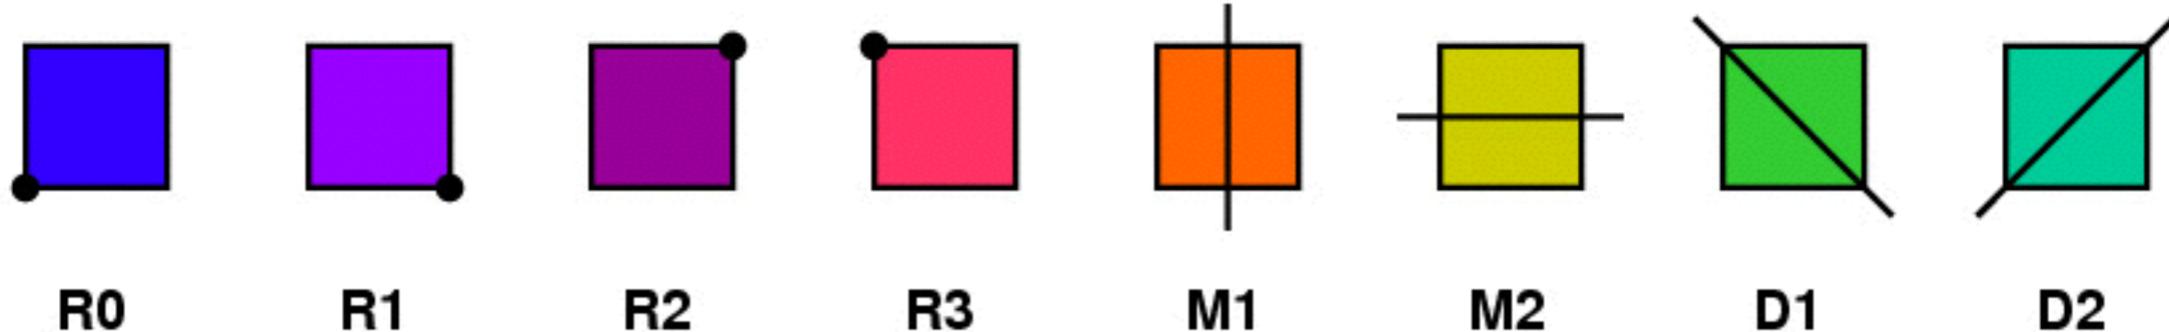


- They form a group

(from <http://www.cs.umb.edu/~eb/>)

Discrete symmetries


- Which transformations leave this square unchanged?



- The set of all symmetries forms a group G :
 - group operation: $\forall g_1, g_2 \in G, g_1 \cdot g_2 \in G$.
 - identity element: $\exists e \in G$ s.t. $g \cdot e = e \cdot g = g \quad \forall g \in G$.
 - inverse: $\forall g \in G \exists g^{-1} \in G$ s.t. $g \cdot g^{-1} = e$.

Discrete symmetries

- Which transformations leave this square unchanged?



- Discrete groups are completely characterized by their multiplication table:

	R0	R1	R2	R3	M1	M2	D1	D2
R0	[Blue]	[Purple]	[Magenta]	[Red]	[Orange]	[Yellow]	[Green]	[Cyan]
R1	[Purple]	[Magenta]	[Red]	[Blue]	[Green]	[Yellow]	[Orange]	[Cyan]
R2	[Magenta]	[Red]	[Blue]	[Purple]	[Yellow]	[Orange]	[Cyan]	[Green]
R3	[Red]	[Blue]	[Purple]	[Magenta]	[Cyan]	[Green]	[Orange]	[Yellow]
M1	[Orange]	[Orange]	[Green]	[Yellow]	[Blue]	[Purple]	[Red]	[Cyan]
M2	[Yellow]	[Yellow]	[Green]	[Orange]	[Purple]	[Blue]	[Cyan]	[Red]
D1	[Green]	[Green]	[Orange]	[Yellow]	[Purple]	[Blue]	[Red]	[Cyan]
D2	[Cyan]	[Cyan]	[Yellow]	[Green]	[Orange]	[Purple]	[Red]	[Blue]

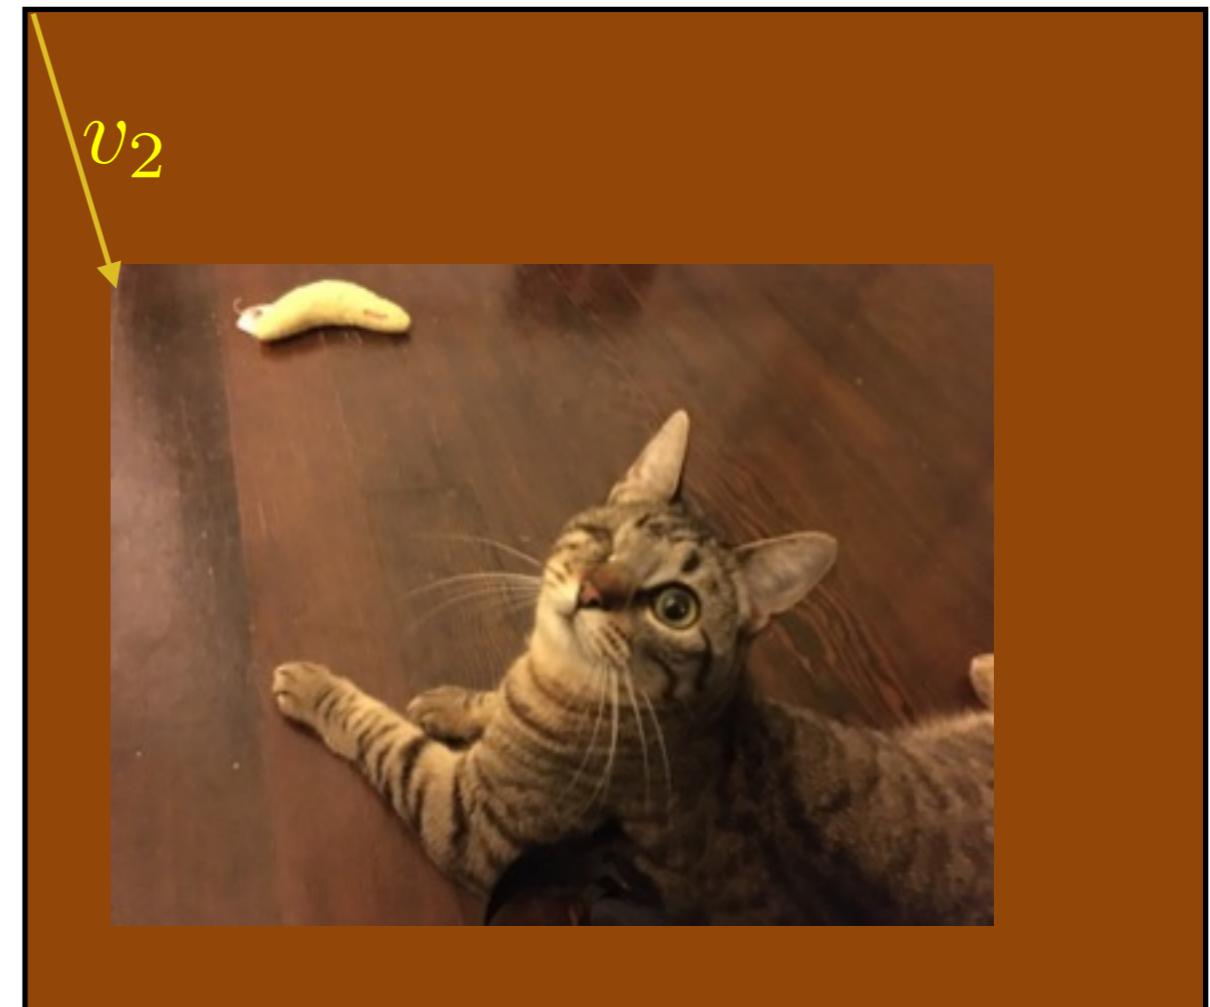

(from <http://www.cs.umb.edu/~eb/>)

Rigid transformation symmetries

- Which symmetries are we likely to find in image recognition problems?

Rigid transformation symmetries

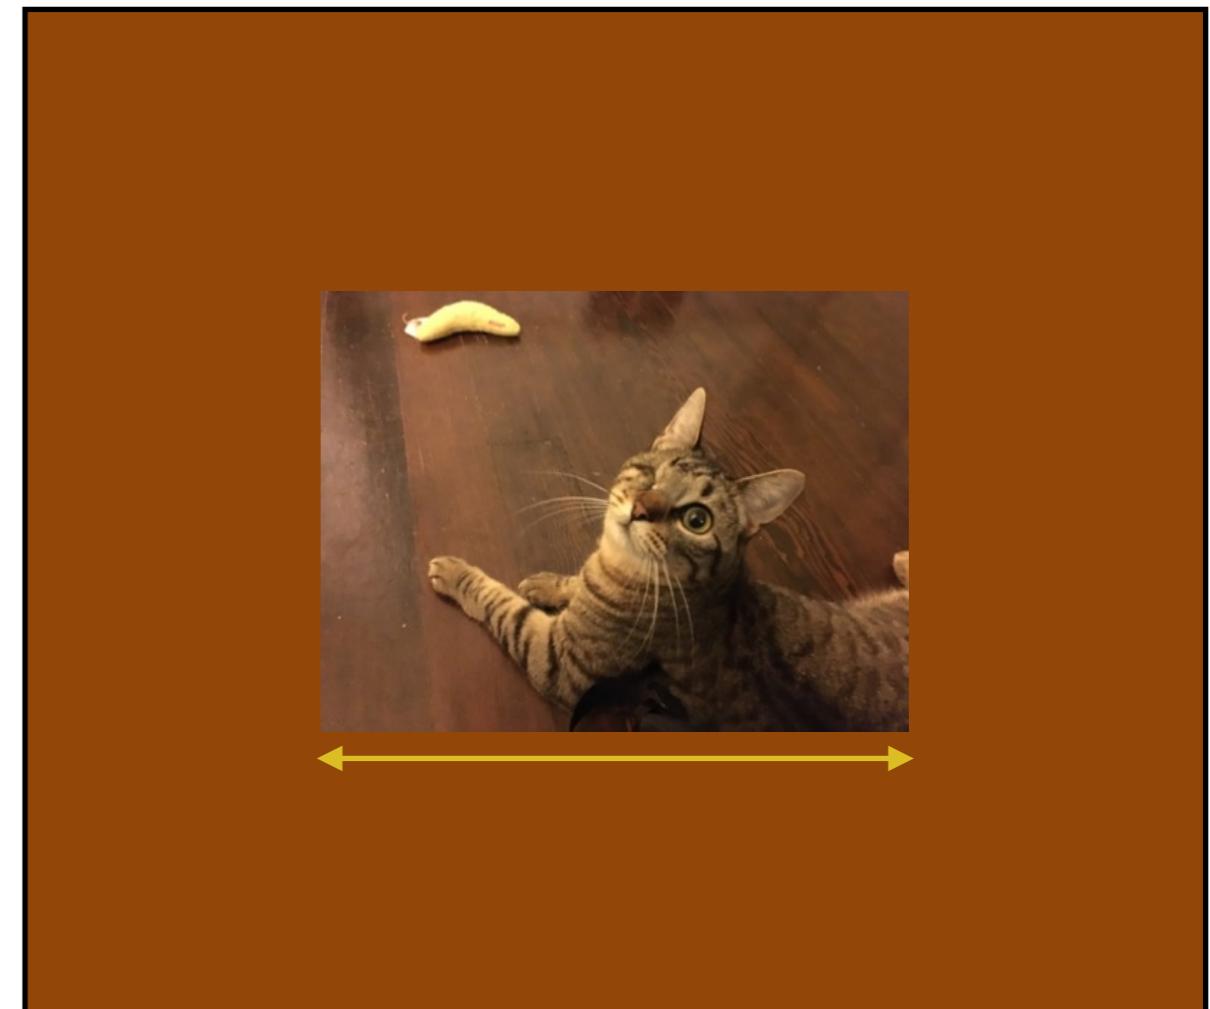
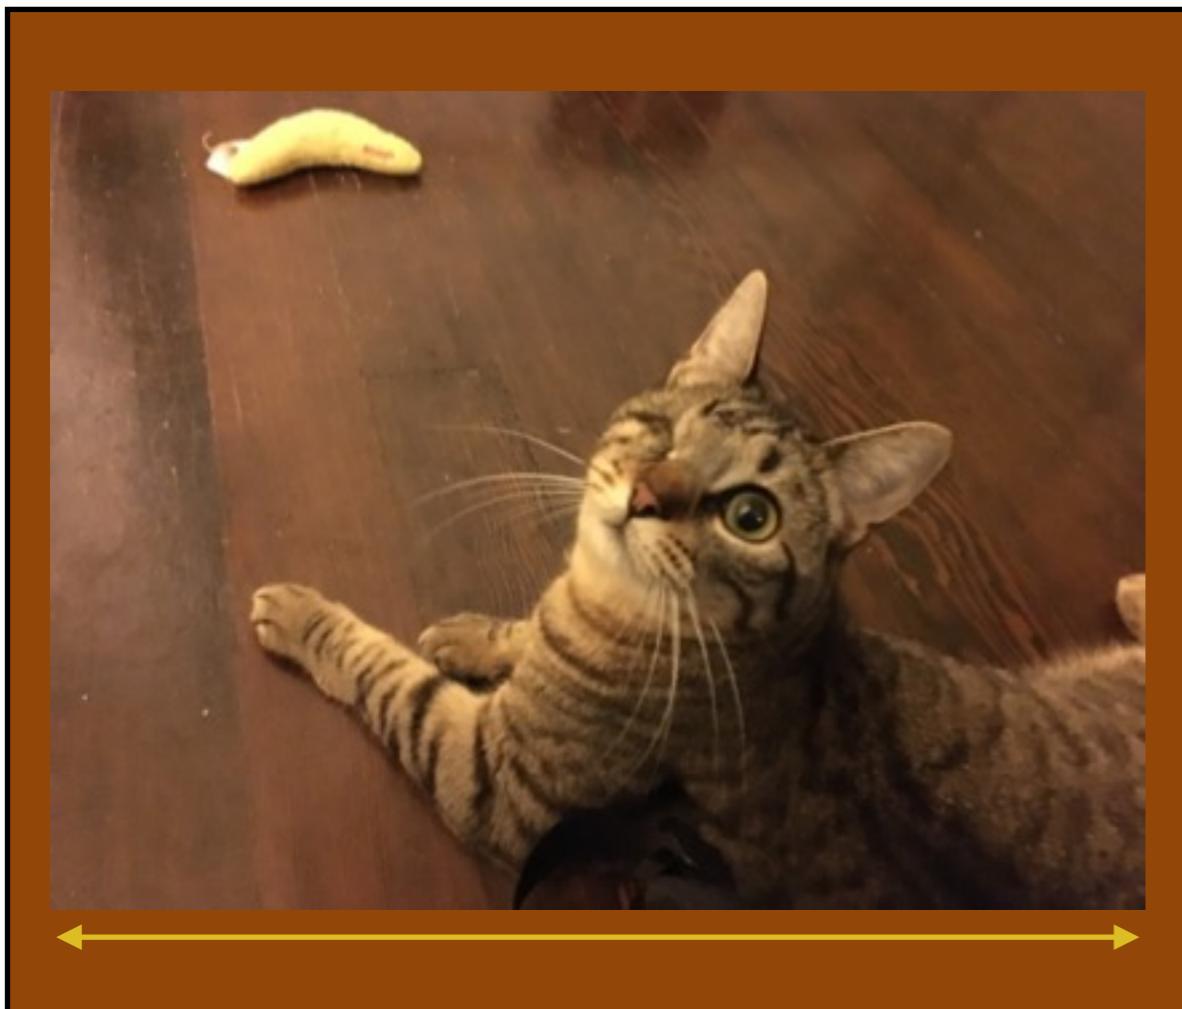
- Which symmetries are we likely to find in image recognition problems?



Translations: $\{\varphi_v ; v \in \mathbb{R}^2\}$, with $\varphi_v(x)(u) = x(u - v)$.

Rigid transformation symmetries


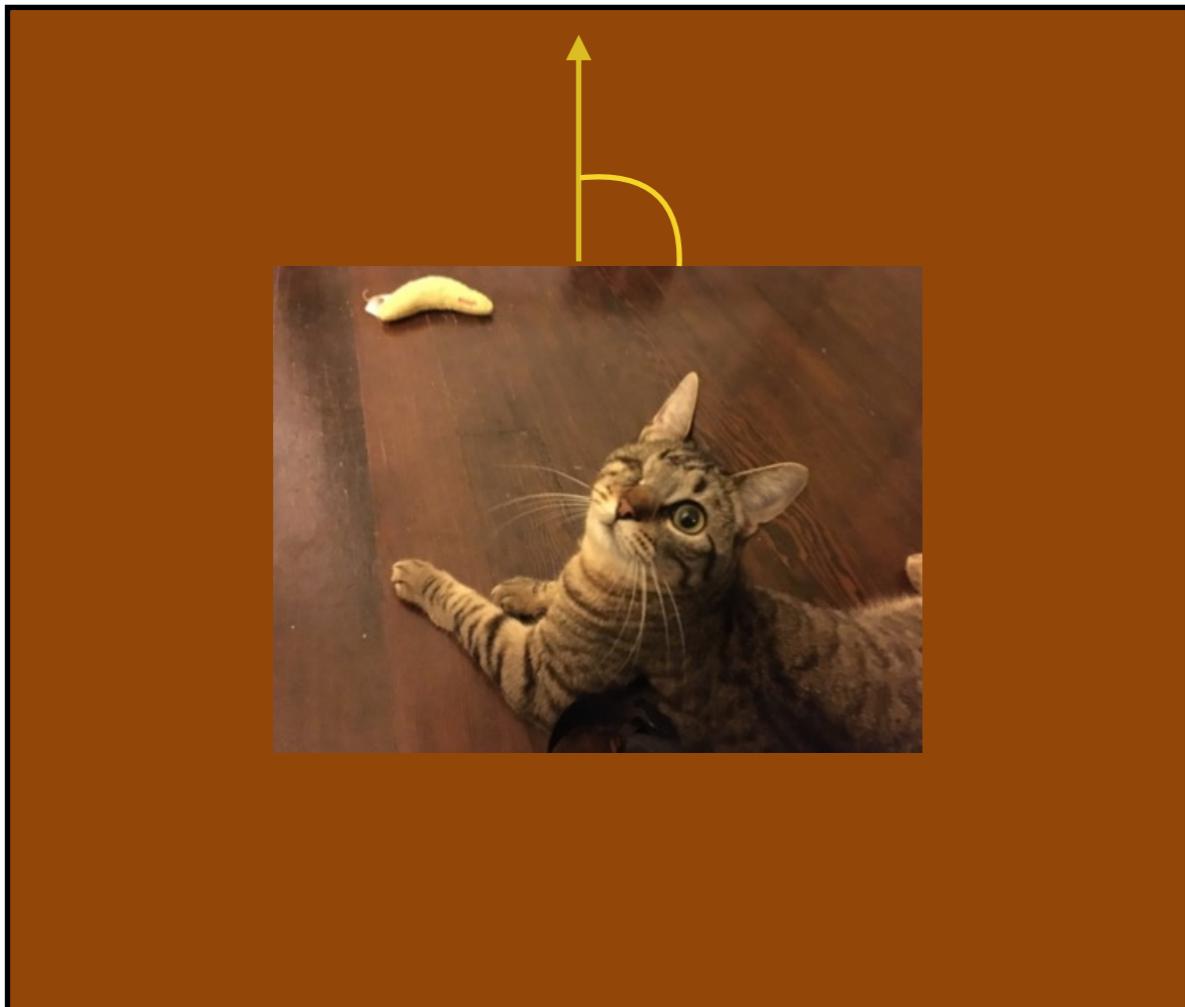
- Which symmetries are we likely to find in image recognition problems?



Dilations: $\{\varphi_s ; s \in \mathbb{R}_+\}$, with $\varphi_s(x)(u) = s^{-1}x(s^{-1}u)$.

Rigid transformation symmetries


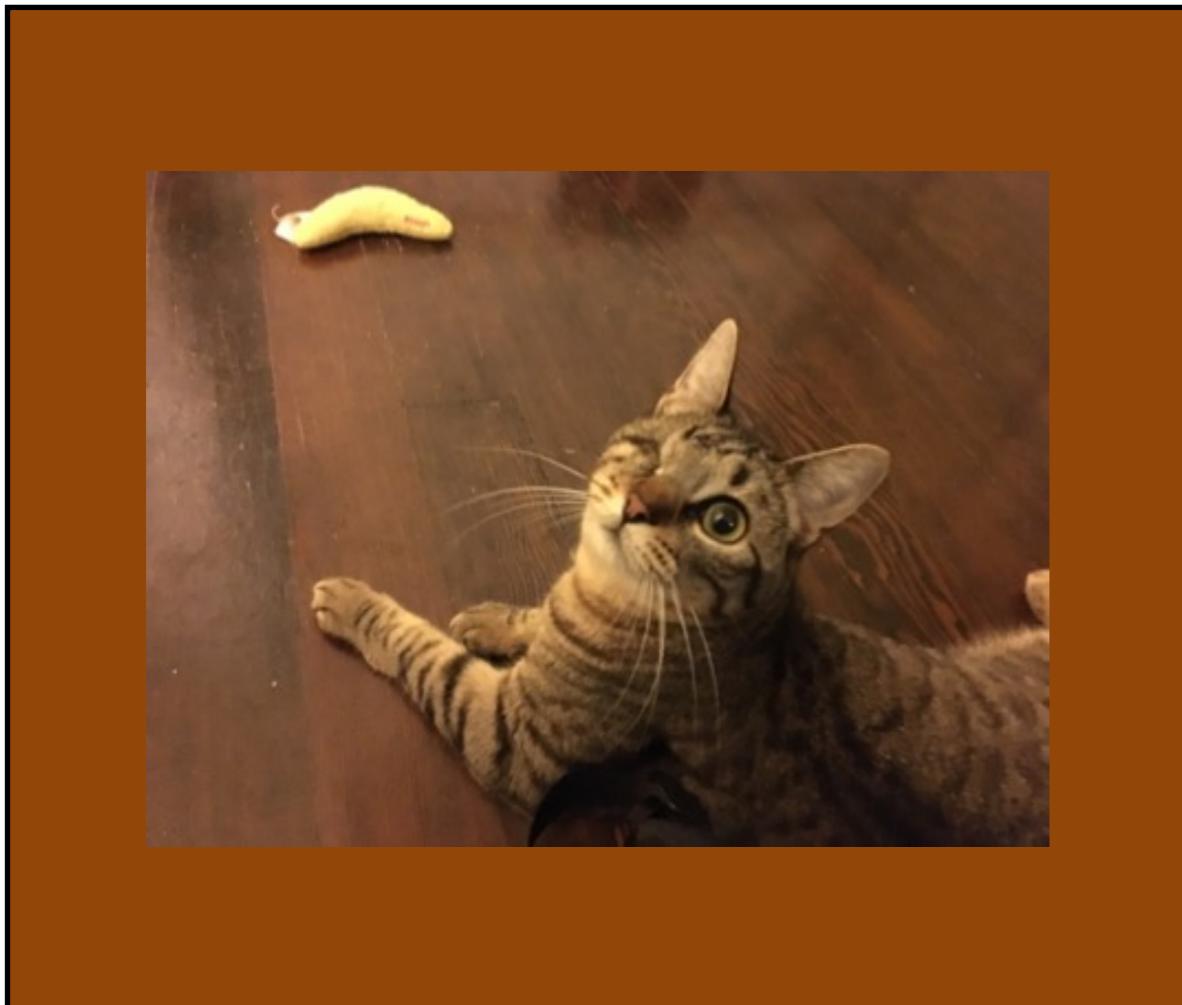
- Which symmetries are we likely to find in image recognition problems?



Rotations: $\{\varphi_\theta ; \theta \in [0, 2\pi)\}$, with $\varphi_\theta(x)(u) = x(R_\theta u)$.

Rigid transformation symmetries

- Which symmetries are we likely to find in image recognition problems?



Mirror symmetry: $\{e, M\}$, with $Mx(u_1, u_2) = x(-u_1, u_2)$.

Rigid transformation symmetries

- We can combine all these transformations into a single group, the Affine Group $\text{Aff}(\mathbb{R}^2)$.
- It has 6 degrees of freedom; in the representation

$$\begin{pmatrix} u_1 \\ u_2 \end{pmatrix} \mapsto \begin{pmatrix} v_1 \\ v_2 \end{pmatrix} + \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \end{pmatrix}$$

$$g = (v_1, v_2, a_1, a_2, a_3, a_4)$$

Rigid transformation symmetries

- We can combine all these transformations into a single group, the Affine Group $\text{Aff}(\mathbb{R}^2)$.

- It has 6 degrees of freedom; in the representation

$$\begin{pmatrix} u_1 \\ u_2 \end{pmatrix} \mapsto \begin{pmatrix} v_1 \\ v_2 \end{pmatrix} + \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \end{pmatrix}$$

$$g = (v_1, v_2, a_1, a_2, a_3, a_4)$$

- Note that this is in general a *non-commutative* group.
- For some groups, we might only observe partial invariance (e.g. rotation and dilation).

Rigid transformation symmetries

- We can combine all these transformations into a single group, the Affine Group $\text{Aff}(\mathbb{R}^2)$.

- It has 6 degrees of freedom; in the representation

$$\begin{pmatrix} u_1 \\ u_2 \end{pmatrix} \mapsto \begin{pmatrix} v_1 \\ v_2 \end{pmatrix} + \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \end{pmatrix}$$


$$g = (v_1, v_2, a_1, a_2, a_3, a_4)$$

- Note that this is in general a non-commutative group.
- For some groups, we might only observe partial invariance (e.g. rotation and dilation).
- In speech, the underlying group modeling time-frequency shifts is the Heisenberg group.

Invariant Representations

- Given a transformation group G and an input x , the *action* of G onto x is called an *orbit*:

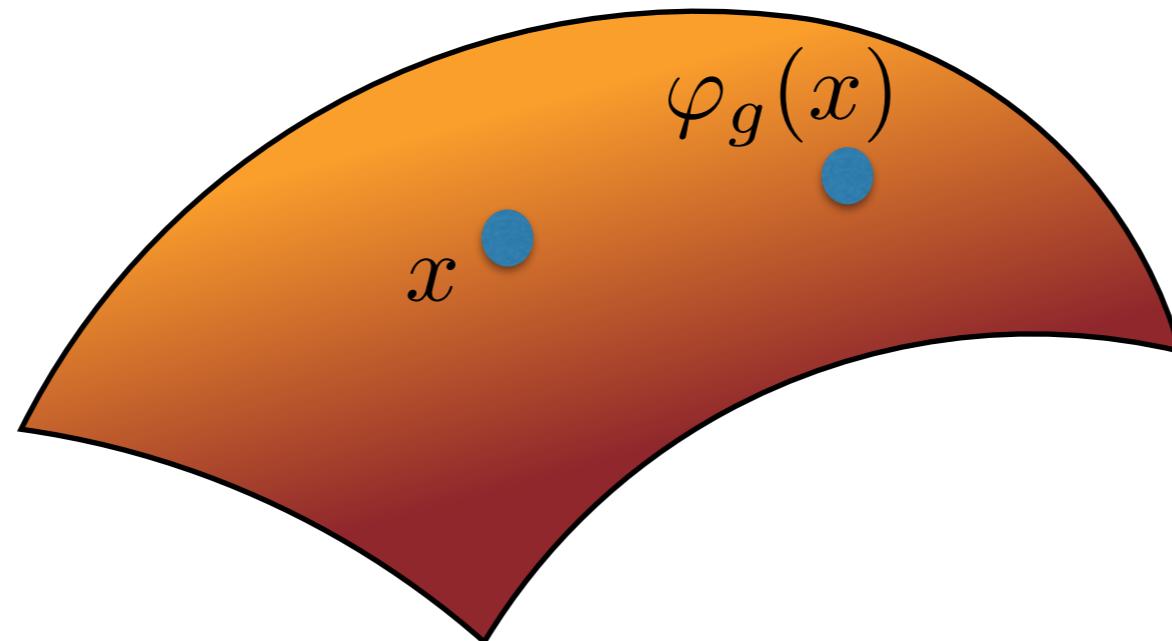
$$G \cdot x = \{\varphi_g(x), g \in G\}$$



Invariant Representations

- Given a transformation group G and an input x , the *action* of G onto x is called an *orbit*:

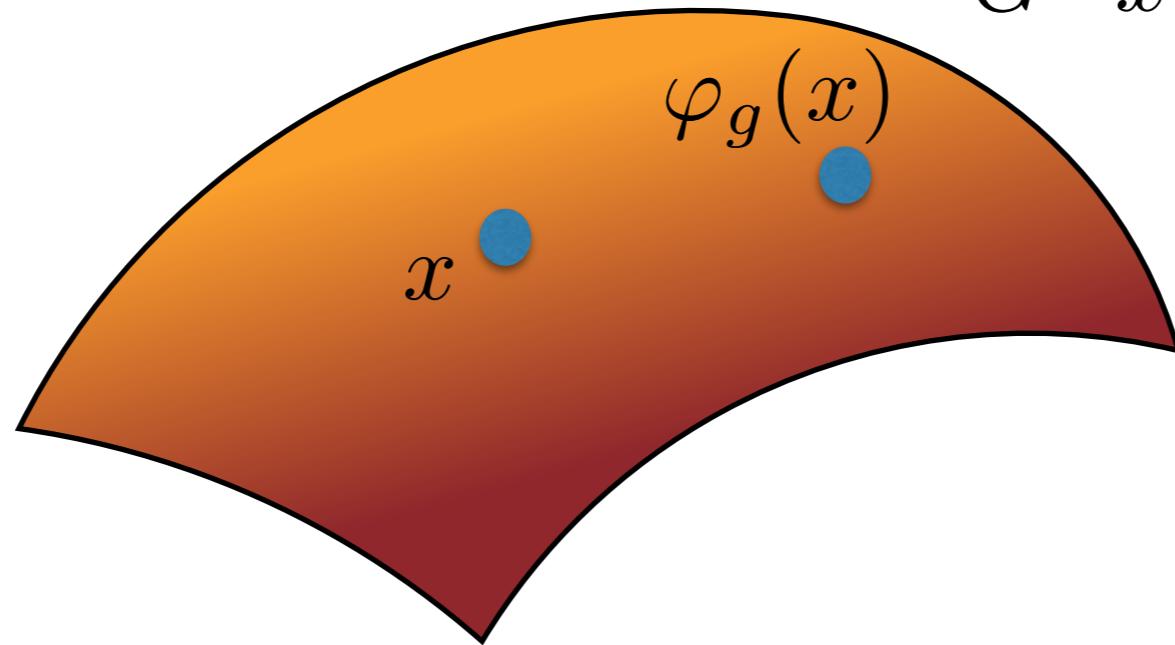
$$G \cdot x = \{\varphi_g(x), g \in G\}$$



- Impact on the learning task?
- Since our estimator is linear in $\Phi(x)$, $\Phi(G \cdot x)$ should be “flat”.

Invariant Representations


$$G \cdot x = \{\varphi_g(x), g \in G\}$$



- Problem?

Invariant Representations

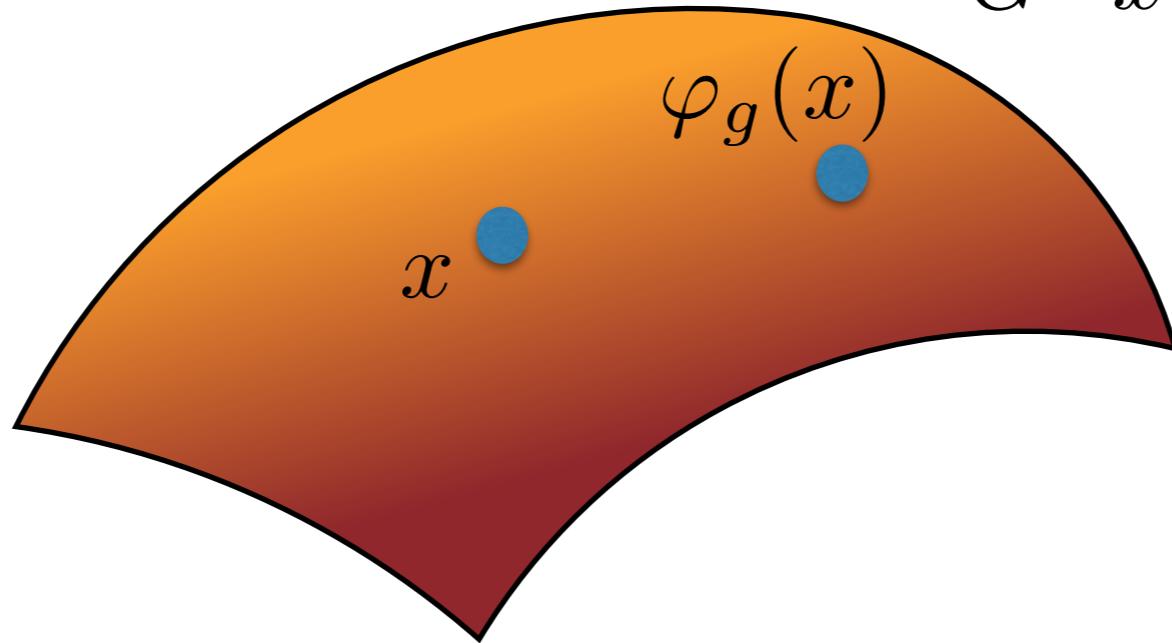
$$G \cdot x = \{\varphi_g(x) , g \in G\}$$



- Problem? A 6-dimensional curvy space looks flat in a high-dimensional space.

Invariant Representations

$$G \cdot x = \{\varphi_g(x), g \in G\}$$



- Problem? A 6-dimensional curvy space looks flat in a high-dimensional space.
- Group symmetries are not sufficient to beat the curse of dimensionality.

From Invariance to Stability

- Symmetry is a very strict criteria. Can we relax it?

From Invariance to Stability

- Symmetry is a very strict criteria. Can we relax it?
- Although image and audio recognition does not have high-dimensional symmetry groups, it is *stable* to local deformations.

$$x \in L^2(\mathbb{R}^m) , \tau : \mathbb{R}^m \rightarrow \mathbb{R}^m \text{ diffeomorphism}$$

$$x_\tau = \varphi_\tau(x) , x_\tau(u) = x(u - \tau(u))$$

φ_τ is a change of variables: (think of x_τ as adding noise to the pixel *locations* rather than to the pixel values)

From Invariance to Stability



- Informally, if $\|\tau\|$ measures the amount of deformation, many recognition tasks satisfy

$$\forall x, \tau, |f(x) - f(x_\tau)| \lesssim \|\tau\|$$

From Invariance to Stability



- Informally, if $\|\tau\|$ measures the amount of deformation, many recognition tasks satisfy

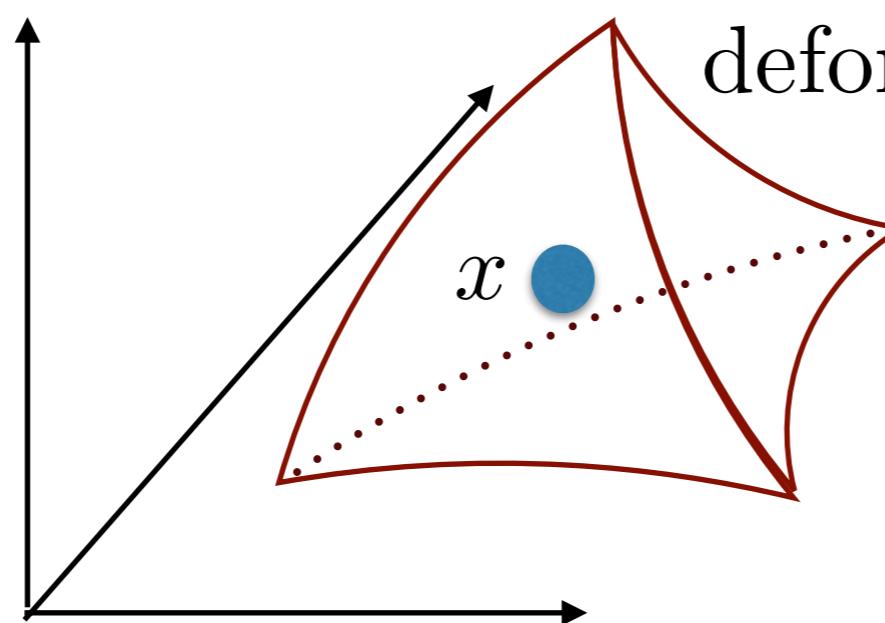
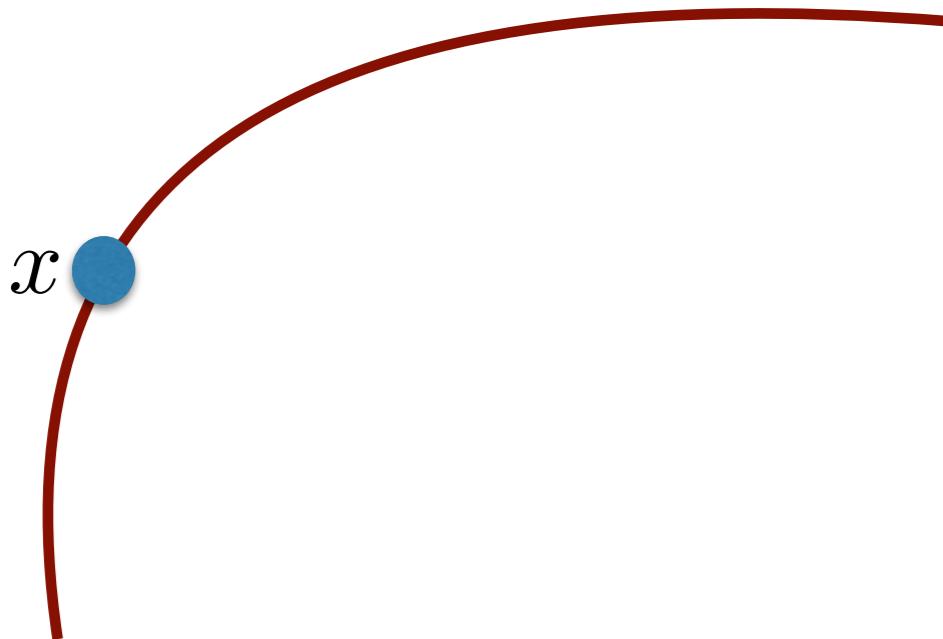
$$\forall x, \tau, |f(x) - f(x_\tau)| \lesssim \|\tau\|$$

- If our representation is stable, then

$$\forall x, \tau, \|\Phi(x) - \Phi(x_\tau)\| \leq C\|\tau\| \implies |\hat{f}(x) - \hat{f}(x_\tau)| \leq \tilde{C}\|\tau\|$$

Filling the space with deformations

symmetry group: low dimension




deformations fill the space

Deformations in Image/Audio Recognition

- Can model 3D viewpoint changes, changes in pitch/timbre in speech recognition.

Deformations in Image/Audio Recognition

- Can model 3D viewpoint changes, changes in pitch/timbre in speech recognition.
- Deformable parts model [Feltzenszwab et al, '10]



- State-of-the-art on object detection pre-CNN.

Deformations in Image/Audio Recognition

- Can model 3D viewpoint changes, changes in pitch/timbre in speech recognition.
- Deformable templates [Grenader, Younes, Trouve, Amit et al.]
 - Equip deformable templates with differentiable structure

Deformations in Image/Audio Recognition

- Can model 3D viewpoint changes, changes in pitch/timbre in speech recognition.
- Deformable templates [Grenader, Younes, Trouve, Amit et al.]
 - Equip deformable templates with differentiable structure
- Data augmentation in Object classification
 - Mostly rigid transformations (random shifts, flips).

Stability Condition

- We introduced the stability condition

$$\forall x, \tau, \|\Phi(x) - \Phi(x_\tau)\| \lesssim \|\tau\| .$$

Stability Condition

- We introduced the stability condition

$$\forall x, \tau, \|\Phi(x) - \Phi(x_\tau)\| \lesssim \|\tau\| .$$

- If we fix the ‘template’ x and consider the mapping

$$F : \tau \mapsto \Phi(x_\tau)$$

the previous condition becomes

$$\|F(\tau) - F(0)\| \leq C\|\tau\| ,$$

thus F is Lipschitz with respect to the deformation metric
 $\|\tau\|$ uniformly on x .

Stationarity Prior

- Two clips. Goal: distinguish which is which.

clip1

clip2

clip ?

Stationarity Prior

- Same experiment. Goal: distinguish which is which.

clip3

clip4

clip ?

Stationarity Prior

- Same experiment. Goal: distinguish which is which.

clip3

clip4

clip ?

- Typically, the latter is harder. Reasons?

Stationarity Prior

- Same experiment. Goal: distinguish which is which.

clip3

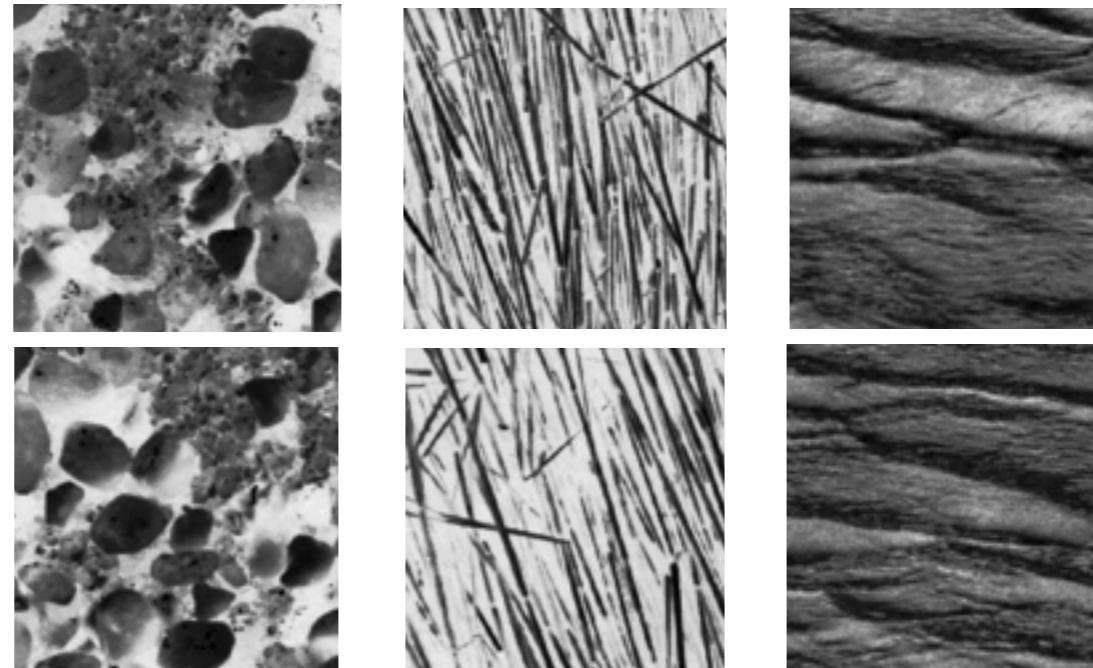
clip4

clip ?

- Typically, the latter is harder. Reasons?
- Despite having more information, the discrimination is worse because we construct temporal averages in presence of *stationary* inputs.


Representation of Stationary Processes

$x(u)$: realizations of a stationary process $X(u)$ (not Gaussian)



Representation of Stationary Processes

$x(u)$: realizations of a stationary process $X(u)$ (not Gaussian)



$$\Phi(X) = \{E(f_i(X))\}_i$$

Estimation from samples $x(n)$: $\widehat{\Phi}(X) = \left\{ \frac{1}{N} \sum_n f_i(x)(n) \right\}_i$

Discriminability: need to capture high-order moments
Stability: $E(\|\widehat{\Phi}(X) - \Phi(X)\|^2)$ small

Ergodicity

- Which class of processes satisfy the following?

$$\forall i, \frac{1}{N} \sum_n f_i(x)(n) \rightarrow \mathbf{E}(f_i(X)) \quad (N \rightarrow \infty)$$

Ergodicity

- Which class of processes satisfy the following?

$$\forall i, \frac{1}{N} \sum_n f_i(x)(n) \rightarrow \mathbf{E}(f_i(X)) \quad (N \rightarrow \infty)$$

- These are called *ergodic* processes.
 - In statistical physics, a process with an *Integral Scale* is ergodic.
 - In statistics, *linear* processes are ergodic (provided the moments are finite).

Class-specific variability

- Besides deformations and stationary variability, object recognition is exposed to much more complex variability:



- clutter
- class-specific diversity